



Resources and Providers Reference

A resource is a key part of a recipe. A resource defines the actions that can be taken, such as when a package should be installed, whether a service should be enabled or restarted, which groups, users, or groups of users should be created, where to put a collection of files, what the name of a new directory should be, and so on. During a chef-client run, each resource is identified and then associated with a provider. The provider then does the work to complete the action defined by the resource. Each resource is processed in the same order as they appear in a recipe. The chef-client ensures that the same actions are taken the same way everywhere and that actions produce the same result every time. A resource is implemented within a recipe using Ruby.

Where a resource represents a piece of the system (and its desired state), a provider defines the steps that are needed to bring that piece of the system from its current state into the desired state. These steps are de-coupled from the request itself. The request is made in a recipe and is defined by a lightweight resource. The steps are then defined by a lightweight provider.

The `Chef::Platform` class maps providers to platforms (and platform versions). Ohai, as part of every chef-client run, verifies the `platform` and `platform_version` attributes on each node. The chef-client then uses those values to identify the correct provider, build an instance of that provider, identify the current state of the resource, do the specified action, and then mark the resource as updated (if changes were made). For example, given the following resource:

```
directory "/tmp/folder" do
  owner "root"
  group "root"
  mode 0755
  action :create
end
```

The chef-client will look up the provider for the `directory` resource, which happens to be `Chef::Provider::Directory`, call `load_current_resource` to create a new resource called `directory["/tmp/folder"]`, and then, based on the current state of the directory, do the specified action, which in this case is to create a directory called `/tmp/folder`. If the directory already exists, nothing will happen. If the directory was changed in any way, the resource is marked as updated.

This reference describes each of the resources available to the chef-client, including the list of actions available for the resource, the attributes that can be used, the providers that will do the work (and the provider's shortcut resource name), and examples of using each resource.

Common Functionality for all Resources

The attributes and actions in this section apply to all resources.

Actions

The following actions are common to every resource:

Action	Description
<code>:nothing</code>	Use to do nothing. In the absence of another default action, <code>nothing</code> is the default. This action can be useful to specify a resource so that it can be notified of other actions.

Examples

The following examples show how to use common actions in a recipe.

Use the `:nothing` action

```
service "memcached" do
  action :nothing
  supports :status => true, :start => true, :stop => true, :restart => true
end
```

Attributes

The following attributes are common to every resource:

Parameter	Description
<code>ignore_failure</code>	Use to continue running a recipe if a resource fails for any reason. Default value: <code>false</code> .
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>retries</code>	Use to specify the number of times to catch exceptions and retry the resource. Default value: <code>0</code> .
<code>retry_delay</code>	Use to specify the retry delay (in seconds). Default value: <code>2</code> .
<code>supports</code>	Use to specify a hash of options that contains hints about the capabilities of a resource. The chef-client may use

these hints to help identify the correct provider. This attribute is only used by a small number of providers, including `User` and `Service`.

Examples

The following examples show how to use common attributes in a recipe.

Use the `ignore_failure` common attribute

```
gem_package "syntax" do
  action :install
  ignore_failure true
end
```

Use the `provider` common attribute

```
package "some_package" do
  provider Chef::Provider::Package::Rubygems
end
```

Use the `supports` common attribute

```
service "apache" do
  supports :restart => true, :reload => true
  action :enable
end
```

Use the `supports` and `providers` common attributes

```
service "some_service" do
  provider Chef::Provider::Service::Upstart
  supports :status => true, :restart => true, :reload => true
  action [ :enable, :start ]
end
```

Guards

A guard can be used to evaluate the state of a node during the execution phase of the chef-client run. Based on the results of this evaluation, a guard is then used to tell the chef-client if it should continue executing a resource. A guard accepts either a string value or a Ruby block value:

- A string is executed as a shell command. If the command returns `0`, the guard is applied. If the command returns any other value, then the guard is not applied.
- A block is executed as Ruby code that must return either `true` or `false`. If the block returns `true`, the guard is applied. If the block returns `false`, the guard is not applied.

A guard is useful for ensuring that a resource is idempotent by allowing a resource to test for the desired state as it is being executed, and then if the desired state is present, for the chef-client to do nothing.

Attributes

The following guards can be used to define a condition to be evaluated during the execution phase of the chef-client run:

Guard	Description
<code>not_if</code>	Use to prevent a resource from executing when the condition returns <code>true</code> .
<code>only_if</code>	Use to allow a resource to execute only if the condition returns <code>true</code> .

Arguments

The following arguments can be used with the `not_if` or `only_if` guard:

Argument	Description
<code>:user</code>	Use to specify the user that a command will run as. For example: <pre>not_if "grep adam /etc/passwd", :user => 'adam'</pre>
<code>:group</code>	Use to specify the group that a command will run as. For example: <pre>not_if "grep adam /etc/passwd", :group => 'adam'</pre>
<code>:environment</code>	Use to specify a Hash of environment variables to be set. For example: <pre>not_if "grep adam /etc/passwd", :environment => { 'HOME' => "/home/adam" }</pre>
<code>:cwd</code>	Use to set the current working directory before running a command. For example: <pre>not_if "grep adam passwd", :cwd => '/etc'</pre>

`:timeout` Use to set a timeout for a command. For example:

```
not_if "sleep 10000", :timeout => 10
```

not_if Examples

The following examples show how to use `not_if` as a condition in a recipe:

Create a file, but not if an attribute has a specific value

The following example shows how to use the `not_if` condition to create a file based on a template and using the presence of an attribute on the node to specify the condition:

```
template "/tmp/somefile" do
  mode 00644
  source "somefile.erb"
  not_if { node[:some_value] }
end
```

Create a file with a Ruby block, but not if “/etc/passwd” exists

The following example shows how to use the `not_if` condition to create a file based on a template and then Ruby code to specify the condition:

```
template "/tmp/somefile" do
  mode 00644
  source "somefile.erb"
  not_if do
    File.exists?("/etc/passwd")
  end
end
```

Create a file with Ruby block that has curly braces, but not if “/etc/passwd” exists

The following example shows how to use the `not_if` condition to create a file based on a template and using a Ruby block (with curly braces) to specify the condition:

```
template "/tmp/somefile" do
  mode 00644
  source "somefile.erb"
  not_if { File.exists?("/etc/passwd") }
end
```

Create a file using a string, but not if “/etc/passwd” exists

The following example shows how to use the `not_if` condition to create a file based on a template and using a string to specify the condition:

```
template "/tmp/somefile" do
  mode 00644
  source "somefile.erb"
  not_if "test -f /etc/passwd"
end
```

Install a file from a remote location using bash

The following is an example of how to install the `foo123` module for Nginx. This module adds shell-style functionality to an Nginx configuration file and does the following:

- Declares three variables
- Gets the Nginx file from a remote location
- Installs the file using Bash to the path specified by the `src_filepath` variable

the following code sample is similar to the ``upload_progress_module`` recipe in the ``nginx`` cookbook

```
src_filename = "foo123-nginx-module-v#{node['nginx']['foo123']['version']}.tar.gz"
src_filepath = "#{Chef::Config['file_cache_path']}/#{src_filename}"
extract_path = "#{Chef::Config['file_cache_path']}/nginx_foo123_module/#{node['nginx']['foo123']['checksum']}"

remote_file src_filepath do
  source node['nginx']['foo123']['url']
  checksum node['nginx']['foo123']['checksum']
  owner 'root'
  group 'root'
  mode 00644
end

bash 'extract_module' do
  cwd ::File.dirname(src_filepath)
  code <<-EOH
  mkdir -p #{extract_path}
  tar xzf #{src_filename} -C #{extract_path}
  mv #{extract_path}/*/* #{extract_path}/
  EOH
  not_if { ::File.exists?(extract_path) }
end
```

only_if Examples

The following examples show how to use `only_if` as a condition in a recipe:

Create a file, but only if an attribute has a specific value

The following example shows how to use the `only_if` condition to create a file based on a template and using the presence of an attribute on the node to specify the condition:

```
template "/tmp/somefile" do
  mode 00644
  source "somefile.erb"
  only_if { node[:some_value] }
end
```

Create a file with a Ruby block, but only if "/etc/passwd" does not exist

The following example shows how to use the `only_if` condition to create a file based on a template, and then use Ruby to specify a condition:

```
template "/tmp/somefile" do
  mode 00644
  source "somefile.erb"
  only_if do ! File.exists?("/etc/passwd") end
end
```

Create a file using a string, but only if "/etc/passwd" exists

The following example shows how to use the `only_if` condition to create a file based on a template and using a string to specify the condition:

```
template "/tmp/somefile" do
  mode 00644
  source "somefile.erb"
  only_if "test -f /etc/passwd"
end
```

Lazy Attribute Evaluation

In some cases, the value for an attribute cannot be known until the execution phase of a chef-client run. In this situation, using lazy evaluation of attribute values can be helpful. Instead of an attribute being assigned a value, it may instead be assigned a code block. The syntax for using lazy evaluation is as follows:

```
attribute_name lazy { code_block }
```

where `lazy` is used to tell the chef-client to evaluate the contents of the code block later on in the resource evaluation process (instead of immediately) and `{ code_block }` is arbitrary Ruby code that provides the value.

For example, a resource that is not doing lazy evaluation:

```
template "template_name" do
  # some attributes
  path "/foo/bar"
end
```

and a resource that is doing lazy evaluation:

```
template "template_name" do
  # some attributes
  path lazy { " some Ruby code " }
end
```

In the previous examples, the first resource uses the value `/foo/bar` and the second resource uses the value provided by the code block, as long as the contents of that code block are a valid resource attribute.

Notifications

The following notifications can be used with any resource:

Notification	Description
<code>notifies</code>	Use to notify another resource to take an action if this resource's state changes for any reason.
<code>subscribes</code>	Use to take action on this resource if another resource's state changes. This is similar to <code>notifies</code> , but reversed.

Notifications Timers

The following timers can be used to define when a notification is triggered:

Timer	Description
<code>:delayed</code>	Use to specify that a notification should be queued up and then executed at the very end of a chef-client run.
<code>:immediately</code>	Use to specify that a notification be run immediately.

Notifies Syntax

The basic syntax of a `notifies` notification is:

```
resource "name" do
  notifies :notification, "resource_type[resource_name]", :timer
end
```

Examples

The following examples show how to use the `notifies` notification in a recipe.

Delay notifications

```
template "/etc/nagios3/configures-nagios.conf" do
  # other parameters
  notifies :run, "execute[test-nagios-config]", :delayed
end
```

Notify immediately

By default, notifications are `:delayed`, that is they are queued up as they are triggered, and then executed at the very end of a chef-client run. To run an action immediately, use `:immediately`:

```
template "/etc/nagios3/configures-nagios.conf" do
  # other parameters
  notifies :run, "execute[test-nagios-config]", :immediately
end
```

and then the chef-client would immediately run the following:

```
execute "test-nagios-config" do
  command "nagios3 --verify-config"
  action :nothing
end
```

Enable a service after a restart or reload

```
service "apache" do
  supports :restart => true, :reload => true
  action :enable
end
```

Notify multiple resources

```
template "/etc/chef/server.rb" do
  source "server.rb.erb"
  owner "root"
  group "root"
  mode "644"
  notifies :restart, "service[chef-solr]", :delayed
  notifies :restart, "service[chef-solr-indexer]", :delayed
  notifies :restart, "service[chef-server]", :delayed
end
```

Notify in a specific order

To notify multiple resources, and then have these resources run in a certain order, do something like the following:

```
execute 'foo' do
  command '...'
  notifies :run, 'template[baz]', :immediately
  notifies :install, 'package[bar]', :immediately
  notifies :run, 'execute[final]', :immediately
end

template 'baz' do
  ...
  notifies :run, 'execute[restart_baz]', :immediately
end

package 'bar'

execute 'restart_baz'

execute 'final' do
  command '...'
end
```

where the sequencing will be in the same order as the resources are listed in the recipe: `execute 'foo'`, `template 'baz'`, `execute [restart_baz]`, `package 'bar'`, and `execute 'final'`.

Reload a service

```
template "/tmp/somefile" do
  mode "0644"
  source "somefile.erb"
  notifies :reload, "service[apache]"
end
```

Restart a service when a template is modified

```
template "/etc/www/configures-apache.conf" do
  notifies :restart, "service[apache]"
end
```

```
end
```

Send notifications to multiple resources

To send notifications to multiple resources, just use multiple attributes. Multiple attributes will get sent to the notified resources in the order specified.

```
template "/etc/netatalk/netatalk.conf" do
  notifies :restart, "service[afpd]", :immediately
  notifies :restart, "service[cnid]", :immediately
end

service "afpd"
service "cnid"
```

Execute a command using a template

The following example shows how to set up IPv4 packet forwarding using the **execute** resource to run a command named "forward_ipv4" that uses a template defined by the **template** resource:

```
execute "forward_ipv4" do
  command "echo > /proc/.../ipv4/ip_forward"
  action :nothing
end

template "/etc/file_name.conf" do
  source "routing/file_name.conf.erb"
  notifies :run, 'execute[forward_ipv4]', :delayed
end
```

where the `command` attribute for the **execute** resource contains the command that is to be run and the `source` attribute for the **template** resource specifies which template to use. The `notifies` attribute for the **template** specifies that the `execute[forward_ipv4]` (which is defined by the **execute** resource) should be queued up and run at the end of the chef-client run.

Restart a service, and then notify a different service

The following example shows how start a service named "example_service" and immediately notify the Nginx service to restart.

```
service "example_service" do
  action :start
  provider Chef::Provider::Service::Init
  notifies :restart, "service[nginx]", :immediately
end
```

where by using the default `provider` for the **service**, the recipe is telling the chef-client to determine the specific provider to be used during the chef-client run based on the platform of the node on which the recipe will run.

Notify when a remote source changes

```
remote_file "/tmp/couch.png" do
  source "http://couchdb.apache.org/img/sketch.png"
  action :nothing
end

http_request "HEAD http://couchdb.apache.org/img/sketch.png" do
  message ""
  url "http://couchdb.apache.org/img/sketch.png"
  action :head
  if File.exists?("/tmp/couch.png")
    headers "If-Modified-Since" => File.mtime("/tmp/couch.png").httpdate
  end
  notifies :create, "remote_file[/tmp/couch.png]", :immediately
end
```

Subscribes Syntax

The basic syntax of a `subscribes` notification is:

```
resource "name" do
  subscribes :notification, "resource_type[resource_name]", :timer
end
```

Examples

The following examples show how to use the `subscribes` notification in a recipe.

Prevent restart and reconfigure if configuration is broken

Use the `:nothing` common action to prevent an application from restarting, and then use the `subscribes` notification to ask the broken configuration to be reconfigured immediately:

```
execute "test-nagios-config" do
  command "nagios3 --verify-config"
  action :nothing
  subscribes :run, "template[/etc/nagios3/configures-nagios.conf]", :immediately
end
```

Reload a service using a template

To reload a service based on a template, use the **template** and **service** resources together in the same recipe, similar to the following:

```

template "/tmp/somefile" do
  mode "0644"
  source "somefile.erb"
end

service "apache" do
  supports :restart => true, :reload => true
  action :enable
  subscribes :reload, "template[/tmp/somefile]", :immediately
end

```

where the `subscribes` notification is used to reload the service using the template specified by the **template** resource.

Stash a file in a data bag

The following example shows how to use the **ruby_block** resource to stash a BitTorrent file in a data bag so that it can be distributed to nodes in the organization.

```

# the following code sample comes from the ``seed`` recipe in the following cookbook: https://github.com
ruby_block "share the torrent file" do
  block do
    f = File.open(node['bittorrent']['torrent'], 'rb')
    #read the .torrent file and base64 encode it
    enc = Base64.encode64(f.read)
    data = {
      'id' => bittorrent_item_id(node['bittorrent']['file']),
      'seed' => node.ipaddress,
      'torrent' => enc
    }
    item = Chef::DataBagItem.new
    item.data_bag('bittorrent')
    item.raw_data = data
    item.save
  end
  action :nothing
  subscribes :create, "bittorrent_torrent[#{node['bittorrent']['torrent']}]"
end

```

Relative Paths

The following relative paths can be used with any resource:

Relative Path	Description
<code>#</code>	Use to return the <code>~</code> path in Linux and Mac OS X or the <code>%HOMEPATH%</code> in Microsoft Windows.
<code>{ENV['HOME']}</code>	

Examples

```

template "#{ENV['HOME']}/chef-getting-started.txt" do
  source "chef-getting-started.txt.erb"
  mode 00644
end

```

Run Resources from the Resource Collection

The chef-client processes recipes in two phases:

1. First, each resource in the node object is identified and a resource collection is built. All recipes are loaded in a specific order, and then the actions specified within each of them are identified.
2. Next, the chef-client configures the system based on the order of the resources in the resource collection. Each resource is mapped to a provider, which then examines the node and then does the steps necessary to complete the action.

Sometimes, it may be necessary to ensure that a specific resource is run during the phase that builds the resource collection. For example:

- A resource may need to run first so that it can download a package that will be used by other resources in the resource collection
- Several resources need to install a package; rather than having the package installer run several times, it can be configured to run only once

To support these types of uses cases, it is possible to tell the chef-client to run a resource at the beginning and/or the end of the resource collection phase. Effectively, run a resource before all other resources are added to the resource collection and/or after all resources have been added, but before the chef-client configures the system.

Before other resources

To run a resource at the start of the resource collection phase of the chef-client run, set up a `Chef::Resource` object, and then call the method that runs the action.

Update a package cache

It is important to make sure that an operating system's package cache is up to date before installing packages, otherwise there may be references to versions that no longer exist. For example, on Debian or Ubuntu systems, the APT cache needs to be updated. Use code similar to the following:

```
e = execute "apt-get update" do
  action :nothing
end

e.run_action(:run)
```

where `e` is created as a `Chef::Resource::Execute` Ruby object. The `action` attribute is set to `:nothing` so that the `run_action` method can be used to tell the chef-client to run the specified command. The **apt** (for Debian and Ubuntu) and **pacman** (for Arch Linux) cookbooks can be used for this purpose. The preceding recipe can be placed at the top of a node's run list to ensure it is run before the chef-client tries to install any packages.

An anti-pattern

Unfortunately, resources that are executed when the resource collection is being built cannot notify any resource that has yet to be added to the resource collection. For example:

```
execute "ifconfig"

p = package 'vim-enhanced' do
  action :nothing
  notifies :run, "execute[ifconfig]", :immediately
end
p.run_action(:install)
```

In some cases, the better approach may be to install the package before the resource collection is built to ensure that it is available to other resources later on. Or, something like the following can be used:

```
p = package "foo" do
  #parameters
end
p.run_action(:install)

if p.updated_by_last_action?
  #Call the resource that we want to "notify"
end
```

After the resource collection is built

To run a resource at the end of the resource collection phase of the chef-client run, use the `:delayed` timer on a notification.

Windows File Security

To support Microsoft Windows security, the **template**, **file**, **remote_file**, **cookbook_file**, **directory**, and **remote_directory** resources support the use of inheritance and access control lists (ACLs) within recipes.

Access Control Lists (ACLs)

The `rights` attribute can be used in a recipe to manage access control lists (ACLs), which allow permissions to be given to multiple users and groups. The syntax for the `rights` attribute is as follows:

```
rights permission, principal, option_type => value
```

where

- `permission` is used to specify which rights will be granted to the `principal`. The possible values are: `:read`, `:write`, `:full_control`, `:modify`, and `:deny`. These permissions are cumulative. If `:write` is specified, then it includes `:read`. If `:full_control` is specified, then it includes both `:write` and `:read`. If `:deny` is specified, then the user or group will not have rights to the object. (For those who know the Microsoft Windows API: `:read` corresponds to `GENERIC_READ` and `GENERIC_EXECUTE`; `:write` corresponds to `GENERIC_WRITE`, `GENERIC_READ`, and `GENERIC_EXECUTE`; `:full_control` corresponds to `GENERIC_ALL`, which allows a user to change the owner and other metadata about a file.)
- `principal` is used to specify a group or user name. This is identical to what is entered in the login box for Microsoft Windows, such as `user_name`, `domain\user_name`, or `user_name@fully_qualified_domain_name`. The chef-client does not need to know if a principal is a user or a group.
- `option_type` is a hash that contains advanced rights options. For example, the rights to a directory that only applies to the first level of children might look something like: `rights :write, "domain\group_name", :one_level_deep => true`. Possible option types:

Option Type	Description
<code>:applies_to_children</code>	Use to specify how permissions are applied to children. Possible values: <code>true</code> to inherit both child directories and files; <code>false</code> to not inherit any child directories or files; <code>:containers_only</code> to inherit only child directories (and not files); <code>:objects_only</code> to recursively inherit files (and not child directories).
<code>:applies_to_self</code>	Indicates whether a permission is applied to the parent directory. Possible values: <code>true</code> to apply to the parent directory or file and its children; <code>false</code> to not apply only to child directories and files.
<code>:one_level_deep</code>	Indicates the depth to which permissions will be applied. Possible values: <code>true</code> to apply only to the first level of children; <code>false</code> to apply to all children.

The `rights` attribute can be used as many times as necessary; the chef-client will apply them to the file or directory as required. For

example:

```
resource "x.txt" do
  rights :read, "Everyone"
  rights :write, "domain\group"
  rights :full_control, "group_name_or_user_name"
  rights :full_control, "user_name", :applies_to_children => true
end
```

or:

```
rights :read, ["Administrators", "Everyone"]
rights :deny, ["Julian", "Lewis"]
rights :full_control, "Users", :applies_to_children => true
rights :write, "Sally", :applies_to_children => :containers_only, :applies_to_self => false, :one_level_d
```

Some other important things to know when using the `rights` attribute:

- Order independence. It doesn't matter if `rights :deny, ["Julian", "Lewis"]` is placed before or after `rights :read, ["Julian", "Lewis"]`, both Julian and Lewis will be unable to read the document.
- Only inherited rights remain. All existing explicit rights on the object are removed and replaced.
- If rights are not specified, nothing will be changed. The chef-client does not clear out the rights on a file or directory if rights are not specified.
- Changing inherited rights can be expensive. Microsoft Windows will propagate rights to all children recursively due to inheritance. This is a normal aspect of Microsoft Windows, so consider the frequency with which this type of action is necessary and take steps to control this type of action if performance is the primary consideration.

Inheritance

By default, a file or directory inherits rights from its parent directory. Most of the time this is the preferred behavior, but sometimes it may be necessary to take steps to more specifically control rights. The `inherits` attribute can be used to specifically tell the chef-client to apply (or not apply) inherited rights from its parent directory.

For example, the following example specifies the rights for a directory:

```
directory 'C:\mordor' do
  rights :read, 'MORDOR\Minions'
  rights :full_control, 'MORDOR\Sauron'
end
```

and then the following example specifies how to use inheritance to deny access to the child directory:

```
directory 'C:\mordor\mount_doom' do
  rights :full_control, 'MORDOR\Sauron'
  inherits false # Sauron is the only person who should have any sort of access
end
```

If the `:deny` permission were to be used instead, something could slip through unless all users and groups were denied.

Another example also shows how to specify rights for a directory:

```
directory 'C:\mordor' do
  rights :read, 'MORDOR\Minions'
  rights :full_control, 'MORDOR\Sauron'
  rights :write, 'SHIRE\Frodo' # Who put that there I didn't put that there
end
```

but then not use the `inherits` attribute to deny those rights on a child directory:

```
directory 'C:\mordor\mount_doom' do
  rights :deny, 'MORDOR\Minions' # Oops, not specific enough
end
```

Because the `inherits` attribute is not specified, the chef-client will default it to `true`, which will ensure that security settings for existing files remain unchanged.

Resources

The following resources are platform resources with built-in providers:

- `apt_package` (based on the `package` resource)
- `bash`
- `chef_gem` (based on the `package` resource)
- `chef_handler` (available from the `chef_handler` cookbook)
- `cookbook_file`
- `cron`
- `csh`
- `deploy` (including `git` and `Subversion`)
- `directory`
- `dpkg_package` (based on the `package` resource)
- `easy_install_package` (based on the `package` resource)
- `env`
- `erl_call`

- `execute`
- `file`
- `freebsd_package` (based on the `package` resource)
- `gem_package` (based on the `package` resource)
- `git`
- `group`
- `http_request`
- `ifconfig`
- `ips_package` (based on the `package` resource)
- `link`
- `log`
- `macports_package` (based on the `package` resource)
- `mdadm`
- `mount`
- `ohai`
- `package`
- `pacman_package` (based on the `package` resource)
- `perl`
- `portage_package` (based on the `package` resource)
- `powershell_script`
- `python`
- `registry_key`
- `remote_directory`
- `remote_file`
- `rpm_package` (based on the `package` resource)
- `route`
- `ruby`
- `ruby_block`
- `script`
- `service`
- `smart_o_s_package` (based on the `package` resource)
- `solaris_package` (based on the `package` resource)
- `subversion`
- `template`
- `user`
- `yum` (based on the `package` resource)

See below for more information about each of these resources, their related actions and attributes, the providers they rely on, and examples of how these resources can be used in recipes.

apt_package

The `apt_package` resource is used to manage packages for the Debian and Ubuntu platforms.

Note

In many cases, it is better to use the `package` resource instead of this one. This is because when the `package` resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the `package` resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

Syntax

The syntax for using the `apt_package` resource in a recipe is as follows:

```
apt_package "name" do
  attribute "value" # see attributes section below
  ...
  action :action # see actions section below
end
```

where

- `apt_package` tells the chef-client to use the `Chef::Provider::Apt` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.

<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:reconfig</code>	Use to reconfigure a package. This action requires a response file.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

Attributes

This resource has the following attributes:

Attribute	Description
<code>arch</code>	The architecture of the package that will be installed or upgraded. (This value can also be passed as part of the package name.)
<code>options</code>	One (or more) additional options that are passed to the command. For example, common apt-get directives, such as <code>--no-install-recommends</code> . See the apt-get man page for the full list.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The direct path to a dpkg or deb package.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Apt</code>	<code>apt_package</code>	The provider that is used with the Debian and Ubuntu platforms.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package using package manager

```
apt_package "name of package" do
  action :install
end
```

Install a package using local file

```
apt_package "jwhois" do
  action :install
  source '/path/to/jwhois.deb'
end
```

Install without using recommend packages as a dependency

```
package "apache2" do
  options "--no-install-recommends"
end
```

bash

The **bash** resource is used to execute scripts using the Bash interpreter and includes all of the actions and attributes that are available to the **execute** resource.

Note

The **bash** script resource (which is based on the **script** resource) is different from the **ruby_block** resource because Ruby code that is run with this resource is created as a temporary file and executed like other script resources, rather than run inline. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the `not_if` and `only_if` meta parameters to guard the use of this resource for idempotence.

Syntax

The syntax for using the **bash** resource in a recipe is as follows:

```
bash "name" do
  attribute "value" # see attributes section below
  ...
  action :action # see actions section below
end
```

where

- **bash** tells the chef-client to use the `Chef::Resource::Script::Bash` provider during the chef-client run
- **name** is the name of the resource block; when the `command` attribute is not specified as part of a recipe, **name** is also the name of the command to be executed
- **attribute** is zero (or more) of the attributes that are available for this resource
- **:action** is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:run</code>	Default. Use to run a script.
<code>:nothing</code>	Indicates that the command should not be run. This action is used to specify that a command is run only when another resource notifies it.

Attributes

This resource has the following attributes:

Attribute	Description
<code>code</code>	A quoted (" ") string of code to be executed.
<code>command</code>	The name of the command to be executed. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>creates</code>	Indicates that a command to create a file will not be run when that file already exists.
<code>cwd</code>	The current working directory.
<code>environment</code>	A Hash of environment variables in the form of <code>{ "ENV_VARIABLE" => "VALUE" }</code> . (These variables must exist for a command to be run successfully.)
<code>flags</code>	One (or more) command line flags that are passed to the interpreter when a command is invoked.
<code>group</code>	The group name or group ID that must be changed before running a command.
<code>path</code>	An array of paths to use when searching for a command. These paths are not added to the command's environment <code>\$PATH</code> . The default value uses the system path.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>returns</code>	The return value for a command. This may be an array of accepted values. An exception is raised when the return value(s) do not match. Default value: <code>0</code> .
<code>timeout</code>	The amount of time (in seconds) a command will wait before timing out. Default value: <code>3600</code> .
<code>user</code>	The user name or user ID that should be changed before running a command.
<code>umask</code>	The file mode creation mask, or umask.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Script</code>	<code>script</code>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<code>Chef::Provider::Script::Bash</code>	<code>bash</code>	The provider that is used with the Bash command interpreter.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Use a named provider to run a script

```
bash "install_something" do
  user "root"
  cwd  "/tmp"
  code <<-EOH
  wget http://www.example.com/tarball.tar.gz
  tar -zxf tarball.tar.gz
  cd tarball
  ./configure
  make
  make install
  EOH
end
```

Install a file from a remote location using bash

The following is an example of how to install the `foo123` module for Nginx. This module adds shell-style functionality to an Nginx configuration file and does the following:

- Declares three variables
- Gets the Nginx file from a remote location
- Installs the file using Bash to the path specified by the `src_filepath` variable

the following code sample is similar to the `upload_progress_module` recipe in the `nginx` cookbook.

```
src_filename = "foo123-nginx-module-v#{node['nginx']['foo123']['version']}.tar.gz"
src_filepath = "#{Chef::Config['file_cache_path']}/#{src_filename}"
extract_path = "#{Chef::Config['file_cache_path']}/nginx_foo123_module/#{node['nginx']['foo123']['checksum']}"

remote_file src_filepath do
  source node['nginx']['foo123']['url']
  checksum node['nginx']['foo123']['checksum']
  owner 'root'
  group 'root'
  mode 00644
end

bash 'extract_module' do
  cwd ::File.dirname(src_filepath)
  code <<-EOH
  mkdir -p #{extract_path}
  tar xzf #{src_filename} -C #{extract_path}
  mv #{extract_path}/*/* #{extract_path}/
  EOH
  not_if { ::File.exists?(extract_path) }
end
```

Install an application from git using bash

The following example shows how Bash can be used to install a plug-in for `rbenv` named "ruby-build", which is located in git version source control. First, the application is synchronized, and then Bash changes its working directory to the location in which "ruby-build" is located, and then runs a command.

```
git "#{Chef::Config[:file_cache_path]}/ruby-build" do
  repository "git://github.com/sstephenson/ruby-build.git"
  reference "master"
  action :sync
end

bash "install_ruby_build" do
  cwd "#{Chef::Config[:file_cache_path]}/ruby-build"
  user "rbenv"
  group "rbenv"
  code <<-EOH
  ./install.sh
  EOH
  environment 'PREFIX' => "/usr/local"
end
```

To read more about `ruby-build`, see here: <https://github.com/sstephenson/ruby-build>.

Store certain settings

The following recipe shows how an attributes file can be used to store certain settings. An attributes file is located in the `attributes/` directory in the same cookbook as the recipe which calls the attributes file. In this example, the attributes file specifies certain settings for Python that are then used across all nodes against which this recipe will run.

Python packages have versions, installation directories, URLs, and checksum files. An attributes file that exists to support this type of recipe would include settings like the following:

```
default['python']['version'] = '2.7.1'

if python['install_method'] == 'package'
  default['python']['prefix_dir'] = '/usr'
else
  default['python']['prefix_dir'] = '/usr/local'
end

default['python']['url'] = 'http://www.python.org/ftp/python'
default['python']['checksum'] = '80e387...85fd61'
```

and then the methods in the recipe may refer to these values. A recipe that is used to install Python will need to do the following:

- Identify each package to be installed (implied in this example, not shown)
- Define variables for the package version and the `install_path`
- Get the package from a remote location, but only if the package does not already exist on the target system
- Use the **bash** resource to install the package on the node, but only when the package is not already installed

the following code sample comes from the ``oc-nginx`` cookbook on [github]: <https://github.com/cookbook>

```
version = node['python']['version']
install_path = "#{node['python']['prefix_dir']}/lib/python#{version.split(/(\d+\.\d+)/)[1]}"

remote_file "#{Chef::Config[:file_cache_path]}/Python-#{version}.tar.bz2" do
  source "#{node['python']['url']}/#{version}/Python-#{version}.tar.bz2"
  checksum node['python']['checksum']
  mode "0644"
  not_if { ::File.exists?(install_path) }
end

bash "build-and-install-python" do
  cwd Chef::Config[:file_cache_path]
  code <<-EOF
  tar -jxvf Python-#{version}.tar.bz2
  (cd Python-#{version} && ./configure #{configure_options})
  (cd Python-#{version} && make && make install)
  EOF
  not_if { ::File.exists?(install_path) }
end
```

chef_gem

The **chef_gem** resource is used to install a gem only for the instance of Ruby that is dedicated to the chef-client. When a package is installed from a local file, it must be added to the node using the **remote_file** or **cookbook_file** resources.

The **chef_gem** resource works with all of the same attributes and options as the **gem_package** resource, but does not accept the `gem_binary` attribute because it always uses the `CurrentGemEnvironment` under which the chef-client is running. In addition to performing actions similar to the **gem_package** resource, the **chef_gem** resource does the following:

- Runs its actions immediately, before convergence, allowing a gem to be used in a recipe immediately after it is installed
- Runs `Gem.clear_paths` after the action, ensuring that gem is aware of changes so that it can be required immediately after it is installed

Warning

The **chef_gem** and **gem_package** resources are both used to install Ruby gems. For any machine on which the chef-client is installed, there are two instances of Ruby. One is the standard, system-wide instance of Ruby and the other is a dedicated instance that is available only to the chef-client. Use the **chef_gem** resource to install gems into the instance of Ruby that is dedicated to the chef-client. Use the **gem_package** resource to install all other gems (i.e. install gems system-wide).

Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

Syntax

The syntax for using the **chef_gem** resource in a recipe is as follows:

```
chef_gem "name" do
  attribute "value" # see attributes section below
  ...
  action :action # see actions section below
end
```

where

- **chef_gem** tells the chef-client to use the `Chef::Provider::Rubygems` provider during the chef-client run
- **name** is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, **name** is also the name of the package
- **attribute** is zero (or more) of the attributes that are available for this resource
- **:action** is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:reconfig</code>	Use to reconfigure a package. This action requires a response file.

<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

Attributes

This resource has the following attributes:

Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Rubygems</code>	<code>chef_gem</code>	Can be used with the <code>options</code> attribute.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a gems file for use in recipes

```
chef_gem "right_aws" do
  action :install
end

require 'right_aws'
```

Install MySQL for Chef

```
execute "apt-get update" do
  ignore_failure true
  action :nothing
end.run_action(:run) if node['platform_family'] == "debian"

node.set['build_essential']['compiletime'] = true
include_recipe "build-essential"
include_recipe "mysql::client"

node['mysql']['client']['packages'].each do |mysql_pack|
  resources("package[#{mysql_pack}]").run_action(:install)
end

chef_gem "mysql"
```

chef_handler

A **resource** is a key part of a **recipe** that defines the actions that can be taken against a piece of the system. These actions are identified during each **chef-client** run as the resource collection is compiled. Once identified, each resource (in turn) is mapped to a provider, which then configures each piece of the system.

The **chef_handler** resource is used to enable handlers during a chef-client run. The resource allows arguments to be passed to the chef-client, which then applies the conditions defined by the custom handler to the node attribute data collected during the chef-client run, and then processes the handler based on that data.

The **chef_handler** resource is typically defined early in a node's run-list (often being the first item). This ensures that all of the handlers will be available for the entire chef-client run.

The **chef_handler** resource is included with the **chef_handler** cookbook. This cookbook defines the the resource itself and also provides the location in which the chef-client looks for custom handlers. All custom handlers should be added to the `files/default/handlers` directory in the **chef_handler** cookbook.

Handler Types

There are three types of handlers:

Handler	Description
exception	An exception handler is used to identify situations that have caused a chef-client run to fail. An exception handler can be loaded at the start of a chef-client run by adding a recipe that contains the <code>chef_handler</code> resource to a node's run-list. An exception handler runs when the <code>failed?</code> property for the <code>run_status</code> object returns <code>true</code> .
report	A report handler is used when a chef-client run succeeds and reports back on certain details about that chef-client run. A report handler can be loaded at the start of a chef-client run by adding a recipe that contains the <code>chef_handler</code> resource to a node's run-list. A report handler runs when the <code>success?</code> property for the <code>run_status</code> object returns <code>true</code> .
start	A start handler is used to run events at the beginning of the chef-client run. A start handler can be loaded at the start of a chef-client run by adding the start handler to the <code>start_handlers</code> setting in the <code>client.rb</code> file or by installing the gem that contains the start handler by using the <code>chef_gem</code> resource in a recipe in the <code>chef-client</code> cookbook. (A start handler may not be loaded using the <code>chef_handler</code> resource.)

Exception / Report

Exception and report handlers are used to trigger certain behaviors in response to specific situations, typically identified during a chef-client run.

- An exception handler is used to trigger behaviors when a defined aspect of a chef-client run fails.
- A report handler is used to trigger behaviors when a defined aspect of a chef-client run is successful.

Both types of handlers can be used to gather data about a chef-client run and can provide rich levels of data about all types of usage, which can be used later for trending and analysis across the entire organization.

Exception and report handlers are made available to the chef-client run in one of the following ways:

- By adding the `chef_handler` resource to a recipe, and then adding that recipe to the run-list for a node. (The `chef_handler` resource is available from the `chef_handler` cookbook.)
- By adding the handler to one of the following settings in the node's `client.rb` file: `exception_handlers` and/or `report_handlers`

The `chef_handler` resource allows exception and report handlers to be enabled from within recipes, which can then added to the run-list for any node on which the exception or report handler should run. The `chef_handler` resource is available from the `chef_handler` cookbook.

To use the `chef_handler` resource in a recipe, add code similar to the following:

```
chef_handler "name_of_handler" do
  source "/path/to/handler/handler_name"
  action :enable
end
```

For example, a handler for Growl needs to be enabled at the beginning of the chef-client run:

```
.. code-block:: ruby

  chef_gem "chef-handler-growl"
```

and then is activated in a recipe by using the `chef_handler` resource:

```
chef_handler "Chef::Handler::Growl" do
  source "chef/handler/growl"
  action :enable
end
```

Start

A start handler is not loaded into the chef-client run from a recipe, but is instead listed in the `client.rb` file using the `start_handlers` attribute. The start handler must be installed on the node and be available to the chef-client prior to the start of the chef-client run. Use the `chef-client` cookbook to install the start handler.

Start handlers are made available to the chef-client run in one of the following ways:

- By adding a start handler to the `chef-client` cookbook, which installs the handler on the node so that it is available to the chef-client at the start of the chef-client run
- By adding the handler to one of the following settings in the node's `client.rb` file: `start_handlers`

The `chef-client` cookbook can be configured to automatically install and configure gems that are required by a start handler. For example:

```
node.set['chef_client']['load_gems']['chef-reporting'] = {
  :require_name => 'chef-reporting',
  :action => :install
}

node.set['chef_client']['start_handlers'] = [
  {
    :class => "Chef::Reporting::StartHandler",
```



```
      :arguments => []
    }
  ]

  include_recipe "chef-client::config"
```

Syntax

The syntax for using the **chef_handler** resource in a recipe is as follows:

```
chef_handler "name" do
  attribute "value" # see attributes section below
  ...
  action :action # see actions section below
end
```

Actions

This resource has the following actions:

Action	Description
:enable	Use to enable the handler for the current chef-client run on the current node.
:disable	Use to disable the handler for the current chef-client run on the current node.

Attributes

This resource has the following attributes:

Attribute	Description
class_name	The name of the handler class. This can be module name-spaced.
source	The full path to the handler file or the path to a gem (if the handler ships as part of a Ruby gem).
arguments	An array of arguments that are passed to the initializer for the handler class. Default value: []. For example: arguments :key1 => 'val1' or: arguments [:key1 => 'val1', :key2 => 'val2']
supports	The type of handler. Possible values: :exception , :report , :both (exception and report handlers), or :start . Default value: { :report => true, :exception => true }.

Custom Handlers

A custom handler can be created to support any situation. The easiest way to build a custom handler:

- 1. Download the **chef_handler** cookbook
- 2. Create a custom handler
- 3. Write a recipe using the **chef_handler** resource
- 4. Add that recipe to a node's run-list, often as the first recipe in that run-list

Syntax

The syntax for a handler can vary, depending on what the the situations the handler is being asked to track, the type of handler being used, and so on. All custom exception and report handlers are defined using Ruby and must be a subclass of the `Chef::Handler` class.

```
require "chef/log"

module ModuleName
  class HandlerName < Chef::Handler
    def report
      # Ruby code goes here
    end
  end
end
```

where:

- `require` ensures that the logging functionality of the chef-client is available to the handler
- `ModuleName` is the name of the module as it exists within the `Chef` library
- `HandlerName` is the name of the handler as it is used in a recipe
- `report` is an interface that is used to define the custom handler

For example, the following shows a custom handler that sends an email that contains the exception data when a chef-client run fails:

```
require "net/smtp"

module OrgName
  class SendEmail < Chef::Handler
    def report
      message = "From: sender_name <sender@example.com>\n"
      message << "To: recipient_address <recipient@example.com>\n"
      message << "Subject: chef-client Run Failed\n"
      message << "Date: #{Time.now.rfc2822}\n\n"
      message << "Chef run failed on #{node.name}\n"
      message << "#{run_status.formatted_exception}\n"
      message << Array(backtrace).join("\n")
      Net::SMTP.start('your.smtp.server', 25) do |smtp|
        smtp.send_message message, 'sender@example', 'recipient@example'
      end
    end
  end
end
```

and then is used in a recipe like:

```
send_email "blah"
# recipe code
end
```

report Interface

The `report` interface is used to define how a handler will behave and is a required part of any custom handler. The syntax for the `report` interface is as follows:

```
def report
  # Ruby code
end
```

The Ruby code used to define a custom handler will vary significantly from handler to handler. The chef-client includes two default handlers: `error_report` and `json_file`. Their use of the `report` interface is shown below.

The `error_report` handler:

```
require 'chef/handler'
require 'chef/resource/directory'

class Chef
  class Handler
    class ErrorReport < ::Chef::Handler
      def report
        Chef::FileCache.store("failed-run-data.json", Chef::JSONCompat.to_json_pretty(data), 0640)
        Chef::Log.fatal("Saving node information to #{Chef::FileCache.load("failed-run-data.json", false)}")
      end
    end
  end
end
```

The `json_file` handler:

```
require 'chef/handler'
require 'chef/resource/directory'

class Chef
  class Handler
    class JsonFile < ::Chef::Handler
      attr_reader :config
      def initialize(config={})
        @config = config
        @config[:path] ||= "/var/chef/reports"
      end
      def report
        if exception
          Chef::Log.error("Creating JSON exception report")
        else
          Chef::Log.info("Creating JSON run report")
        end
        build_report_dir
        savetime = Time.now.strftime("%Y%m%d%H%M%S")
        File.open(File.join(config[:path], "chef-run-report-#{savetime}.json"), "w") do |file|
          run_data = data
          run_data[:start_time] = run_data[:start_time].to_s
          run_data[:end_time] = run_data[:end_time].to_s
          file.puts Chef::JSONCompat.to_json_pretty(run_data)
        end
      end
      def build_report_dir
        unless File.exists?(config[:path])
          FileUtils.mkdir_p(config[:path])
          File.chmod(00700, config[:path])
        end
      end
    end
  end
end
```

Optional Interfaces

The following interfaces may be used in a handler in the same way as the `report` interface to override the default handler behavior in the chef-client. That said, the following interfaces are not typically used in a handler and, for the most part, are completely unnecessary for a handler to work properly and/or as desired.

```data```

The `data` method is used to return the Hash representation of the `run_status` object. For example:

```
def data
 @run_status.to_hash
end
```

```run_report_safely```

The `run_report_safely` method is used to run the report handler, rescuing and logging errors that may arise as the handler runs and ensuring that all handlers get a chance to run during the chef-client run (even if some handlers fail during that run). In general, this method should never be used as an interface in a custom handler unless this default behavior simply must be overridden.

```
def run_report_safely(run_status)
  run_report_unsafe(run_status)
rescue Exception => e
  Chef::Log.error("Report handler #{self.class.name} raised #{e.inspect}")
  Array(e.backtrace).each { |line| Chef::Log.error(line) }
ensure
  @run_status = nil
end
```

```run_report_unsafe```

The `run_report_unsafe` method is used to run the report handler without any error handling. This method should never be used directly in any handler, except during testing of that handler. For example:

```
def run_report_unsafe(run_status)
 @run_status = run_status
 report
end
```

#### `run_status` Object

The `run_status` object is initialized by the chef-client before the `report` interface is run for any handler. The `run_status` object keeps track of the status of the chef-client run and will contain some (or all) of the following properties:

Property	Description
<code>all_resources</code>	A list of all resources that are included in the <code>resource_collection</code> property for the current chef-client run.
<code>backtrace</code>	A backtrace associated with the uncaught exception data which caused a chef-client run to fail, if present; <code>nil</code> for a successful chef-client run.
<code>elapsed_time</code>	The amount of time between the start ( <code>start_time</code> ) and end ( <code>end_time</code> ) of a chef-client run.
<code>end_time</code>	The time at which a chef-client run ended.
<code>exception</code>	The uncaught exception data which caused a chef-client run to fail; <code>nil</code> for a successful chef-client run.
<code>failed?</code>	Indicates that a chef-client run failed; <code>true</code> when uncaught exceptions were raised during a chef-client run. An exception handler runs when the <code>failed?</code> indicator is <code>true</code> .
<code>node</code>	The node on which the chef-client run occurred.
<code>run_context</code>	An instance of the <code>Chef::RunContext</code> object; used by the chef-client to track the context of the run; provides access to the <code>cookbook_collection</code> , <code>resource_collection</code> , and <code>definitions</code> properties.
<code>start_time</code>	The time at which a chef-client run started.
<code>success?</code>	Indicates that a chef-client run succeeded; <code>true</code> when uncaught exceptions were not raised during a chef-client run. A report handler runs when the <code>success?</code> indicator is <code>true</code> .
<code>updated_resources</code>	A list of resources that were marked as updated as a result of the chef-client run.

#### Note

These properties are not always available. For example, a start handler runs at the beginning of the chef-client run, which means that properties like `end_time` and `elapsed_time` are still unknown and will be unavailable to the `run_status` object.

#### Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Enable the CloudkickHandler handler

The following example shows how to enable the `CloudkickHandler` handler, which adds it to the default handler path and passes the `oauth` key/secret to the handler's initializer:

```
chef_handler "CloudkickHandler" do
 source "#{node['chef_handler']['handler_path']}/cloudkick_handler.rb"
 arguments [node['cloudkick']['oauth_key'], node['cloudkick']['oauth_secret']]
 action :enable
end
```

### Enable handlers during the compile phase

```
chef_handler "Chef::Handler::JsonFile" do
 source "chef/handler/json_file"
 arguments :path => '/var/chef/reports'
 action :nothing
end.run_action(:enable)
```

### Handle only exceptions

```
chef_handler "Chef::Handler::JsonFile" do
 source "chef/handler/json_file"
 arguments :path => '/var/chef/reports'
 supports :exception => true
 action :enable
end
```

### Cookbook Versions (a custom handler)

Community member [juliandunn](#) created a custom report handler that logs all of the cookbooks and cookbook versions that were used during the chef-client run, and then reports after the run is complete. This handler requires the `chef_handler` resource (which is available from the `chef_handler` cookbook).

cookbook\_versions.rb:

The following custom handler defines how cookbooks and cookbook versions that are used during the chef-client run will be compiled into a report using the `Chef::Log` class in the chef-client:

```
require 'chef/log'

module Opscode
 class CookbookVersionsHandler < Chef::Handler
 def report
 cookbooks = run_context.cookbook_collection
 Chef::Log.info("Cookbooks and versions run: #{cookbooks.keys.map {|x| cookbooks[x].name.to_s + " "}.join}")
 end
 end
end
```

default.rb:

The following recipe is added to the run-list for every node on which a list of cookbooks and versions will be generated as report output after every chef-client run.

```
include_recipe "chef_handler"

cookbook_file "#{node['chef_handler']['handler_path']}/cookbook_versions.rb" do
 source "cookbook_versions.rb"
 owner "root"
 group "root"
 mode 00755
 action :create
end

chef_handler "Opscode::CookbookVersionsHandler" do
 source "#{node['chef_handler']['handler_path']}/cookbook_versions.rb"
 supports :report => true
 action :enable
end
```

This recipe will generate report output similar to the following:

```
[2013-11-26T03:11:06+00:00] INFO: Chef Run complete in 0.300029878 seconds
[2013-11-26T03:11:06+00:00] INFO: Running report handlers
[2013-11-26T03:11:06+00:00] INFO: Cookbooks and versions run: ["chef_handler 1.1.4", "cookbook_versions_h
[2013-11-26T03:11:06+00:00] INFO: Report handlers complete
```

### JsonFile Handler

The `json_file` handler is available from the `chef_handler` cookbook and can be used with exceptions and reports. It serializes run status data to a JSON file. This handler needs to be enabled by adding the following lines of Ruby code to either `client.rb` or `solo.rb`:

```
require 'chef/handler/json_file'
report_handlers << Chef::Handler::JsonFile.new(:path => "/var/chef/reports")
exception_handlers << Chef::Handler::JsonFile.new(:path => "/var/chef/reports")
```

and then is added to a recipe:

```
chef_handler "Chef::Handler::JsonFile" do
 source "chef/handler/json_file"
 arguments :path => '/var/chef/reports'
 action :enable
end
```

After it has run, the run status data can be loaded and inspected via Interactive Ruby (IRb):

```
irb(main):001:0> require 'rubygems' => true
irb(main):002:0> require 'json' => true
irb(main):003:0> require 'chef' => true
irb(main):004:0> r = JSON.parse(IO.read("/var/chef/reports/chef-run-report-20110322060731.json")) => ...
irb(main):005:0> r.keys => ["end_time", "node", "updated_resources", "exception", "all_resources", "succe
irb(main):006:0> r['elapsed_time'] => 0.00246
```

Register the JsonFile handler

```
chef_handler "Chef::Handler::JsonFile" do
 source "chef/handler/json_file"
 arguments :path => '/var/chef/reports'
 action :enable
end
```

ErrorReport Handler

The `error_report` handler is built into the chef-client and can be used for both exceptions and reports. It serializes error report data to a JSON file. This handler needs to be enabled by adding the following lines of Ruby code to either the `client.rb` file or the `solo.rb` file, depending on how the chef-client is being run:

```
require 'chef/handler/error_report'
report_handlers << Chef::Handler::ErrorReport.new(:path => "/var/chef/reports")
exception_handlers << Chef::Handler::ErrorReport.new(:path => "/var/chef/reports")
```

and then is added to a recipe:

```
chef_handler "Chef::Handler::ErrorReport" do
 source "chef/handler/error_report"
 arguments :path => '/var/chef/reports'
 action :enable
end
```

cookbook\_file

The `cookbook_file` resource is used to transfer files from a sub-directory of the `files/` directory in a cookbook to a specified path that is located on the host running the chef-client or chef-solo. The file in a cookbook is selected according to file specificity, which allows different source files to be used based on the hostname, host platform (operating system, distro, or as appropriate), or platform version. Files that are located under `COOKBOOK_NAME/files/default` can be used on any platform.

Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a file.
<code>:create_if_missing</code>	Use to create a file only if the file does not exist. (When the file exists, nothing happens.)
<code>:delete</code>	Use to delete a file.
<code>:touch</code>	Use to touch a file. This updates the access (atime) and file modification (mtime) times for a file. (This action may be used with this resource, but is typically only used with the <code>file</code> resource.)

Attributes

This resource has the following attributes:

Attribute	Description
<code>atomic_update</code>	Indicates whether atomic file updates are used on a per-resource basis. Set to <code>true</code> for atomic file updates. Set to <code>false</code> for non-atomic file updates. (This setting overrides <code>file_atomic_update</code> , which is a global setting found in the <code>client.rb</code> file.) Default value: <code>true</code> .
<code>backup</code>	The number of backups to be kept. Set to <code>false</code> to prevent backups from being kept. Default value: 5.
<code>cookbook</code>	The cookbook in which a file is located (if it is not located in the current cookbook). The default value is the current cookbook.
<code>force_unlink</code>	Use to specify how the chef-client handles certain situations when the target file turns out not to be a file. For example, when a target file is actually a symlink. Set to <code>true</code> to have the chef-client delete the non-file target and replace it with the specified file. Set to <code>false</code> for the chef-client to

raise an error. Default value: `false`.

<code>group</code>	A string or ID that identifies the group owner by group name, including fully qualified group names such as <code>domain\group</code> or <code>group@domain</code> . If this value is not specified, existing groups will remain unchanged and new group assignments will use the default POSIX group (if available).
<code>inherits</code>	Microsoft Windows only. Indicates that a file inherits rights from its parent. Default value: <code>true</code> .
<code>manage_symlink_source</code>	Indicates that the chef-client will detect and manage the source file for a symlink. Possible values: <code>nil</code> , <code>true</code> , or <code>false</code> . When this value is set to <code>nil</code> , the chef-client will manage a symlink's source file and emit a warning. When this value is set to <code>true</code> , the chef-client will manage a symlink's source file and not emit a warning. Default value: <code>nil</code> . The default value will be changed to <code>false</code> in a future version.
<code>mode</code>	<p>The octal mode for a file. If <code>mode</code> is not specified and if the file already exists, the existing mode on the file is used. If <code>mode</code> is not specified, the file does not exist, and the <code>:create</code> action is specified, the chef-client will assume a mask value of <code>0777</code> and then apply the umask for the system on which the file will be created to the mask value. For example, if the umask on a system is <code>022</code>, the chef-client would use the default value of <code>0755</code>.</p> <p>The behavior is different depending on the platform.</p> <p>UNIX- and Linux-based systems: The octal mode that is passed to <code>chmod</code>. If the value is specified as a quoted string, it will work exactly as if the <code>chmod</code> command was passed. If the value is specified as an integer, prepend a zero (0) to the value to ensure it is interpreted as an octal number. For example, to assign read, write, and execute rights for all users, use <code>0777</code> or <code>'777'</code>; for the same rights, plus the sticky bit, use <code>01777</code> or <code>'1777'</code>.</p> <p>Microsoft Windows: The octal mode that is translated into rights for Microsoft Windows security. Values up to <code>0777</code> are allowed (no sticky bits) and mean the same in Microsoft Windows as they do in UNIX, where 4 equals <code>GENERIC_READ</code>, 2 equals <code>GENERIC_WRITE</code>, and 1 equals <code>GENERIC_EXECUTE</code>. This attribute cannot be used to set <code>:full_control</code>. This attribute has no effect if not specified, but when this attribute and <code>rights</code> are both specified, the effects will be cumulative.</p>
<code>owner</code>	A string or ID that identifies the group owner by user name, including fully qualified user names such as <code>domain\user</code> or <code>user@domain</code> . If this value is not specified, existing owners will remain unchanged and new owner assignments will use the current user (when necessary).
<code>path</code>	<p>The path to the location in which a file will be created.</p> <p>Microsoft Windows: A path that begins with a forward slash (/) will point to the root of the current working directory of the chef-client process. This path can vary from system to system. Therefore, using a path that begins with a forward slash (/) is not recommended.</p>
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>rights</code>	Microsoft Windows only. The permissions for users and groups in a Microsoft Windows environment. For example: <code>rights &lt;permissions&gt;, &lt;principal&gt;, &lt;options&gt;</code> where <code>&lt;permissions&gt;</code> specifies the rights granted to the principal, <code>&lt;principal&gt;</code> is the group or user name, and <code>&lt;options&gt;</code> is a Hash with one (or more) advanced rights options.
<code>source</code>	The location of a file in the <code>/files</code> directory in a cookbook located in the chef-repo. Can be used to distribute specific files to specific platforms (see the section "File Specificity", below). Default value: the <code>name</code> of the resource block (see Syntax section above).

#### Note

Use the `owner` and `right` attributes and avoid the `group` and `mode` attributes whenever possible. The `group` and `mode` attributes are not true Microsoft Windows concepts and are provided more for backward compatibility than for best practice.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::CookbookFile</code>	<code>cookbook_file</code>	The default provider for all platforms.

## Syntax

The syntax for using the `cookbook_file` resource in a recipe is as follows:

```
cookbook_file "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `cookbook_file` tells the chef-client to use the `Chef::Provider::CookbookFile` provider during the chef-client run
- `name` is the name of the resource block; when the `path` attribute is not specified as part of a recipe, `name` is also the path to the file
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the `cookbook_file` resource can work when used in a recipe. In this example, because the `source` attribute is unspecified, the name of the resource ("`cookbook_test_file`") defines the name the source file. The chef-client will look for this source file in the `/cookbook_name/files/default/` directory. The `path` attribute defines the location in which the file will be created. The `:create_if_missing` action ensures that nothing happens if the file already exists.

```
cookbook_file "cookbook_test_file" do
 path "/tmp/test_file"
 action :create_if_missing
end
```

## File Specificity

A cookbook will frequently be designed to work across many platforms and will often be required to distribute a specific file to a specific platform. A cookbook can be designed to support distributing files across platforms, but ensuring that the right file ends up on each system.

The pattern for file specificity is as follows:

1. `host-node[:fqdn]`
2. `node[:platform]-node[:platform_version]`
3. `node[:platform]-version_components`: The version string is split on decimals and searched from greatest specificity to least; for example, if the location from the last rule was `centos-5.7.1`, then `centos-5.7` and `centos-5` would also be searched.
4. `node[:platform]`
5. `default`

A cookbook may have a `/files` directory structure like this:

```
files/
 host-foo.example.com
 ubuntu-10.04
 ubuntu-10
 ubuntu
 redhat-5.8
 redhat-6.4
 ...
 default
```

and a resource that looks something like the following:

```
resource_type "/usr/local/bin/apache2_module_conf_generate.pl" do
 source "apache2_module_conf_generate.pl"
 mode 0755
 owner "root"
 group "root"
end
```

where `resource_type` is the `cookbook_file` or `remote_file` resource. This resource would be matched in the same order as the `/files` directory structure. For a node that is running Ubuntu 10.04, the second item would be the matching item and the location to which the file identified in the `cookbook_file` resource would be distributed:

```
host-foo.example.com/apache2_module_conf_generate.pl
ubuntu-10.04/apache2_module_conf_generate.pl
ubuntu-10/apache2_module_conf_generate.pl
ubuntu/apache2_module_conf_generate.pl
default/apache2_module_conf_generate.pl
```

If the `apache2_module_conf_generate.pl` file was located in the cookbook directory under `files/host-foo.example.com/`, the specified file(s) would only be copied to the machine with the domain name `foo.example.com`.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Transfer a file

```
cookbook_file "/tmp/testfile" do
 source "testfile"
 mode 00644
end
```

### Handle `cookbook_file` and `yum_package` resources in the same recipe

When a `cookbook_file` resource and a `yum_package` resource are both called from within the same recipe, dump the cache and use the new repository immediately to ensure that the correct package is installed:

```
cookbook_file "/etc/yum.repos.d/custom.repo" do
 source "custom"
 mode 00644
```

```
end
```

```
yum_package "only-in-custom-repo" do
 action :install
 flush_cache [:before]
end
```

#### Install repositories from a file, trigger a command, and force the internal cache to reload

The following example shows how to install new yum repositories from a file, where the installation of the repository triggers a creation of the yum cache that forces the internal cache for the chef-client to reload:

```
execute "create-yum-cache" do
 command "yum -q makecache"
 action :nothing
end

ruby_block "reload-internal-yum-cache" do
 block do
 Chef::Provider::Package::Yum::YumCache.instance.reload
 end
 action :nothing
end

cookbook_file "/etc/yum.repos.d/custom.repo" do
 source "custom"
 mode 00644
 notifies :run, "execute[create-yum-cache]", :immediately
 notifies :create, "ruby_block[reload-internal-yum-cache]", :immediately
end
```

#### Use a case statement

The following example shows how a case statement can be used to handle a situation where an application needs to be installed on multiple platforms, but the where the install directories are different, depending on the platform:

```
cookbook_file "application.pm" do
 case node[:platform]
 when "centos", "redhat"
 path "/usr/lib/version/1.2.3/dir/application.pm"
 when "arch"
 path "/usr/share/version/core_version/dir/application.pm"
 else
 path "/etc/version/dir/application.pm"
 end
 source "application-#{node[:languages][:perl][:version]}.pm"
 owner "root"
 group "root"
 mode 0644
end
```

## cron

The **cron** resource is used to manage cron entries for time-based job scheduling. Attributes for a schedule will default to `*` if not provided. The **cron** resource requires access to a crontab program, typically `cron`.

### Warning

The **cron** resource should only be used to modify an entry in a crontab file. Use the **cookbook\_file** or **template** resources to add a crontab file to the `cron.d` directory. The **cron\_d** lightweight resource (found in the **cron** cookbook) is another option for managing crontab files.

## Syntax

The syntax for using the **cron** resource in a recipe is as follows:

```
cron "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `cron` tells the chef-client to use the `Chef::Provider::Cron` provider during the chef-client run
- `"name"` is the name of the cron entry
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

For example, the following example runs weekly cookbook reports:

```
cron "cookbooks_report" do
 action node.tags.include?('cookbooks-report') ? :create : :delete
 minute "0"
 hour "0"
 weekday "1"
 user "opscode"
 mailto "nharvey@opscode.com"
 home "/srv/opscode-community-site/shared/system"
 command %Q{
 cd /srv/opscode-community-site/current &&
 env RUBYLIB="/srv/opscode-community-site/current/lib"
 RAILS_ASSET_ID=`git rev-parse HEAD` RAILS_ENV="#{rails_env}"
```



```
 }
 bundle exec rake cookbooks_report
 end
```

Actions

This resource has the following actions:

Action	Description
<a href="#">:create</a>	Default. Use to create an entry in a cron table file ("crontab"). If an entry already exists with the same name, use to update that entry.
<a href="#">:delete</a>	Use to delete an entry from a cron table file ("crontab").

Attributes

This resource has the following attributes:

Attribute	Description
<a href="#">command</a>	<p>The command to be run or the path to a file that contains the command to be run.</p> <p>Some examples:</p> <pre>command if [ -x /usr/share/mdadm/checkarray ] &amp;&amp; [ \$(date +%d) -le 7 ]; then /usr/share/mdadm/checkarray --cron --all --idle --quiet; fi</pre> <p>and:</p> <pre>command %Q{   cd /srv/opscode-community-site/current &amp;&amp;   env RUBYLIB="/srv/opscode-community-site/current/lib"   RAILS_ASSET_ID=`git rev-parse HEAD` RAILS_ENV="#{rails_env}"   bundle exec rake cookbooks_report }</pre> <p>and:</p> <pre>command "/srv/app/scripts/daily_report"</pre>
<a href="#">day</a>	The day of month at which the cron entry should run (1 - 31). Default value: <a href="#">*</a> .
<a href="#">home</a>	Use to set the HOME environment variable.
<a href="#">hour</a>	The hour at which the cron entry should run (0 - 23). Default value: <a href="#">*</a> .
<a href="#">mailto</a>	Use to set the MAILTO environment variable.
<a href="#">minute</a>	The minute at which the cron entry should run (0 - 59). Default value: <a href="#">*</a> .
<a href="#">month</a>	The month in the year on which a cron entry should run (1 - 12). Default value: <a href="#">*</a> .
<a href="#">path</a>	Use to set the PATH environment variable.
<a href="#">provider</a>	Optional. Use to specify a provider by using its long name. For example: <a href="#">provider Chef::Provider::Long::Name</a> . See the Providers section below for the list of providers available to this resource.
<a href="#">shell</a>	Use to set the SHELL environment variable.
<a href="#">user</a>	The name of the user that runs the command. If the <a href="#">user</a> attribute is changed, the original <a href="#">user</a> for the crontab program will continue to run until that crontab program is deleted. Default value: <a href="#">root</a> .
<a href="#">weekday</a>	The day of the week on this entry should run (0 - 6), where Sunday = 0. Default value: <a href="#">*</a> . May be entered as a symbol, e.g. <a href="#">:monday</a> or <a href="#">:friday</a> .

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<a href="#">Chef::Provider::Cron</a>	<a href="#">cron</a>	The default provider for all platforms.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Run a program at a specified interval

```
cron "noop" do
 hour "5"
```

```

minute "0"
command "/bin/true"
end

```

#### Run an entry if a folder exists

```

cron "ganglia_tomcat_thread_max" do
 command "/usr/bin/gmetric -n 'tomcat threads max' -t uint32 -v `/usr/local/bin/tomcat-stat --thread-max`"
 only_if do File.exist?("/home/jboss") end
end

```

#### Run every Saturday, 8:00 AM

The following example shows a schedule that will run every hour at 8:00 each Saturday morning, and will then send an email to "admin@opscode.com" after each run.

```

cron "name_of_cron_entry" do
 minute "0"
 hour "8"
 weekday "6"
 mailto "admin@opscode.com"
 action :create
end

```

#### Run only in November

The following example shows a schedule that will run at 8:00 PM, every weekday (Monday through Friday), but only in November:

```

cron "name_of_cron_entry" do
 minute "0"
 hour "20"
 day "*"
 month "11"
 weekday "1-5"
 action :create
end

```

## csh

The **csh** resource is used to execute scripts using the csh interpreter and includes all of the actions and attributes that are available to the **execute** resource.

Note
The <b>csh</b> script resource (which is based on the <b>script</b> resource) is different from the <b>ruby_block</b> resource because Ruby code that is run with this resource is created as a temporary file and executed like other script resources, rather than run inline. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the <code>not_if</code> and <code>only_if</code> meta parameters to guard the use of this resource for idempotence.

## Syntax

The syntax for using the **csh** resource in a recipe is as follows:

```

csh "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end

```

where

- `csh` tells the chef-client to use the `Chef::Resource::Script::Csh` provider during the chef-client run
- `name` is the name of the resource block; when the `command` attribute is not specified as part of a recipe, `name` is also the name of the command to be executed
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:

Action	Description
<code>:run</code>	Default. Use to run a script.
<code>:nothing</code>	Indicates that the command should not be run. This action is used to specify that a command is run only when another resource notifies it.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>code</code>	A quoted ( " ") string of code to be executed.

<u>command</u>	The name of the command to be executed. Default value: the <u>name</u> of the resource block (see Syntax section above).
<u>creates</u>	Indicates that a command to create a file will not be run when that file already exists.
<u>cwd</u>	The current working directory.
<u>environment</u>	A Hash of environment variables in the form of {"ENV_VARIABLE" => "VALUE"}. (These variables must exist for a command to be run successfully.)
<u>flags</u>	One (or more) command line flags that are passed to the interpreter when a command is invoked.
<u>group</u>	The group name or group ID that must be changed before running a command.
<u>path</u>	An array of paths to use when searching for a command. These paths are not added to the command's environment \$PATH. The default value uses the system path.
<u>provider</u>	Optional. Use to specify a provider by using its long name. For example: <u>provider</u> Chef::Provider::Long::Name. See the Providers section below for the list of providers available to this resource.
<u>returns</u>	The return value for a command. This may be an array of accepted values. An exception is raised when the return value(s) do not match. Default value: 0.
<u>timeout</u>	The amount of time (in seconds) a command will wait before timing out. Default value: 3600.
<u>user</u>	The user name or user ID that should be changed before running a command.
<u>umask</u>	The file mode creation mask, or umask.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<u>Chef::Provider::Script</u>	<u>script</u>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<u>Chef::Provider::Script::Csh</u>	<u>csh</u>	The provider that is used with the csh command interpreter.

Examples

None.

deploy

The **deploy** resource is used to manage and control deployments. This is a popular resource, but is also complex, having the most attributes, multiple providers, the added complexity of callbacks, plus four attributes that support layout modifications from within a recipe.

The **deploy** resource is modeled after Capistrano, a utility and framework for executing commands in parallel on multiple remote machines via SSH. The **deploy** resource is designed to behave in a way that is similar to the deploy and deploy:migration tasks in Capistrano.

Syntax

The syntax for using the **deploy** resource in a recipe is as follows:

```
deploy "name" do
 attribute "value" # see attributes section below
 ...
 callback do
 # callback, including release_path or new_resource
 end
 ...
 purge_before_symlink
 create_dirs_before_symlink
 symlink
 action :action # see actions section below
end
```

where

- deploy tells the chef-client to use either the Chef::Provider::Deploy::Revision or Chef::Provider::Deploy::Timestamped provider during the chef-client run. More specific short names —timestamped\_deploy, deploy\_revision, or deploy\_branch—can be used instead of the deploy short name.
- name is the name of the resource block; when the deploy\_to attribute is not specified as part of a recipe, name is also the location in which the deployment steps will occur
- attribute is zero (or more) of the attributes that are available for this resource
- callback represents additional Ruby code that is used to pass a block or to specify a file, and then provide additional information to the chef-client at specific times during the deployment process
- purge\_before\_symlink, create\_dirs\_before\_symlink, and symlink are attributes that are used to link configuration files,

remove directories, create directories, or map files and directories during the deployment process

- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the `deploy_revision` resource can work when used in a recipe. In this example, an application will be deployed to a folder named `/path/to/application`:

```
deploy_revision "/path/to/application" do
 repo 'ssh://name-of-git-repo/repos/repo.git'
 migrate false
 purge_before_symlink %w{one two folder/three}
 create_dirs_before_symlink []
 symlinks(
 "one" => "one",
 "two" => "two",
 "three" => "folder/three"
)
 before_restart do
 # some Ruby code
 end
 notifies :restart, "service[foo]"
 notifies :restart, "service[bar]"
end
```

For the example shown above:

- Because an action is not explicitly specified, the chef-client will use the default action: `:deploy`
- The `purge_before_symlink` application layout is an array of paths that will be cleared before the `symlinks` attribute is run
- The `create_dirs_before_symlink` attribute is empty, which is different from the default
- The `symlinks` attribute is creating three symbolic links
- The `before_restart` callback is being used to add custom actions that will occur at the end of the deployment process, but before any services have been notified
- At the end, the recipe is using the `notifies` attribute—a common attribute available to all resources—to alert two services (named “foo” and “bar”) that they should restart.

## Deploy Strategies

In the `deploy` directory, a sub-directory named `shared` must be created. This sub-directory is where configuration and temporary files will be kept. A typical Ruby on Rails application will have `config`, `log`, `pids`, and `system` directories within the `shared` directory to keep the files stored there independent of the code in the source repository.

In addition to the `shared` sub-directory, the deploy process will create sub-directories named `releases` and `current` (also in the `deploy` directory). The `release` directory holds (up to) five most recently deployed versions of an application. The `current` directory holds the currently-released version.

For example:

```
deploy_directory/
 current/
 releases/
 shared/
 config/
 log/
 pids/
 system/
```

## Deploy Cache File

The chef-client uses a cache file to keep track of the order in which each revision of an application is deployed. By default, the cache file is located at `/var/chef/cache/revision-deploys/APPNAME/`. To force a re-deploy, delete the deployment directory or delete the cache file.

## Deploy Phases

A deployment happens in four phases:

1. **Checkout**—the chef-client uses the `scm` resource to get the specified application revision, placing a clone or checkout in the sub-directory of the `deploy` directory named `cached-copy`. A copy of the application is then placed in a sub-directory under `releases`.
2. **Migrate**—If a migration is to be run, the chef-client symlinks the database configuration file into the checkout (`config/database.yml` by default) and runs the migration command. For a Ruby on Rails application, the `migration_command` is usually set to `rake db:migrate`.
3. **Symlink**—Directories for shared and temporary files are removed from the checkout (`log`, `tmp/pids`, and `public/system` by default). After this step, any needed directories (`tmp`, `public`, and `config` by default) are created if they don't already exist. This step is completed by symlinking shared directories into the current `release`, `public/system`, `tmp/pids`, and `log` directories, and then symlinking the `release` directory to `current`.
4. **Restart**—The application is restarted according to the restart command set in the recipe.

## Callbacks

In-between each step in a deployment process, callbacks can be run using arbitrary Ruby code, including recipes. All callbacks support embedded recipes given in a block, but each callback assumes a shell command (instead of a deploy hook filename) when given a string.

The following callback types are available:

Callback	Description
<code>after_restart</code>	A block of code or a path to a file that contains code that is run after restarting. Default value: <code>deploy/after_restart.rb</code> .
<code>before_migrate</code>	A block of code (or a path to a file that contains code) that is run before a migration. Default value: <code>deploy/before_migrate.rb</code> .
<code>before_restart</code>	A block of code (or a path to a file that contains code) that is run before restarting. Default value: <code>deploy/before_restart.rb</code> .
<code>before_symlink</code>	A block of code (or a path to a file that contains code) that is run before symbolic linking. Default value: <code>deploy/before_symlink.rb</code> .

Each of these callback types can be used in one of three ways:

- To pass a block of code, such as Ruby or Python
- To specify a file
- To do neither; the chef-client will look for a callback file named after one of the callback types (`before_migrate.rb`, for example) and if the file exists, to evaluate it as if it were a specified file

Within a callback, there are two ways to get access to information about the deployment:

- `release_path` can be used to get the path to the current release
- `new_resource` can be used to access the deploy resource, including environment variables that have been set there (using `new_resource` is a preferred approach over using the `@configuration` variable)

Both of these options must be available at the top-level within the callback, along with any assigned values that will be used later in the callback.

Callbacks and Capistrano

If you are familiar with Capistrano, the following examples should help you know when to use the various callbacks that are available. If you are not familiar with Capistrano, then follow the semantic names of these callbacks to help you determine when to use each of the callbacks within a recipe that is built with the **deploy** resource.

The following example shows where callbacks fit in relation to the steps taken by the `deploy` process in Capistrano:



and the following example shows the same comparison, but with the `deploy:migrations` process:



Layout Modifiers

The **deploy** resource expects an application to be structured like a Ruby on Rails application, but the layout can be modified to meet custom requirements as needed. Use the following attributes within a recipe to modify the layout of a recipe that is using the **deploy** resource:

Layout Modifiers	Description
<code>create_dirs_before_symlink</code>	Use this attribute to create directories before symbolic links are created. This attribute runs after <code>purge_before_symlink</code> and before <code>symlink</code> .
<code>purge_before_symlink</code>	Use this attribute to specify an array of directories (relative to the application root) that should be removed from a checkout before symbolic links are created. This attribute runs before <code>create_dirs_before_symlink</code> and before <code>symlink</code> .
<code>symlink_before_migrate</code>	Use this attribute to map files in a shared directory to the current release directory. The symbolic links for these files will be created before any migration

is run. Use `symlink_before_migrate({})` or `symlink_before_migrate nil` instead of `symlink_before_migrate {}` because `{}` will be interpreted as a block rather than an empty Hash. Set to `nil` to prevent the creation of default symbolic links.

## symlinks

Use this attribute to map files in a shared directory to their paths in the current release directory. This attribute runs after `create_dirs_before_symlink` and `purge_before_symlink`.

## Actions

This resource has the following actions:

Action	Description
<code>:deploy</code>	Default. Use to deploy an application.
<code>:force_deploy</code>	Use to remove any existing release of the same code version and re-deploy a new one in its place.
<code>:rollback</code>	Use to roll an application back to the previous release.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>after_restart</code>	A block of code or a path to a file that contains code that is run after restarting. Default value: <code>deploy/after_restart.rb</code> .
<code>before_migrate</code>	A block of code (or a path to a file that contains code) that is run before a migration. Default value: <code>deploy/before_migrate.rb</code> .
<code>before_restart</code>	A block of code (or a path to a file that contains code) that is run before restarting. Default value: <code>deploy/before_restart.rb</code> .
<code>before_symlink</code>	A block of code (or a path to a file that contains code) that is run before symbolic linking. Default value: <code>deploy/before_symlink.rb</code> .
<code>branch</code>	The alias for the revision.
<code>create_dirs_before_symlink</code>	Use this attribute to create directories before symbolic links are created. This attribute runs after <code>purge_before_symlink</code> and before <code>symlink</code> . Default value: <code>%w{tmp public config}</code> (or the same as <code>["tmp", "public", "config"]</code> ).
<code>deploy_to</code>	The "meta root" for the application, if different from the path that is used to specify the name of a resource. Default value: the name of the resource block (see Syntax section above).
<code>environment</code>	A Hash of environment variables in the form of <code>{"ENV_VARIABLE" =&gt; "VALUE"}</code> . (These variables must exist for a command to be run successfully.)
<code>group</code>	The system group that is responsible for the checked-out code.
<code>keep_releases</code>	The number of releases for which a backup is kept. Default value: 5.
<code>migrate</code>	Indicates that the migration command will be run. Default value: <code>false</code> .
<code>migration_command</code>	A string that contains a shell command that can be executed to run a migration operation.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>purge_before_symlink</code>	Use this attribute to specify an array of directories (relative to the application root) that should be removed from a checkout before symbolic links are created. This attribute runs before <code>create_dirs_before_symlink</code> and before <code>symlink</code> . Default value: <code>%w{log tmp/pids public/system}</code> (or the same as <code>["log", "tmp/pids", "public/system"]</code> ).
<code>repo</code>	The alias for the repository.
<code>repository</code>	The URI for the repository.
<code>repository_cache</code>	The name of the sub-directory in which the pristine copy of an application's source is kept. Default value: <code>cached-copy</code> .
<code>restart_command</code>	A string that contains a shell command that can be executed to run a restart operation.
<code>revision</code>	The revision to be checked out. This can be symbolic, like <code>HEAD</code> or it can be a source control management-specific revision identifier. Default value: <code>HEAD</code> .

<code>rollback_on_error</code>	Indicates whether a resource will roll back to a previously-deployed release if an error occurs when deploying a new release. Default value: <code>false</code> .
<code>scm_provider</code>	The name of the source control management provider. Default value: <code>Chef::Provider::Git</code> . Optional values: <code>Chef::Provider::Subversion</code> .
<code>symlinks</code>	Use this attribute to map files in a shared directory to their paths in the current release directory. This attribute runs after <code>create_dirs_before_symlink</code> and <code>purge_before_symlink</code> . Default value: <code>{"system" =&gt; "public/system", "pids" =&gt; "tmp/pids", "log" =&gt; "log"}</code> .
<code>symlink_before_migrate</code>	Use this attribute to map files in a shared directory to the current release directory. The symbolic links for these files will be created before any migration is run. Use <code>symlink_before_migrate({})</code> or <code>symlink_before_migrate nil</code> instead of <code>symlink_before_migrate {}</code> because <code>{}</code> will be interpreted as a block rather than an empty Hash. Set to <code>nil</code> to prevent the creation of default symbolic links. Default value: <code>{"config/database.yml" =&gt; "config/database.yml"}</code> .
<code>user</code>	The system user that is responsible for the checked-out code.

The following attributes are for use with git only:

Attribute	Description
<code>enable_submodules</code>	Use to perform a sub-module initialization and update. Default value: <code>false</code> .
<code>git_ssh_wrapper</code>	The alias for the <code>ssh_wrapper</code> .
<code>remote</code>	The remote repository to be used when synchronizing an existing clone. Default value: <code>origin</code> .
<code>shallow_clone</code>	Indicates that the clone depth is set to 5. Default value: <code>false</code> .
<code>ssh_wrapper</code>	The path to the wrapper script used when running SSH with git. The <code>GIT_SSH</code> environment variable is set to this.

The following attributes are for use with Subversion only:

Attribute	Description
<code>svn_arguments</code>	The extra arguments that are passed to the Subversion command.
<code>svn_password</code>	The password for the user that has access to the Subversion repository.
<code>svn_username</code>	The user name for a user that has access to the Subversion repository.

For example:

```
deploy "/my/deploy/dir" do
 repo "git@github.com:whoami/project"
 revision "abc123" # or "HEAD" or "TAG_for_1.0" or (subversion) "1234"
 user "deploy_ninja"
 enable_submodules true
 migrate true
 migration_command "rake db:migrate"
 environment "RAILS_ENV" => "production", "OTHER_ENV" => "foo"
 shallow_clone true
 keep_releases 10
 action :deploy # or :rollback
 restart_command "touch tmp/restart.txt"
 git_ssh_wrapper "wrap-ssh4git.sh"
 scm_provider Chef::Provider::Git # is the default, for svn: Chef::Provider::Subversion
end
```

Providers

The `deploy` resource providers are used to determine whether to deploy based on whether the release directory in which the deployment is to be made actually exists. The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Deploy</code>	<code>deploy</code>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<code>Chef::Provider::Deploy::Branch</code>	<code>deploy_branch</code>	See below for more information.
<code>Chef::Provider::Deploy::Revision</code>	<code>deploy_revision</code>	See below for more information.
<code>Chef::Provider::Deploy::TimestampedDeploy</code>	<code>timestamped_deploy</code>	The default provider for all platforms. See below for more information.

`deploy_branch`

The `deploy_branch` resource is used in the same way as the `deploy_resource` resource. It uses the `Deploy::Revision` provider

and has uses the same set of actions and attributes.

#### deploy\_revision

The `deploy_revision` provider is the recommended provider, even if it is not listed as the default. The `deploy_revision` provider is used to ensure that the name of a release sub-directory is based on a revision identifier. For users of git, this will be the familiar SHA checksum. For users of Subversion, it will be the integer revision number. If a name other than a revision identifier is provided—branch names, tags, and so on—the chef-client will ignore the alternate names and will look up the revision identifier and use it to name the release sub-directory. When the `deploy_revision` provider is given an exact revision to deploy, it will behave in an idempotent manner.

The `deploy_revision` provider results in deployed components under the destination location that is owned by the user who runs the application. This is sometimes an issue for certain workflows. If issues arise, consider the following:

- Incorporate changing permissions to the desired end state from within a recipe
- Add a `before_restart` block to fix up the permissions
- Have an unprivileged user (for example: `opscod`) be the owner of the `deploy` directory and another unprivileged user (for example: `opscodapp`) run the application. Most often, this is the solution that works best

When using the `deploy_revision` provider, and when the deploy fails for any reason, and when the same code is used to re-deploy, the action should be set manually to `:force_deploy`. Forcing the re-deploy will remove the old release directory, after which the deploy can proceed as usual. (Forcing a re-deploy over the current release can cause some downtime.) Deployed revisions are stored in `(file_cache_path)/revision-deploys/(deploy_path)`.

#### timestamped\_deploy

The `timestamped_deploy` provider is the default **deploy** provider. It is used to name release directories with a timestamp in the form of `YYYYMMDDHHMMSS`. For example: `/my/deploy/dir/releases/20121120162342`. The **deploy** resource will determine whether or not to deploy code based on the existence of the release directory in which it is attempting to deploy. Because the timestamp is different for every chef-client run, the `timestamped_deploy` provider is not idempotent. When the `timestamped_deploy` provider is used, it requires that the action setting on a resource be managed manually in order to prevent unintended continuous deployment.

### Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

#### Modify the layout of a Ruby on Rails application

The layout of the **deploy** resource matches a Ruby on Rails app by default, but this can be customized. To customize the layout, do something like the following:

```
deploy "/my/apps/dir/deploy" do
 # Use a local repo if you prefer
 repo "/path/to/gitrepo/typo/"
 environment "RAILS_ENV" => "production"
 revision "HEAD"
 action :deploy
 migration_command "rake db:migrate --trace"
 migrate true
 restart_command "touch tmp/restart.txt"
 create_dirs_before_symlink %w{tmp public config deploy}

 # You can use this to customize if your app has extra configuration files
 # such as amqp.yml or app_config.yml
 symlink_before_migrate "config/database.yml" => "config/database.yml"

 # If your app has extra files in the shared folder, specify them here
 symlinks {
 "system" => "public/system",
 "pids" => "tmp/pids",
 "log" => "log",
 "deploy/before_migrate.rb" => "deploy/before_migrate.rb",
 "deploy/before_symlink.rb" => "deploy/before_symlink.rb",
 "deploy/before_restart.rb" => "deploy/before_restart.rb",
 "deploy/after_restart.rb" => "deploy/after_restart.rb"
 }
end
```

#### Use resources within callbacks

Using resources from within your callbacks as blocks or within callback files distributed with your application's source code. To use embedded recipes for callbacks:

```
deploy "#{node['tmpdir']}/deploy" do
 repo "#{node['tmpdir']}/gitrepo/typo/"
 environment "RAILS_ENV" => "production"
 revision "HEAD"
 action :deploy
 migration_command "rake db:migrate --trace"
 migrate true

 # Callback awesomeness:
 before_migrate do
 current_release = release_path

 directory "#{current_release}/deploy" do
 mode 00755
 end

 # creates a callback for before_symlink
 template "#{current_release}/deploy/before_symlink_callback.rb" do
 source "embedded_recipe_before_symlink.rb.erb"
 end
 end
end
```



```

 mode 00644
 end

end

This file can contain Chef recipe code, plain ruby also works
before_symlink "deploy/before_symlink_callback.rb"

restart do
 current_release = release_path
 file "#{release_path}/tmp/restart.txt" do
 mode 00644
 end
end

end
end

```

### Deploy from a private git repository without using the application cookbook

To deploy from a private git repository without using the **application** cookbook, first ensure that:

- the private key does not have a passphrase, as this will pause a chef-client run to wait for input
- an SSH wrapper is being used
- a private key has been added to the node

and then use code like the following to remove a passphrase from a private key:

```
ssh-keygen -p -P 'YOURPASSPHRASE' -N '' -f id_deploy
```

### Use an SSH wrapper

To write a recipe that uses an SSH wrapper:

1. Create a file in the `cookbooks/COOKBOOK_NAME/files/default` directory that is named `wrap-ssh4git.sh` and which contains the following:

```
#!/usr/bin/env bash
/usr/bin/env ssh -o "StrictHostKeyChecking=no" -i "/tmp/private_code/.ssh/id_deploy" $1 $2
```

2. Set up the cookbook file.
3. Add a recipe to the cookbook file similar to the following:

```

directory "/tmp/private_code/.ssh" do
 owner "ubuntu"
 recursive true
end

cookbook_file "/tmp/private_code/wrap-ssh4git.sh" do
 source "wrap-ssh4git.sh"
 owner "ubuntu"
 mode 00700
end

deploy "private_repo" do
 repo "git@github.com:acctname/private-repo.git"
 user "ubuntu"
 deploy_to "/tmp/private_code"
 action :deploy
 ssh_wrapper "/tmp/private_code/wrap-ssh4git.sh"
end

```

This will deploy the git repository at `git@github.com:acctname/private-repo.git` in the `/tmp/private_code` directory.

### Use a callback to include a file that will be passed as a code block

The code in a file that is included in a recipe using a callback is evaluated exactly as if the code had been put in the recipe as a block. Files are searched relative to the current release.

To specify a file that contains code to be used as a block:

```

deploy "/deploy/dir/" do
 # ...

 before_migrate "callbacks/do_this_before_migrate.rb"
end

```

### Use a callback to pass a code block

To pass a block of Python code before a migration is run:

```

deploy_revision "/deploy/dir/" do
 # other attributes
 # ...

 before_migrate do
 # release_path is the path to the timestamp dir
 # for the current release
 current_release = release_path

 # Create a local variable for the node so we'll have access to
 # the attributes
 deploy_node = node

 # A local variable with the deploy resource.

```

```

deploy_resource = new_resource

python do
 cwd current_release
 user "myappuser"
 code<<-PYCODE
 # Woah, callbacks in python!
 # ...
 # current_release, deploy_node, and deploy_resource are all available
 # within the deploy hook now.
PYCODE
end
end
end

```

#### Use the same API for all recipes using the same gem

Any recipes using the `git-deploy` gem can continue using the same API. To include this behavior in a recipe, do something like the following:

```

deploy "/srv/#{appname}" do
 repo "git://github.com/radiant/radiant.git"
 revision "HEAD"
 user "railsdev"
 enable_submodules false
 migrate true
 migration_command "rake db:migrate"
 # Giving a string for environment sets RAILS_ENV, MERB_ENV, RACK_ENV
 environment "production"
 shallow_clone true
 action :deploy
 restart_command "touch tmp/restart.txt"
end

```

#### Deploy without creating symbolic links to a shared folder

To deploy without creating symbolic links to a shared folder:

```

deploy "/my/apps/dir/deploy" do
 symlinks {}
end

```

When deploying code that is not Ruby on Rails and symbolic links to a shared folder are not wanted, use parentheses `()` or `Hash.new` to avoid ambiguity. For example, using parentheses:

```

deploy "/my/apps/dir/deploy" do
 symlinks({})
end

```

or using `Hash.new`:

```

deploy "/my/apps/dir/deploy" do
 symlinks Hash.new
end

```

#### Clear a layout modifier attribute

Using the default attribute values for the various resources is the recommended starting point when working with recipes. Then, depending on what each node requires, these default values can be overridden with `node-`, `role-`, `environment-`, and `cookbook-specific` values. The **deploy** resource has four layout modifiers: `create_dirs_before_symlink`, `purge_before_symlink`, `symlink_before_migrate`, and `symlinks`. Each of these is a Hash that behaves as an attribute of the **deploy** resource. When these layout modifiers are used in a recipe, they appear similar to the following:

```

deploy "name" do
 ...
 symlink_before_migrate {"config/database.yml" => "config/database.yml"}
 create_dirs_before_symlink %w{tmp public config}
 purge_before_symlink %w{log tmp/pids public/system}
 symlinks {
 "system" => "public/system",
 "pids" => "tmp/pids",
 "log" => "log"
 }
 ...
end

```

and then what these layout modifiers look like if they were empty:

```

deploy "name" do
 ...
 symlink_before_migrate nil
 create_dirs_before_symlink []
 purge_before_symlink []
 symlinks nil
 ...
end

```

In most cases, using the empty values for the layout modifiers will prevent the `chef-client` from passing symbolic linking information to a node during the `chef-client` run. However, in some cases, it may be preferable to ensure that one (or more) of these layout modifiers do not pass any symbolic linking information to a node during the `chef-client` run at all. Because each of these layout modifiers are a Hash, the `clear` instance method can be used to clear out these values.

To clear the default values for a layout modifier:

```
deploy "name" do
 ...
 symlink_before_migrate.clear
 create_dirs_before_symlink.clear
 purge_before_symlink.clear
 symlinks.clear
 ...
end
```

In general, use this approach carefully and only after it is determined that nil or empty values won't provide the expected result.

directory

The **directory** resource is used to manage a directory, which is a hierarchy of folders that comprises all of the information stored on a computer. The root directory is the top-level, under which the rest of the directory is organized. The **directory** resource uses the `name` attribute to specify the path to a location in a directory. Typically, permission to access that location in the directory is required.

Syntax

The syntax for using the **directory** resource in a recipe is as follows:

```
directory "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `directory` tells the chef-client to use the `Chef::Provider::Directory` provider during the chef-client run
- `name` is the name of the resource block; when the `path` attribute is not specified as part of a recipe, `name` is also the path to the directory, from the root
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the **directory** resource can work when used in a recipe:

```
directory "/var/lib/foo" do
 owner "root"
 group "root"
 mode 00644
 action :create
end
```

Also, a variable can be used to define the directory, and then that variable can be used within the recipe itself:

```
node.default['apache']['dir'] = '/etc/apache2'

directory node['apache']['dir'] do
 owner 'apache'
 group 'apache'
 action :create
end
```

Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a directory.
<code>:delete</code>	Use to delete a directory.

Attributes

This resource has the following attributes:

Attribute	Description
<code>group</code>	A string or ID that identifies the group owner by group name, including fully qualified group names such as <code>domain\group</code> or <code>group@domain</code> . If this value is not specified, existing groups will remain unchanged and new group assignments will use the default <code>POSIX</code> group (if available).
<code>inherits</code>	Microsoft Windows only. Indicates that a file inherits rights from its parent. Default value: <code>true</code> .
<code>mode</code>	The octal mode for a directory. If <code>mode</code> is not specified and if the directory already exists, the existing mode on the directory is used. If <code>mode</code> is not specified, the directory does not exist, and the <code>:create</code> action is specified, the chef-client will assume a mask value of <code>0777</code> and then apply the umask for the system on which the directory will be created to the <code>mask</code> value. For example, if the umask on a system is <code>022</code> , the chef-client would use the default value of <code>0755</code> .  The behavior is different depending on the platform.  UNIX- and Linux-based systems: The octal mode that is passed to <code>chmod</code> . If the value is specified

as a quoted string, it will work exactly as if the `chmod` command was passed. If the value is specified as an integer, prepend a zero (0) to the value to ensure it is interpreted as an octal number. For example, to assign read, write, and execute rights for all users, use `0777` or `'777'`; for the same rights, plus the sticky bit, use `01777` or `'1777'`.

Microsoft Windows: The octal mode that is translated into rights for Microsoft Windows security. Values up to `0777` are allowed (no sticky bits) and mean the same in Microsoft Windows as they do in UNIX, where `4` equals `GENERIC_READ`, `2` equals `GENERIC_WRITE`, and `1` equals `GENERIC_EXECUTE`. This attribute cannot be used to set `:full_control`. This attribute has no effect if not specified, but when this attribute and `rights` are both specified, the effects will be cumulative.

<code>owner</code>	A string or ID that identifies the group owner by user name, including fully qualified user names such as <code>domain\user</code> or <code>user@domain</code> . If this value is not specified, existing owners will remain unchanged and new owner assignments will use the current user (when necessary).
<code>path</code>	The path to the directory. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>recursive</code>	Use to create or delete parent directories recursively. For the <code>owner</code> , <code>group</code> , and <code>mode</code> attributes, the value of this attribute applies only to the leaf directory. Default value: <code>false</code> .
<code>rights</code>	Microsoft Windows only. The permissions for users and groups in a Microsoft Windows environment. For example: <code>rights &lt;permissions&gt;, &lt;principal&gt;, &lt;options&gt;</code> where <code>&lt;permissions&gt;</code> specifies the rights granted to the principal, <code>&lt;principal&gt;</code> is the group or user name, and <code>&lt;options&gt;</code> is a Hash with one (or more) advanced rights options.

#### Recursive Directories

The **directory** resource can be used to create directory structures, as long as each directory within that structure is created explicitly. This is because the `recursive` attribute only applies `group`, `mode`, and `owner` attribute values to the leaf directory.

A directory structure:

```
/foo
 /bar
 /baz
```

The following example shows a way create a file in the `/baz` directory:

```
directory "/foo/bar/baz" do
 owner "root"
 group "root"
 mode 00755
 action :create
end
```

But with this example, the `group`, `mode`, and `owner` attribute values will only be applied to `/baz`. Which is fine, if that's what you want. But most of the time, when the entire `/foo/bar/baz` directory structure is not there, you must be explicit about each directory. For example:

```
%w[/foo /foo/bar /foo/bar/baz].each do |path|
 directory path do
 owner "root"
 group "root"
 mode 00755
 end
end
```

This approach will create the correct hierarchy—`/foo`, then `/bar` in `/foo`, and then `/baz` in `/bar`—and also with the correct attribute values for `group`, `mode`, and `owner`.

A similar approach is required when changing the access permissions to directory objects, the owner of a file, or the group associated with a directory object. For example:

```
%w["/usr/local/**/*"].each do |path|
 file path do
 owner "root"
 group "root"
 end if File.file?(path)
 directory path do
 owner "root"
 group "root"
 end if File.directory?(path)
end
```

Though it should be noted that the previous example isn't a great approach when there are a large number of actions that will take place. Consider using the **execute** resource and/or a definition to handle use cases that need to support a large number of recursive actions.

#### Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Directory</code>	<code>directory</code>	The default provider for all platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Create a directory

```
directory "/tmp/something" do
 owner "root"
 group "root"
 mode 00755
 action :create
end
```

### Create a directory in Microsoft Windows

```
directory "C:\\tmp\\something.txt" do
 rights :full_control, "DOMAIN\\User"
 inherits false
 action :create
end
```

or:

```
directory 'C:\\tmp\\something.txt' do
 rights :full_control, 'DOMAIN\\User'
 inherits false
 action :create
end
```

#### Note

The difference between the two previous examples is the single- versus double-quoted strings, where if the double quotes are used, the backslash character (\\) must be escaped using the Ruby escape character (which is a backslash).

### Create a directory recursively

```
%w{dir1 dir2 dir3}.each do |dir|
 directory "/tmp/mydirs/#{dir}" do
 mode 00775
 owner "root"
 group "root"
 action :create
 recursive true
 end
end
```

### Delete a directory

```
directory "/tmp/something" do
 recursive true
 action :delete
end
```

### Set directory permissions using a variable

The following example shows how read/write/execute permissions can be set using a variable named `user_home` and then for owners and groups on any matching node:

```
user_home = "#{node[:matching_node][:user]}"

directory user_home do
 owner node[:matching_node][:user]
 group node[:matching_node][:group]
 mode "0755"
 action :create
end
```

where `matching_node` represents a type of node. For example, if the `user_home` variable specified `{node[:nginx]...}`, a recipe might look something like this:

```
user_home = "#{node[:nginx][:user]}"

directory user_home do
 owner node[:nginx][:user]
 group node[:nginx][:group]
 mode "0755"
 action :create
end
```

### Set directory permissions for a specific type of node

The following example shows how permissions can be set for the `/certificates` directory on any node that is running Nginx. In this example, permissions are being set for the owner and group as `"root"`, and then read/write permissions are granted to the root.

```
directory "#{node[:nginx][:dir]}/shared/certificates" do
```

```

owner "root"
group "root"
mode "700"
recursive true
end

```

### Reload the configuration

The following example shows how to reload the configuration of a chef-client using the **remote\_file** resource to:

- using an **if** statement to check whether the plugins on a node are the latest versions
- identify the location from which Ohai plugins are stored
- using the **notifies** attribute and a **ruby\_block** resource to trigger an update (if required) and to then reload the client.rb file.

```

directory node[:ohai][:plugin_path] do
 owner "chef"
 recursive true
end

ruby_block "reload_config" do
 block do
 Chef::Config.from_file("/etc/chef/client.rb")
 end
 action :nothing
end

if node[:ohai].key?(:plugins)
 node[:ohai][:plugins].each do |plugin|
 remote_file node[:ohai][:plugin_path] + "/" + plugin do
 source plugin
 owner "chef"
 notifies :create, resources(:ruby_block => "reload_config")
 end
 end
end
end

```

## dpkg\_package

The **dpkg\_package** resource is used to manage packages for the dpkg platform. When a package is installed from a local file, it must be added to the node using the **remote\_file** or **cookbook\_file** resources.

### Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

## Syntax

The syntax for using the **dpkg\_package** resource in a recipe is as follows:

```

dpkg_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end

```

where

- **dpkg\_package** tells the chef-client to use the `Chef::Provider::Dpkg` provider during the chef-client run
- **name** is the name of the resource block; when the **package\_name** attribute is not specified as part of a recipe, **name** is also the name of the package
- **attribute** is zero (or more) of the attributes that are available for this resource
- **:action** is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.

<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Dpkg</code>	<code>dpkg_package</code>	The provider that is used with the dpkg platform. Can be used with the <code>options</code> attribute.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package

```
dpkg_package "name of package" do
 action :install
end
```

easy\_install\_package

The `easy_install_package` resource is used to manage packages for the Python platform.

Note

In many cases, it is better to use the `package` resource instead of this one. This is because when the `package` resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the `package` resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

Syntax

The syntax for using the `easy_install_package` resource in a recipe is as follows:

```
easy_install_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `easy_install_package` tells the chef-client to use the `Chef::Provider::EasyInstall` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>easy_install_binary</code>	The location of the Easy Install binary.
<code>module_name</code>	The name of the module.
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>python_binary</code>	The location of the Python binary.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::EasyInstall</code>	<code>easy_install_package</code>	The provider that is used with Python platform.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Install a package

```
easy_install_package "name of package" do
 action :install
end
```

## env

The **env** resource is used to manage environment keys in Microsoft Windows. After an environment key is set, Microsoft Windows must be restarted before the environment key will be available to the Task Scheduler.

### Note

On UNIX-based systems, the best way to manipulate environment keys is with the `ENV` variable in Ruby; however, this approach does not have the same permanent effect as using the **env** resource.

## Syntax

The syntax for using the **env** resource in a recipe is as follows:

```
env "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `env` tells the chef-client to use the `Chef::Provider::Env::Windows` provider during the chef-client run
- `name` is the name of the resource block; when the `key_name` attribute is not specified as part of a recipe, `name` is also the name of the environment key that is created, deleted, or modified
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:



Action	Description
<code>:create</code>	Default. Use to create a new environment variable.
<code>:delete</code>	Use to delete an environment variable.
<code>:modify</code>	Use to modify an existing environment variable. This will append the new value to the existing value, using the delimiter specified by the <code>de<del>l</del>im</code> attribute.

Attributes

This resource has the following attributes:

Attribute	Description
<code>delim</code>	The delimiter that is used to separate multiple values for a single key.
<code>key_name</code>	The name of the key that will be created, deleted, or modified. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>value</code>	The value with which <code>key_name</code> is set.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Env::Windows</code>	<code>env</code>	The default provider for all Microsoft Windows platforms.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Set an environment variable

```
env "ComSpec" do
 value "C:\\Windows\\system32\\cmd.exe"
end
```

erl\_call

The `erl_call` resource is used to connect to a node located within a distributed Erlang system. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the `not_if` and `only_if` meta parameters to guard the use of this resource for idempotence.

Note

The `erl_call` command needs to be on the path for this resource to work properly.

Syntax

The syntax for using the `erl_call` resource in a recipe is as follows:

```
erl_call "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `erl_call` tells the chef-client to use the `Chef::Provider::ErlCall` provider during the chef-client run
- `"name"` is the name of the call
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:run</code>	Default. Indicates that the Erlang call should be run.
<code>:nothing</code>	Indicates that the Erlang call should not be run.

Attributes

This resource has the following attributes:

Attribute	Description
<code>code</code>	The code to be executed on a node located within a distributed Erlang system. Default value: <code>q()</code> .
<code>cookie</code>	The magic cookie for the node to which a connection is made.
<code>distributed</code>	Indicates that a node is a distributed Erlang node. Default value: <code>false</code> .
<code>name_type</code>	Indicates whether the <code>node_name</code> attribute is a short node name ( <code>sname</code> ) or a long node name ( <code>name</code> ). A node with a long node name cannot communicate with a node with a short node name. Default value: <code>sname</code> .
<code>node_name</code>	The hostname to which the node will connect. Default value: <code>chef@localhost</code> .
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::ErlCall</code>	<code>erl_call</code>	The default provider for all platforms.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Run a command

```
erl_call "list names" do
 code "net_adm:names()."
 distributed true
 node_name "chef@latte"
end
```

execute

The **execute** resource is used to execute a command. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the `not_if` and `only_if` meta parameters to guard the use of this resource for idempotence.

Note

Use the **script** resource to execute a script using a specific interpreter (Ruby, Python, Perl, csh, or Bash).

Syntax

The syntax for using the **execute** resource in a recipe is as follows:

```
execute "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `execute` tells the chef-client to use the `Chef::Provider::Execute` provider during the chef-client run
- `name` is the name of the resource block; when the `command` attribute is not specified as part of a recipe, `name` is also the command to be executed
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the **execute** resource can work when used in a recipe. In this example, a whitespace array is used to identify the names of the pets that will then be fed (by the command that is run):

```
%w{rover fido bubbers}.each do |pet_name|
 execute "feed_pet_#{pet_name}" do
 command "echo 'Feeding: #{pet_name}'; touch '/tmp/#{pet_name}'"
 not_if { ::File.exists?("/tmp/#{pet_name}") }
 end
end
```

## Actions

This resource has the following actions:

Action	Description
<code>:run</code>	Default. Indicates that the command should be run.
<code>:nothing</code>	Indicates that the command should not be run. This action is used to specify that a command is run only when another resource notifies it.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>command</code>	The name of the command to be executed. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>creates</code>	Indicates that a command to create a file will not be run when that file already exists.
<code>cwd</code>	The current working directory from which a command is run.
<code>environment</code>	A Hash of environment variables in the form of <code>{"ENV_VARIABLE" =&gt; "VALUE"}</code> . (These variables must exist for a command to be run successfully.)
<code>group</code>	The group name or group ID that must be changed before running a command.
<code>path</code>	An array of paths to use when searching for a command. These paths are not added to the command's environment \$PATH. The default value uses the system path.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>returns</code>	The return value for a command. This may be an array of accepted values. An exception is raised when the return value(s) do not match. Default value: <code>0</code> .
<code>timeout</code>	The amount of time (in seconds) a command will wait before timing out. Default value: <code>3600</code> .
<code>user</code>	The user name or user ID that should be changed before running a command.
<code>umask</code>	The file mode creation mask, or umask.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Execute</code>	<code>execute</code>	The default provider for all platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Run a command upon notification

```
execute "slapadd" do
 command "slapadd < /tmp/something.ldif"
 creates "/var/lib/slapd/uid.bdb"
 action :nothing
end
```

```
template "/tmp/something.ldif" do
 source "something.ldif"
 notifies :run, "execute[slapadd]"
end
```

### Run a touch file only once while running a command

```
execute "upgrade script" do
 command "php upgrade-application.php && touch /var/application/.upgraded"
 creates "/var/application/.upgraded"
 action :run
end
```

### Run a command which requires an environment variable

```
execute "slapadd" do
 command "slapadd < /tmp/something.ldif"
 creates "/var/lib/slapd/uid.bdb"
 action :run
 environment ({ 'HOME' => '/home/myhome' })
```

```
end
```

#### Delete a repository using yum to scrub the cache

```
the following code sample thanks to gaffney @ https://gist.github.com/918711

execute "clean-yum-cache" do
 command "yum clean all"
 action :nothing
end

file "/etc/yum.repos.d/bad.repo" do
 action :delete
 notifies :run, "execute[clean-yum-cache]", :immediately
 notifies :create, "ruby_block[reload-internal-yum-cache]", :immediately
end
```

#### Install repositories from a file, trigger a command, and force the internal cache to reload

The following example shows how to install new yum repositories from a file, where the installation of the repository triggers a creation of the yum cache that forces the internal cache for the chef-client to reload:

```
execute "create-yum-cache" do
 command "yum -q makecache"
 action :nothing
end

ruby_block "reload-internal-yum-cache" do
 block do
 Chef::Provider::Package::Yum::YumCache.instance.reload
 end
 action :nothing
end

cookbook_file "/etc/yum.repos.d/custom.repo" do
 source "custom"
 mode 00644
 notifies :run, "execute[create-yum-cache]", :immediately
 notifies :create, "ruby_block[reload-internal-yum-cache]", :immediately
end
```

#### Prevent restart and reconfigure if configuration is broken

Use the `:nothing` common action to prevent an application from restarting, and then use the `subscribes` notification to ask the broken configuration to be reconfigured immediately:

```
execute "test-nagios-config" do
 command "nagios3 --verify-config"
 action :nothing
 subscribes :run, "template[/etc/nagios3/configures-nagios.conf]", :immediately
end
```

#### Notify in a specific order

To notify multiple resources, and then have these resources run in a certain order, do something like the following:

```
execute 'foo' do
 command '...'
 notifies :run, 'template[baz]', :immediately
 notifies :install, 'package[bar]', :immediately
 notifies :run, 'execute[final]', :immediately
end

template 'baz' do
 ...
 notifies :run, 'execute[restart_baz]', :immediately
end

package 'bar'

execute 'restart_baz'

execute 'final' do
 command '...'
end
```

where the sequencing will be in the same order as the resources are listed in the recipe: `execute 'foo'`, `template 'baz'`, `execute [restart_baz]`, `package 'bar'`, and `execute 'final'`.

#### Execute a command using a template

The following example shows how to set up IPv4 packet forwarding using the `execute` resource to run a command named "forward\_ipv4" that uses a template defined by the `template` resource:

```
execute "forward_ipv4" do
 command "echo > /proc/.../ipv4/ip_forward"
 action :nothing
end

template "/etc/file_name.conf" do
 source "routing/file_name.conf.erb"
 notifies :run, 'execute[forward_ipv4]', :delayed
end
```

where the `command` attribute for the `execute` resource contains the command that is to be run and the `source` attribute for the `template`

resource specifies which template to use. The `notifies` attribute for the **template** specifies that the `execute[forward_ipv4]` (which is defined by the **execute** resource) should be queued up and run at the end of the chef-client run.

#### Add a rule to an IP table

The following example shows how to add a rule named "test\_rule" to an IP table using the **execute** resource to run a command using a template that is defined by the **template** resource:

```
execute 'test_rule' do
 command "command_to_run"
 --option value
 ...
 --option value
 --source #{node[:name_of_node][:ipsec][:local][:subnet]}
 -j test_rule"
 action :nothing
end

template "/etc/file_name.local" do
 source "routing/file_name.local.erb"
 notifies :run, 'execute[test_rule]', :delayed
end
```

where the `command` attribute for the **execute** resource contains the command that is to be run and the `source` attribute for the **template** resource specifies which template to use. The `notifies` attribute for the **template** specifies that the `execute[test_rule]` (which is defined by the **execute** resource) should be queued up and run at the end of the chef-client run.

#### Stop a service, do stuff, and then restart it

The following example shows how to use the **execute**, **service**, and **mount** resources together to ensure that a node running on Amazon EC2 is running MySQL. This example does the following:

- Checks to see if the Amazon EC2 node has MySQL
- If the node has MySQL, stops MySQL
- Installs MySQL
- Mounts the node
- Restarts MySQL

*# the following code sample comes from the ``server\_ec2`` recipe in the following cookbook: <https://github.com>*

```
if (node.attribute?('ec2') && ! FileTest.directory?(node['mysql']['ec2_path']))

 service "mysql" do
 action :stop
 end

 execute "install-mysql" do
 command "mv #{node['mysql']['data_dir']} #{node['mysql']['ec2_path']}"
 not_if do FileTest.directory?(node['mysql']['ec2_path']) end
 end

 [node['mysql']['ec2_path'], node['mysql']['data_dir']].each do |dir|
 directory dir do
 owner "mysql"
 group "mysql"
 end
 end

 mount node['mysql']['data_dir'] do
 device node['mysql']['ec2_path']
 fstype "none"
 options "bind,rw"
 action [:mount, :enable]
 end

 service "mysql" do
 action :start
 end
end
```

where

- the two **service** resources are used to stop, and then restart the MySQL service
- the **execute** resource is used to install MySQL
- the **mount** resource is used to mount the node and enable MySQL

#### Use the `platform_family?` method

The following is an example of using the `platform_family?` method in the Recipe DSL to create a variable that can be used with other resources in the same recipe. In this example, `platform_family?` is being used to ensure that a specific binary is used for a specific platform before using the **remote\_file** resource to download a file from a remote location, and then using the **execute** resource to install that file by running a command.

```
if platform_family?('rhel')
 pip_binary = "/usr/bin/pip"
else
 pip_binary = "/usr/local/bin/pip"
end

remote_file "#{Chef::Config[:file_cache_path]}/distribute_setup.py" do
 source "http://python-distribute.org/distribute_setup.py"
 mode "0644"
end
```

```

not_if { ::File.exists?(pip_binary) }
end

execute "install-pip" do
 cwd Chef::Config[:file_cache_path]
 command <<-EOF
 # command for installing Python goes here
 EOF
 not_if { ::File.exists?(pip_binary) }
end

```

where a command for installing Python might look something like:

```

#{node['python']['binary']} distribute_setup.py
#{::File.dirname(pip_binary)}/easy_install pip

```

#### Control a service using the execute resource

##### Warning

This is an example of something that should NOT be done. Use the **service** resource to control a service, not the **execute** resource.

Do something like this:

```

service "tomcat" do
 action :start
end

```

and NOT something like this:

```

execute "start-tomcat" do
 command "/etc/init.d/tomcat6 start"
 action :run
end

```

There is no reason to use the **execute** resource to control a service because the **service** resource exposes the `start_command` attribute directly, which gives a recipe full control over the command issued in a much cleaner, more direct manner.

#### Use the search recipe DSL method to find users

The following example shows how to use the `search` method in the Recipe DSL to search for users:

*# the following code sample comes from the openvpn cookbook: <https://github.com/opscode-cookbooks/openvpn>*

```

search("users", ".*:") do |u|
 execute "generate-openvpn-#{u['id']}" do
 command "./pktool #{u['id']}"
 cwd "/etc/openvpn/easy-rsa"
 environment(
 'EASY_RSA' => '/etc/openvpn/easy-rsa',
 'KEY_CONFIG' => '/etc/openvpn/easy-rsa/openssl.cnf',
 'KEY_DIR' => node["openvpn"]["key_dir"],
 'CA_EXPIRE' => node["openvpn"]["key"]["ca_expire"].to_s,
 'KEY_EXPIRE' => node["openvpn"]["key"]["expire"].to_s,
 'KEY_SIZE' => node["openvpn"]["key"]["size"].to_s,
 'KEY_COUNTRY' => node["openvpn"]["key"]["country"],
 'KEY_PROVINCE' => node["openvpn"]["key"]["province"],
 'KEY_CITY' => node["openvpn"]["key"]["city"],
 'KEY_ORG' => node["openvpn"]["key"]["org"],
 'KEY_EMAIL' => node["openvpn"]["key"]["email"]
)
 not_if { ::File.exists?(node["openvpn"]["key_dir"] + "/" + u['id'] + ".crt") }
 end

 %w{ conf ovpn }.each do |ext|
 template "#{node["openvpn"]["key_dir"]}/#{u['id']}.#{ext}" do
 source "client.conf.erb"
 variables :username => u['id']
 end
 end

 execute "create-openvpn-tar-#{u['id']}" do
 cwd node["openvpn"]["key_dir"]
 command <<-EOH
 tar zcf #{u['id']}.tar.gz \
 ca.crt #{u['id']}.crt #{u['id']}.key \
 #{u['id']}.conf #{u['id']}.ovpn \
 EOH
 not_if { ::File.exists?(node["openvpn"]["key_dir"] + "/" + u['id'] + ".tar.gz") }
 end
end
end

```

where

- the search will use both of the **execute** resources, unless the condition specified by the `not_if` commands are met
- the `environment` attribute in the first **execute** resource is being used to define values that appear as variables in the OpenVPN configuration
- the `template` resource tells the chef-client which template to use

#### Enable remote login for Mac OS X

```

execute "enable ssh" do
 command "/usr/sbin/systemsetup -setremotelogin on"
 not_if "/usr/sbin/systemsetup -getremotelogin | /usr/bin/grep On"
end

```

```
 action :run
end
```

Execute code immediately, based on the template resource

By default, notifications are `:delayed`, that is they are queued up as they are triggered, and then executed at the very end of a chef-client run. To run an action immediately, use `:immediately`:

```
template "/etc/nagios3/configures-nagios.conf" do
 # other parameters
 notifies :run, "execute[test-nagios-config]", :immediately
end
```

and then the chef-client would immediately run the following:

```
execute "test-nagios-config" do
 command "nagios3 --verify-config"
 action :nothing
end
```

file

The **file** resource is used to manage files that are present on a node, including setting or updating the contents of those files.

Note

Other resources should be used to manage files that are not present on a node. Use **cookbook\_file** when copying a file from a cookbook, **template** when using a template, and **remote\_file** when transferring files from remote locations.

Syntax

The syntax for using the **file** resource in a recipe is as follows:

```
file "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `file` tells the chef-client to use the `Chef::Provider::File` provider during the chef-client run
- `name` is the name of the resource block; when the `path` attribute is not specified as part of a recipe, `name` is also the path to the file
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the **file** resource can work when used in a recipe:

```
file "/tmp/something" do
 owner "root"
 group "root"
 mode "0755"
 action :create
end
```

Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a file.
<code>:create_if_missing</code>	Use to create a file only if the file does not exist. (When the file exists, nothing happens.)
<code>:delete</code>	Use to delete a file.
<code>:touch</code>	Use to touch a file. This updates the access (atime) and file modification (mtime) times for a file.

Attributes

This resource has the following attributes:

Attribute	Description
<code>atomic_update</code>	Indicates whether atomic file updates are used on a per-resource basis. Set to <code>true</code> for atomic file updates. Set to <code>false</code> for non-atomic file updates. (This setting overrides <code>file_atomic_update</code> , which is a global setting found in the <code>client.rb</code> file.) Default value: <code>true</code> .
<code>backup</code>	The number of backups to be kept. Set to <code>false</code> to prevent backups from being kept. Default value: <code>5</code> .
<code>content</code>	A string that is written to the file. The contents of this attribute will replace any previous content when this attribute has something other than the default value. The default behavior will not modify content.

<code>force_unlink</code>	Use to specify how the chef-client handles certain situations when the target file turns out not to be a file. For example, when a target file is actually a symlink. Set to <code>true</code> to have the chef-client delete the non-file target and replace it with the specified file. Set to <code>false</code> for the chef-client to raise an error. Default value: <code>false</code> .
<code>group</code>	A string or ID that identifies the group owner by group name, including fully qualified group names such as <code>domain\group</code> or <code>group@domain</code> . If this value is not specified, existing groups will remain unchanged and new group assignments will use the default <code>POSIX</code> group (if available).
<code>inherits</code>	Microsoft Windows only. Indicates that a file inherits rights from its parent. Default value: <code>true</code> .
<code>manage_symlink_source</code>	Indicates that the chef-client will detect and manage the source file for a symlink. Possible values: <code>nil</code> , <code>true</code> , or <code>false</code> . When this value is set to <code>nil</code> , the chef-client will manage a symlink's source file and emit a warning. When this value is set to <code>true</code> , the chef-client will manage a symlink's source file and not emit a warning. Default value: <code>nil</code> . The default value will be changed to <code>false</code> in a future version.
<code>mode</code>	<p>The octal mode for a file. If <code>mode</code> is not specified and if the file already exists, the existing mode on the file is used. If <code>mode</code> is not specified, the file does not exist, and the <code>:create</code> action is specified, the chef-client will assume a mask value of <code>0777</code> and then apply the umask for the system on which the file will be created to the mask value. For example, if the umask on a system is <code>022</code>, the chef-client would use the default value of <code>0755</code>.</p> <p>The behavior is different depending on the platform.</p> <p>UNIX- and Linux-based systems: The octal mode that is passed to <code>chmod</code>. If the value is specified as a quoted string, it will work exactly as if the <code>chmod</code> command was passed. If the value is specified as an integer, prepend a zero (<code>0</code>) to the value to ensure it is interpreted as an octal number. For example, to assign read, write, and execute rights for all users, use <code>0777</code> or <code>'777'</code>; for the same rights, plus the sticky bit, use <code>01777</code> or <code>'1777'</code>.</p> <p>Microsoft Windows: The octal mode that is translated into rights for Microsoft Windows security. Values up to <code>0777</code> are allowed (no sticky bits) and mean the same in Microsoft Windows as they do in UNIX, where <code>4</code> equals <code>GENERIC_READ</code>, <code>2</code> equals <code>GENERIC_WRITE</code>, and <code>1</code> equals <code>GENERIC_EXECUTE</code>. This attribute cannot be used to set <code>:full_control</code>. This attribute has no effect if not specified, but when this attribute and <code>rights</code> are both specified, the effects will be cumulative.</p>
<code>owner</code>	A string or ID that identifies the group owner by user name, including fully qualified user names such as <code>domain\user</code> or <code>user@domain</code> . If this value is not specified, existing owners will remain unchanged and new owner assignments will use the current user (when necessary).
<code>path</code>	<p>The path to the file. Default value: the name of the resource block (see Syntax section above).</p> <p>Microsoft Windows: A path that begins with a forward slash (<code>/</code>) will point to the root of the current working directory of the chef-client process. This path can vary from system to system. Therefore, using a path that begins with a forward slash (<code>/</code>) is not recommended.</p>
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>rights</code>	Microsoft Windows only. The permissions for users and groups in a Microsoft Windows environment. For example: <code>rights &lt;permissions&gt;, &lt;principal&gt;, &lt;options&gt;</code> where <code>&lt;permissions&gt;</code> specifies the rights granted to the principal, <code>&lt;principal&gt;</code> is the group or user name, and <code>&lt;options&gt;</code> is a Hash with one (or more) advanced rights options.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::File</code>	<code>file</code>	The default provider for all platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Create a file

```
file "/tmp/something" do
 owner "root"
 group "root"
 mode 00755
 action :create
end
```

### Create a file in Microsoft Windows



```
file "C:\tmp\something.txt" do
 rights :read, "Everyone"
 rights :full_control, "DOMAIN\User"
 action :create
end
```

#### Remove a file

```
file "/tmp/something" do
 action :delete
end
```

#### Set file modes

```
file "/tmp/something" do
 mode "644"
end
```

or:

```
file "/tmp/something" do
 mode 00644
end
```

#### Delete a repository using yum to scrub the cache

```
the following code sample thanks to gaffney @ https://gist.github.com/918711

execute "clean-yum-cache" do
 command "yum clean all"
 action :nothing
end

file "/etc/yum.repos.d/bad.repo" do
 action :delete
 notifies :run, "execute[clean-yum-cache]", :immediately
 notifies :create, "ruby_block[reload-internal-yum-cache]", :immediately
end
```

## freebsd\_package

The **freebsd\_package** resource is used to manage packages for the FreeBSD platform.

#### Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

## Syntax

The syntax for using the **freebsd\_package** resource in a recipe is as follows:

```
freebsd_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `freebsd_package` tells the chef-client to use the `Chef::Provider::Freebsd` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:remove</code>	Use to remove a package.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.

<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Freebsd</code>	<code>freebsd_package</code>	The provider that is used with the FreeBSD platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package

```
freebsd_package "name of package" do
 action :install
end
```

gem\_package

The **gem\_package** resource is used to manage gem packages that are only included in recipes. When a package is installed from a local file, it must be added to the node using the **remote\_file** or **cookbook\_file** resources.

Warning

The **chef\_gem** and **gem\_package** resources are both used to install Ruby gems. For any machine on which the chef-client is installed, there are two instances of Ruby. One is the standard, system-wide instance of Ruby and the other is a dedicated instance that is available only to the chef-client. Use the **chef\_gem** resource to install gems into the instance of Ruby that is dedicated to the chef-client. Use the **gem\_package** resource to install all other gems (i.e. install gems system-wide).

Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

Syntax

The syntax for using the **gem\_package** resource in a recipe is as follows:

```
gem_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `gem_package` tells the chef-client to use the `Chef::Provider::Rubygems` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Gem Package Options

The RubyGems package provider attempts to use the RubyGems API to install gems without spawning a new process, whenever possible. A gems command to install will be spawned under the following conditions:

- When a `gem_binary` attribute is specified (as a hash, a string, or by a .gemrc file), the provider will run that command to examine its environment settings and then again to install the gem.
- When install options are specified as a string, the provider will span a gems command with those options when installing the gem.

- The omnibus installer will search the PATH for a gem command rather than defaulting to the current gem environment. As part of `enforce_path_sanity`, the `bin` directories area added to the PATH, which means when there are no other proceeding RubyGems, the installation will still be operated against it.

#### Use a Hash

If an explicit `gem_binary` parameter is not being used with the `gem_package` resource, it is preferable to provide the install options as a hash. This approach allows the provider to install the gem without needing to spawn an external gem process.

The following RubyGems options are available for inclusion within a hash and are passed to the RubyGems DependencyInstaller:

- `:env_shebang`
- `:force`
- `:format_executable`
- `:ignore_dependencies`
- `:prerelease`
- `:security_policy`
- `:wrappers`

For more information about these options, see the RubyGems documentation: <http://rubygems.rubyforge.org/rubygems-update/Gem/DependencyInstaller.html>.

#### Example

```
gem_package "bundler" do
 options(:prerelease => true, :format_executable => false)
end
```

#### Use a String

When using an explicit `gem_binary`, options must be passed as a string. When not using an explicit `gem_binary`, the chef-client is forced to spawn a gems process to install the gems (which uses more system resources) when options are passed as a string. String options are passed verbatim to the gems command and should be specified just as if they were passed on a command line. For example, `--prerelease` for a pre-release gem.

#### Example

```
gem_package "nokogiri" do
 gem_binary("/opt/ree/bin/gem")
 options("--prerelease --no-format-executable")
end
```

#### Use a .gemrc File

Options can be specified in a `.gemrc` file. By default the `gem_package` resource will use the Ruby interface to install gems which will ignore the `.gemrc` file. The `gem_package` resource can be forced to use the gems command instead (and to read the `.gemrc` file) by adding the `gem_binary` attribute to a code block.

#### Example

```
gem_package "nokogiri" do
 gem_binary "gem"
end
```

### Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:reconfig</code>	Use to reconfigure a package. This action requires a response file.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

### Attributes

This resource has the following attributes:

Attribute	Description
<code>gem_binary</code>	An attribute for the <code>gem_package</code> provider that is used to specify a gems binary. This attribute is useful when installing Ruby 1.9 gems while running in Ruby 1.8. By default, the same version of Ruby that is used by the chef-client will be installed.

<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The URL at which the gem package is located.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Rubygems</code>	<code>gem_package</code>	Can be used with the <code>options</code> attribute.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a gems file from the local file system

```
gem_package "right_aws" do
 source "/tmp/right_aws-1.11.0.gem"
 action :install
end
```

Use the ignore\_failure common attribute

```
gem_package "syntax" do
 action :install
 ignore_failure true
end
```

git

The `git` resource is used to manage source control resources that exist in a git repository. git version 1.6.5 (or higher) is required to use all of the functionality in the `git` resource.

Note

This resource is often used in conjunction with the `deploy` resource.

Syntax

The syntax for using the git resource in a recipe is as follows:

```
git "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `git` tells the chef-client to use the `Chef::Provider::Git` provider during the chef-client run.
- `"name"` is the location in which the source files will be placed and/or synchronized with the files under source control management
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example shows the git resource:

```
git "#{Chef::Config[:file_cache_path]}/app_name" do
 repository node[:app_name][:git_repository]
 revision node[:app_name][:git_revision]
 action :sync
 notifies :run, "bash[compile_app_name]"
end
```

where

- the name of the resource is `#{Chef::Config[:file_cache_path]}/libvpx`
- the `repository` and `reference` nodes tell the chef-client which repository and revision to use

## Actions

This resource has the following actions:

Action	Description
<code>:sync</code>	Default. Use to update the source to the specified version, or to get a new clone or checkout.
<code>:checkout</code>	Use to clone or check out the source. When a checkout is available, this provider does nothing.
<code>:export</code>	Use to export the source, excluding or removing any version control artifacts.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>additional_remotes</code>	An array of additional remotes that are added to the git repository configuration.
<code>checkout_branch</code>	Use to specify the name of a branch to be checked out. Default value: <code>deploy</code> .
<code>depth</code>	The number of past revisions that will be included in the git shallow clone. The default behavior will do a full clone.
<code>destination</code>	The path to the location to which the source will be cloned, checked out, or exported. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>enable_checkout</code>	Use to check out a repo from master. Default value: <code>true</code> .
<code>enable_submodules</code>	Use to perform a sub-module initialization and update. Default value: <code>false</code> .
<code>group</code>	The system group that is responsible for the checked-out code.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>reference</code>	The alias for revision.
<code>remote</code>	The remote repository to be used when synchronizing an existing clone.
<code>repository</code>	The URI for the git repository.
<code>revision</code>	The revision to be checked out. This can be symbolic, like <code>HEAD</code> or it can be a source control management-specific revision identifier. Default value: <code>HEAD</code> .
<code>ssh_wrapper</code>	The path to the wrapper script used when running SSH with git. The <code>GIT_SSH</code> environment variable is set to this.
<code>user</code>	The system user that is responsible for the checked-out code.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Git</code>	<code>git</code>	This provider works only with git.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Use the git mirror

```
git "/opt/mysources/couch" do
 repository "git://git.apache.org/couchdb.git"
 revision "master"
 action :sync
end
```

### Use different branches

To use different branches, depending on the environment of the node:

```
if node.chef_environment == "QA"
 branch_name = "staging"
else
 branch_name = "master"
end

git "/home/user/deployment" do
```

```
repository "git@github.com:git:site/deployment.git"
revision branch_name
action :sync
user "user"
group "test"
end
```

where the `branch_name` variable is set to `staging` or `master`, depending on the environment of the node. Once this is determined, the `branch_name` variable is used to set the revision for the repository. If the `git status` command is used after running the example above, it will return the branch name as `deploy`, as this is the default value. Run the `chef-client` in debug mode to verify that the correct branches are being checked out:

```
$ sudo chef-client -l debug
```

Install an application from git using bash

The following example shows how `Bash` can be used to install a plug-in for `rvm` named "ruby-build", which is located in `git` version source control. First, the application is synchronized, and then `Bash` changes its working directory to the location in which "ruby-build" is located, and then runs a command.

```
git "#{Chef::Config[:file_cache_path]}/ruby-build" do
 repository "git://github.com/sstephenson/ruby-build.git"
 reference "master"
 action :sync
end

bash "install_ruby_build" do
 cwd "#{Chef::Config[:file_cache_path]}/ruby-build"
 user "rvm"
 group "rvm"
 code <<-EOH
 ./install.sh
 EOH
 environment 'PREFIX' => "/usr/local"
end
```

To read more about `ruby-build`, see here: <https://github.com/sstephenson/ruby-build>.

Upgrade packages from git

The following example shows the `scm` resource using the `git` short name as part of a larger recipe that is used to upgrade packages:

```
the following code sample comes from the `source` recipe in the `libvpx-cookbook` cookbook: https:
git "#{Chef::Config[:file_cache_path]}/libvpx" do
 repository node[:libvpx][:git_repository]
 revision node[:libvpx][:git_revision]
 action :sync
 notifies :run, "bash[compile_libvpx]"
end
```



group

The `group` resource is used to manage a local group.

Syntax

The syntax for using the `group` resource in a recipe is as follows:

```
group "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `group` tells the chef-client to use one of the following providers during the chef-client run: `Chef::Provider::Group`, `Chef::Provider::Group::Aix`, `Chef::Provider::Group::Dsl`, `Chef::Provider::Group::Gpasswd`, `Chef::Provider::Group::Groupadd`, `Chef::Provider::Group::Groupmod`, `Chef::Provider::Group::Pw`, `Chef::Provider::Group::Suse`, `Chef::Provider::Group::Usermod`, or `Chef::Provider::Group::Windows`. The provider that is used by the chef-client depends on the platform of the machine on which the chef-client run is taking place
- `name` is the name of the resource block; when the `group_name` attribute is not specified as part of a recipe, `name` is also the name of the group
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a group.
<code>:remove</code>	Use to remove a group.

<code>:modify</code>	Use to modify an existing group. This action will raise an exception if the group does not exist.
<code>:manage</code>	Use to manage an existing group. This action will do nothing if the group does not exist.

Attributes

This resource has the following attributes:

Attribute	Description
<code>append</code>	Use to specify how members should be appended and/or removed from a group. When <code>true</code> , <code>members</code> will be appended and <code>excluded_members</code> will be removed. When <code>false</code> , group members will be reset to the value of the <code>members</code> attribute. Default value: <code>false</code> .
<code>excluded_members</code>	Use to remove users from a group. May only be used when <code>append</code> is set to <code>true</code> .
<code>gid</code>	The identifier for the group.
<code>group_name</code>	The name of the group. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>members</code>	Indicates which users should be set or appended to a group.
<code>non_unique</code>	Indicates that <code>gid</code> duplication is allowed. May only be used with the <code>Groupadd</code> provider. Default value: <code>false</code> .
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>system</code>	Indicates whether a group is a system group ( <code>true</code> ) or is not a system group ( <code>false</code> ).

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Group</code>	<code>group</code>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<code>Chef::Provider::Group::Aix</code>	<code>group</code>	The provider that is used with the AIX platform.
<code>Chef::Provider::Group::DscL</code>	<code>group</code>	The provider that is used with the Mac OS X platform.
<code>Chef::Provider::Group::Gpasswd</code>	<code>group</code>	The provider that is used with the <code>gpasswd</code> command.
<code>Chef::Provider::Group::Groupadd</code>	<code>group</code>	The provider that is used with the <code>groupadd</code> command.
<code>Chef::Provider::Group::Groupmod</code>	<code>group</code>	The provider that is used with the <code>groupmod</code> command.
<code>Chef::Provider::Group::Pw</code>	<code>group</code>	The provider that is used with the FreeBSD platform.
<code>Chef::Provider::Group::Suse</code>	<code>group</code>	The provider that is used with the SuSE platform.
<code>Chef::Provider::Group::Usermod</code>	<code>group</code>	The provider that is used with the Solaris platform.
<code>Chef::Provider::Group::Windows</code>	<code>group</code>	The provider that is used with the Microsoft Windows platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Append users to groups

```
group "www-data" do
 action :modify
 members "maintenance"
 append true
end
```

http\_request

The `http_request` resource is used to send an HTTP request (GET, PUT, POST, DELETE, HEAD, or OPTIONS) with an arbitrary message. This resource is useful when custom callbacks are necessary.

Syntax

The syntax for using the `http_request` resource in a recipe is as follows:

```
http_request "name" do
```

```

url "http://opscode.com/patn"
attribute "value" # see attributes section below
...
action :action # see actions section below
end

```

where

- `http_request` tells the chef-client to use the `Chef::Provider::HttpRequest` provider during the chef-client run
- `name` is the name of the resource block; when the `message` attribute is not specified as part of a recipe, `name` is also the message that is sent by the HTTP request
- `attribute` is zero (or more) of the attributes that are available for this resource
- `url` is the URL that will precede `?message=` in the HTTP request
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the `http_request` resource can work when used in a recipe. In this example, the following example will send a `DELETE` request to `"http://www.opscode.com/some_page?message=please_delete_me"`.

```

http_request "please_delete_me" do
 url "http://www.opscode.com/some_page"
 action :delete
end

```

## Actions

This resource has the following actions:

Action	Description
<code>:get</code>	Default. Use to send a <code>GET</code> request.
<code>:put</code>	Use to send a <code>PUT</code> request.
<code>:post</code>	Use to send a <code>POST</code> request.
<code>:delete</code>	Use to send a <code>DELETE</code> request.
<code>:head</code>	Use to send a <code>HEAD</code> request.
<code>:options</code>	Use to send an <code>OPTIONS</code> request.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>headers</code>	A Hash of custom headers. Default value: <code>{}</code> .
<code>message</code>	The message that is sent by the HTTP request. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>url</code>	The URL to which an HTTP request is sent.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::HttpRequest</code>	<code>http_request</code>	The default provider for all platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Send a GET request

```

http_request "some_message" do
 url "http://example.com/check_in"
end

```

The message is sent as `"http://example.com/check_in?message=some_message"`.

### Send a POST request

To send a `POST` request as JSON data, convert the message to JSON and include the correct content-type header. For example:



```
http_request "posting data" do
 action :post
 url "http://example.com/check_in"
 message ({:some => "data"}.to_json)
 headers({"AUTHORIZATION" => "Basic #{Base64.encode64("username:password")}", "Content-Type" => "applicat
end
```

Transfer a file only when the remote source changes

```
remote_file "/tmp/couch.png" do
 source "http://couchdb.apache.org/img/sketch.png"
 action :nothing
end

http_request "HEAD http://couchdb.apache.org/img/sketch.png" do
 message ""
 url "http://couchdb.apache.org/img/sketch.png"
 action :head
 if File.exists?("/tmp/couch.png")
 headers "If-Modified-Since" => File.mtime("/tmp/couch.png").httpdate
 end
 notifies :create, "remote_file[/tmp/couch.png]", :immediately
end
```

ifconfig

The **ifconfig** resource is used to manage interfaces.

Syntax

The syntax for using the **ifconfig** resource in a recipe is as follows:

```
ifconfig "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `ifconfig` tells the chef-client to use the `Chef::Provider::Ifconfig` provider during the chef-client run
- `name` is the name of the resource block; when the `target` attribute is not specified as part of a recipe, `name` is also the IP address that will be assigned to the network interface
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:add</code>	Default. Use to run ifconfig to configure a network interface and (on some platforms) write a configuration file for that network interface.
<code>:delete</code>	Use to run ifconfig to disable a network interface and (on some platforms) delete that network interface's configuration file.
<code>:enable</code>	Use to run ifconfig to enable a network interface.
<code>:disable</code>	Use to run ifconfig to disable a network interface.

Attributes

This resource has the following attributes:

Attribute	Description
<code>bcast</code>	The broadcast address for a network interface. On some platforms this attribute is not set using ifconfig, but is instead added to the startup configuration file for the network interface.
<code>bootproto</code>	The boot protocol used by a network interface.
<code>device</code>	The network interface to be configured.
<code>hwaddr</code>	The hardware address for the network interface.
<code>inet_addr</code>	The Internet host address for the network interface.
<code>mask</code>	The decimal representation of the network mask. For example: <code>255.255.255.0</code> .
<code>metric</code>	The routing metric for the interface.
<code>mtu</code>	The maximum transmission unit (MTU) for the network interface.

<u>network</u>	The address for the network interface.
<u>onboot</u>	Indicates that the network interface should be brought up on boot when this value is set to <u>yes</u> .
<u>onparent</u>	Indicates that the network interface should be brought up when its parent interface is brought up when this value is set to <u>yes</u> .
<u>provider</u>	Optional. Use to specify a provider by using its long name. For example: <u>provider</u> <u>Chef::Provider::Long::Name</u> . See the Providers section below for the list of providers available to this resource.
<u>target</u>	The IP address that will be assigned to the network interface. Default value: the <u>name</u> of the resource block (see Syntax section above).

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<u>Chef::Provider::Ifconfig</u>	<u>ifconfig</u>	The default provider for all platforms. Currently, this provider only writes out a start-up configuration file for the interface on Red Hat-based platforms (it writes to <u>/etc/sysconfig/network-scripts/ifcfg-#{device_name}</u> ).

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Configure a network interface

```
ifconfig "192.186.0.1" do
 device "eth0"
end
```

ips\_package

The **ips\_package** resource is used to manage packages (using Image Packaging System (IPS)) on the Solaris 11 platform.

Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

Syntax

The syntax for using the **ips\_package** resource in a recipe is as follows:

```
ips_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- ips\_package tells the chef-client to use the Chef::Provider::Ips provider during the chef-client run
- name is the name of the resource block; when the package\_name attribute is not specified as part of a recipe, name is also the name of the package
- attribute is zero (or more) of the attributes that are available for this resource
- :action is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<u>:install</u>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<u>:upgrade</u>	Use to install a package and/or to ensure that a package is the latest version.
<u>:remove</u>	Use to remove a package.

Attributes

This resource has the following attributes:

Attribute	Description
<code>accept_license</code>	Indicates that an end-user license agreement will be accepted automatically. Default value: <code>false</code> .
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Ips</code>	<code>ips_package</code>	The provider that is used with the ips platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package

```
ips_package "name of package" do
 action :install
end
```

link

The `link` resource is used to create symbolic or hard links.

Syntax

The syntax for using the `link` resource in a recipe is as follows:

```
link "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `link` tells the chef-client to use the `Chef::Provider::Link` provider during the chef-client run
- `name` is the name of the resource block; when the `target_file` attribute is not specified as part of a recipe, `name` is also name of the link
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a link.
<code>:delete</code>	Use to delete a link.

Attributes

This resource has the following attributes:

Attribute	Description
<code>group</code>	A string or ID that identifies the group associated with a symbolic link.
<code>link_type</code>	The type of link: <code>:symbolic</code> or <code>:hard</code> . Default value: <code>symbolic</code> .

<code>owner</code>	The owner associated with a symbolic link.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>target_file</code>	The name of the link. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>to</code>	The actual file to which the link will be created.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Link</code>	<code>link</code>	The default provider for all platforms.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Create symbolic links

```
link "/tmp/passwd" do
 to "/etc/passwd"
end
```

Create hard links

```
link "/tmp/passwd" do
 to "/etc/passwd"
 link_type :hard
end
```

Delete links

```
link "/tmp/mylink" do
 action :delete
 only_if "test -L /tmp/mylink"
end
```

log

The `log` resource is used to create log entries from a recipe.

Syntax

The syntax for using the `log` resource in a recipe is as follows:

```
log "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `log` tells the chef-client to use the `Chef::Provider::Log::ChefLog` provider during the chef-client run
- `name` is the name of the resource block; when the `message` attribute is not specified as part of a recipe, `name` is also the message to be added to a log file
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:write</code>	Default. Use to write to log.

Attributes

This resource has the following attributes:

Attribute	Description
<code>level</code>	The level of logging that will be displayed by the chef-client. The chef-client uses the <code>mixlib-log</code> ( <a href="https://github.com/opscode/mixlib-log">https://github.com/opscode/mixlib-log</a> ) to handle logging behavior. Options (in order of priority): <code>:debug</code> , <code>:info</code> , <code>:warn</code> , <code>:error</code> , and <code>:fatal</code> . Default value: <code>:info</code> .

<code>message</code>	The message to be added to a log file. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Log::Name</code> . See the Providers section below for the list of providers available to this resource.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Log::ChefLog</code>	<code>Log</code>	The default provider for all platforms.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Set default logging level

```
log "your string to log"
```

Set debug logging level

```
log "a debug string" do
 level :debug
end
```

Add a message to a log file

```
log "message" do
 message "This is the message that will be added to the log."
 level :info
end
```

macports\_package

The `macports_package` resource is used to manage packages for the Mac OS X platform.

Note

In many cases, it is better to use the `package` resource instead of this one. This is because when the `package` resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the `package` resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

Syntax

The syntax for using the `macports_package` resource in a recipe is as follows:

```
macports_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `macports_package` tells the chef-client to use the `Chef::Provider::Macports` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Macports</code>	<code>macports_package</code>	The provider that is used with the Mac OS X platform.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Install a package

```
macports_package "name of package" do
 action :install
end
```

## mdadm

The **mdadm** resource is used to manage RAID devices in a Linux environment using the mdadm utility. The **mdadm** provider will create and assemble an array, but it will not create the config file that is used to persist the array upon reboot. If the config file is required, it must be done by specifying a template with the correct array layout, and then by using the **mount** provider to create a file systems table (fstab) entry.

## Syntax

The syntax for using the **mdadm** resource in a recipe is as follows:

```
mdadm "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- **mdadm** tells the chef-client to use the `Chef::Provider::Mdadm` provider during the chef-client run
- `name` is the name of the resource block; when the `raid_device` attribute is not specified as part of a recipe, `name` is also the name of the RAID device
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a new array with per-device superblocks.
<code>:assemble</code>	Use to assemble a previously created array into an active array.
<code>:stop</code>	Use to stop an active array.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>bitmap</code>	The path to a file in which a write-intent bitmap is stored.
<code>chunk</code>	The chunk size. This attribute should not be used for a RAID 1 mirrored pair (i.e. when the <code>level</code> attribute is set to 1). Default value: <code>16</code> .
<code>devices</code>	A comma-separated list of devices to be part of a RAID array. Default value: <code>[]</code> .
<code>exists</code>	Indicates whether the RAID array exists. Default value: <code>false</code> .
<code>level</code>	The RAID level. Default value: <code>1</code> .
<code>metadata</code>	The superblock type for RAID metadata. Default value: <code>0.90</code> .
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>raid_device</code>	The name of the RAID device. Default value: the <code>name</code> of the resource block (see Syntax section above).

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Mdadm</code>	<code>mdadm</code>	The default provider for the Linux platform.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Create and assemble a RAID 0 array

The `mdadm` command can be used to create RAID arrays. For example, a RAID 0 array named `/dev/md0` with 10 devices would have a command similar to the following:

```
$ mdadm --create /dev/md0 --level=0 --raid-devices=10 /dev/s01.../dev/s10
```

where `/dev/s01 .. /dev/s10` represents 10 devices (01, 02, 03, and so on). This same command, when expressed as a recipe using the `mdadm` resource, would be similar to:

```
mdadm "/dev/md0" do
 devices ["/dev/s01", ... "/dev/s10"]
 level 0
 action :create
end
```

(again, where `/dev/s01 .. /dev/s10` represents devices `/dev/s01`, `/dev/s02`, `/dev/s03`, and so on).

### Create and assemble a RAID 1 array

```
mdadm "/dev/md0" do
 devices ["/dev/sda", "/dev/sdb"]
 level 1
 action [:create, :assemble]
end
```

### Create and assemble a RAID 5 array

The `mdadm` command can be used to create RAID arrays. For example, a RAID 5 array named `/dev/sd0` with 4, and a superblock type of `0.90` would be similar to:

```
mdadm "/dev/sd0" do
 devices ["/dev/s1", "/dev/s2", "/dev/s3", "/dev/s4"]
 level 5
 metadata "0.90"
 chunk 32
 action :create
end
```

## mount

The `mount` resource is used to manage a mounted file system.

## Syntax

The syntax for using the **mount** resource in a recipe is as follows:

```
mount "name" do
 attribute "value" # see attributes section below
 ...
 fstype "type"
 action :action # see actions section below
end
```

where

- **mount** tells the chef-client to use the `Chef::Provider::Mount` provider during the chef-client run for all platforms except for Microsoft Windows, which uses the `Chef::Provider::Mount::Windows` provider
- **name** is the name of the resource block; when the `mount_point` attribute is not specified as part of a recipe, **name** is also the directory (or path) in which a device should be mounted
- **attribute** is zero (or more) of the attributes that are available for this resource
- **fstype** is the file system type; this attribute is required
- **:action** is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the **mount** resource can work when used in a recipe:

```
mount node['mysql']['ec2_path'] do
 device ebs_vol_dev
 fstype "xfs"
 action :mount
end
```

## Actions

This resource has the following actions:

Action	Description
<code>:mount</code>	Default. Use to mount a device.
<code>:umount</code>	Use to unmount a device.
<code>:remount</code>	Use to remount a device.
<code>:enable</code>	Use to add an entry to the file systems table (fstab).
<code>:disable</code>	Use to remove an entry from the file systems table (fstab).

### Note

Order matters when passing multiple actions. For example: `action [:mount, :enable]` ensures that the file system is mounted before it is enabled.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>device</code>	Required for <code>:umount</code> and <code>:remount</code> actions (for the purpose of checking the mount command output for presence). The special block device or remote node, a label, or a uuid to be mounted.
<code>device_type</code>	The type of device: <code>:device</code> , <code>:label</code> , or <code>:uuid</code> . Default value: <code>:device</code> .
<code>domain</code>	Microsoft Windows only. Use to specify the domain in which the <code>username</code> and <code>password</code> are located.
<code>dump</code>	The dump frequency (in days) used while creating a file systems table (fstab) entry. Default value: <code>0</code> .
<code>enabled</code>	Use to specify if a mounted file system is enabled. Default value: <code>false</code> .
<code>fstype</code>	Required. The file system type (fstype) of the device.
<code>mount_point</code>	The directory (or path) in which the device should be mounted. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>mounted</code>	Use to specify if a file system is already mounted. Default value: <code>false</code> .
<code>options</code>	An array or string that contains mount options. If this value is a string, it will be converted to an array. Default value: <code>defaults</code> .
<code>pass</code>	The pass number used by the file system check ( <code>fsck</code> ) command while creating a file systems table (fstab) entry. Default value: <code>2</code> .
<code>password</code>	Microsoft Windows only. Use to specify the password for <code>username</code> .
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available



<code>supports</code>	A Hash of options for supported mount features. Default value: <code>{ :remount =&gt; false }</code> .
<code>username</code>	Microsoft Windows only. Use to specify the user name.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Mount</code>	<code>mount</code>	The default provider for all platforms, except for Microsoft Windows.
<code>Chef::Provider::Mount::Windows</code>	<code>mount</code>	The default provider for the Microsoft Windows platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Mount a labeled file system

```
mount "/mnt/volume1" do
 device "volume1"
 device_type :label
 fstype "xfs"
 options "rw"
end
```

Mount a local block drive

```
mount "/mnt/local" do
 device "/dev/sdb1"
 fstype "ext3"
end
```

Mount a non-block file system

```
mount "/mount/tmp" do
 pass 0
 fstype "tmpfs"
 device "/dev/null"
 options "nr_inodes=999k,mode=755,size=500m"
 action [:mount, :enable]
end
```

Mount and add to the file systems table

```
mount "/export/www" do
 device "naslprod:/export/web_sites"
 fstype "nfs"
 options "rw"
 action [:mount, :enable]
end
```

Mount a remote file system

```
mount "/export/www" do
 device "naslprod:/export/web_sites"
 fstype "nfs"
 options "rw"
end
```

Mount a remote folder in Microsoft Windows

```
mount "T:" do
 action :mount
 device "\\hostname.example.com\folder"
end
```

Unmount a remote folder in Microsoft Windows

```
mount "T:" do
 action :umount
 device "\\hostname.example.com\D$"
end
```

Stop a service, do stuff, and then restart it

The following example shows how to use the `execute`, `service`, and `mount` resources together to ensure that a node running on Amazon EC2 is running MySQL. This example does the following:

- Checks to see if the Amazon EC2 node has MySQL
- If the node has MySQL, stops MySQL
- Installs MySQL
- Mounts the node
- Restarts MySQL

# the following code sample comes from the ``server\_ec2`` recipe in the following cookbook: [https://github.com/chef-cookbooks/server\\_ec2](https://github.com/chef-cookbooks/server_ec2)

```

if (node.attribute?('ec2') && ! FileTest.directory?(node['mysql']['ec2_path']))

 service "mysql" do
 action :stop
 end

 execute "install-mysql" do
 command "mv #{node['mysql']['data_dir']} #{node['mysql']['ec2_path']}"
 not_if do FileTest.directory?(node['mysql']['ec2_path']) end
 end

 [node['mysql']['ec2_path'], node['mysql']['data_dir']].each do |dir|
 directory dir do
 owner "mysql"
 group "mysql"
 end
 end

 mount node['mysql']['data_dir'] do
 device node['mysql']['ec2_path']
 fstype "none"
 options "bind,rw"
 action [:mount, :enable]
 end

 service "mysql" do
 action :start
 end
end

```

where

- the two **service** resources are used to stop, and then restart the MySQL service
- the **execute** resource is used to install MySQL
- the **mount** resource is used to mount the node and enable MySQL

## ohai

The **ohai** resource is used to reload the Ohai configuration on a node, which allows recipes that change system attributes (like adding a user) to refer to those attributes later on during the chef-client run.

### Syntax

The syntax for using the **ohai** resource in a recipe is as follows:

```

ohai "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end

```

where

- **ohai** tells the chef-client to use the `Chef::Provider::Ohai` provider during the chef-client run
- **"name"** is a friendly name for the action that is defined in the recipe
- **attribute** is zero (or more) of the attributes that are available for this resource
- **:action** is the step that the resource will ask the provider to take during the chef-client run

### Actions

This resource has the following actions:

Action	Description
<code>:reload</code>	Default. Use to reload the Ohai configuration on a node.

### Attributes

This resource has the following attributes:

Attribute	Description
<code>name</code>	Always the same value as the name of the resource block (see Syntax section above).
<code>plugin</code>	Optional. Indicates that the specified plug-ins are reloaded by Ohai. The default behavior reloads all plug-ins.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Ohai</code>	<code>ohai</code>	The default provider for all platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Reload Ohai

```
ohai "reload" do
 action :reload
end
```

### Reload Ohai after a new user is created

```
ohai "reload_passwd" do
 action :nothing
 plugin "passwd"
end

user "daemonuser" do
 home "/dev/null"
 shell "/sbin/nologin"
 system true
 notifies :reload, "ohai[reload_passwd]", :immediately
end

ruby_block "just an example" do
 block do
 # These variables will now have the new values
 puts node['etc']['passwd']['daemonuser']['uid']
 puts node['etc']['passwd']['daemonuser']['gid']
 end
end
```

## package

The **package** resource is used to manage packages. When the package is installed from a local file (such as with RubyGems, dpkg, or RPM Package Manager), the file must be added to the node using the **remote\_file** or **cookbook\_file** resources.

Note

There are a number of platform-specific resources available for package management. In general, the **package** resource will use the correct package manager based on the platform-specific details collected by Ohai at the start of the chef-client run, which means that the platform-specific resources are often unnecessary. That said, there are cases when using a platform-specific package-based resource is desired. See the following resources for more information about these platform-specific resources: `apt_package`, `chef_gem`, `dpkg_package`, `easy_install_package`, `freebsd_package`, `gem_package`, `ips_package`, `macports_package`, `pacman_package`, `portage_package`, `rpm_package`, `smartos_package`, `solaris_package`, and `yum_package`.

## Syntax

The syntax for using the **package** resource in a recipe is as follows:

```
package "name" do
 some_attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `package` tells the chef-client to use one of sixteen different providers during the chef-client run, where the provider that is used by chef-client depends on the platform of the machine on which the chef-client run is taking place
- `"name"` is the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

## Gem Package Options

The RubyGems package provider attempts to use the RubyGems API to install gems without spawning a new process, whenever possible. A gems command to install will be spawned under the following conditions:

- When a `gem_binary` attribute is specified (as a hash, a string, or by a .gemrc file), the provider will run that command to examine its environment settings and then again to install the gem.
- When install options are specified as a string, the provider will span a gems command with those options when installing the gem.
- The omnibus installer will search the PATH for a gem command rather than defaulting to the current gem environment. As part of `enforce_path_sanity`, the `bin` directories area added to the PATH, which means when there are no other proceeding RubyGems, the installation will still be operated against it.

**Specify Options with a Hash**

If an explicit `gem_binary` parameter is not being used with the `gem_package` resource, it is preferable to provide the install options as a hash. This approach allows the provider to install the gem without needing to spawn an external gem process.

The following RubyGems options are available for inclusion within a hash and are passed to the RubyGems DependencyInstaller:

- `:env_shebang`
- `:force`
- `:format_executable`
- `:ignore_dependencies`
- `:prerelease`
- `:security_policy`
- `:wrappers`

For more information about these options, see the RubyGems documentation: <http://rubygems.rubyforge.org/rubygems-update/Gem/DependencyInstaller.html>.

**Example**

```
gem_package "bundler" do
 options(:prerelease => true, :format_executable => false)
end
```

**Specify Options with a String**

When using an explicit `gem_binary`, options must be passed as a string. When not using an explicit `gem_binary`, the chef-client is forced to spawn a gems process to install the gems (which uses more system resources) when options are passed as a string. String options are passed verbatim to the gems command and should be specified just as if they were passed on a command line. For example, `--prerelease` for a pre-release gem.

**Example**

```
gem_package "nokogiri" do
 gem_binary("/opt/ree/bin/gem")
 options("--prerelease --no-format-executable")
end
```

**Specify Options with a .gemrc File**

Options can be specified in a `.gemrc` file. By default the `gem_package` resource will use the Ruby interface to install gems which will ignore the `.gemrc` file. The `gem_package` resource can be forced to use the gems command instead (and to read the `.gemrc` file) by adding the `gem_binary` attribute to a code block.

**Example**

```
gem_package "nokogiri" do
 gem_binary "gem"
end
```

**Actions**

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:reconfig</code>	Use to reconfigure a package. This action requires a response file.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package. (Debian platform only; for other platforms, use the <code>:remove</code> action.)

**Attributes**

This resource has the following attributes:

Attribute	Description
<code>allow_downgrade</code>	<b>yum_package</b> resource only. Indicates that yum can downgrade a package to satisfy requested version requirements. Default value: <code>false</code> .
<code>arch</code>	The architecture of the package that will be installed or upgraded. (This value can also be passed as part of the package name.) This attribute is only available for the <b>yum_package</b> resource.
<code>flush_cache</code>	<b>yum_package</b> resource only. An array that indicates whether the yum cache should be flushed before or after a yum operation that installs, upgrades, or

	removes a package. Possible values: <code>:before</code> and <code>:after</code> . Default value: <code>{ :before =&gt; false, :after =&gt; false }</code> .
<code>gem_binary</code>	An attribute for the <code>gem_package</code> provider that is used to specify a gems binary. This attribute is useful when installing Ruby 1.9 gems while running in Ruby 1.8.
<code>options</code>	One (or more) additional options that are passed to the command. Can be used with APT, dpkg, Gentoo, RPM Package Manager, and RubyGems.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Apt</code>	<code>apt_package</code>	
<code>Chef::Provider::Package::Dpkg</code>	<code>dpkg_package</code>	Can be used with the <code>options</code> attribute.
<code>Chef::Provider::Package::EasyInstall</code>	<code>easy_install_package</code>	
<code>Chef::Provider::Package::Freebsd</code>	<code>freebsd_package</code>	
<code>Chef::Provider::Package::Ips</code>	<code>ips_package</code>	
<code>Chef::Provider::Package::Macports</code>	<code>macports_package</code>	
<code>Chef::Provider::Package::Pacman</code>	<code>pacman_package</code>	
<code>Chef::Provider::Package::Portage</code>	<code>portage_package</code>	Can be used with the <code>options</code> attribute.
<code>Chef::Provider::Package::Rpm</code>	<code>rpm_package</code>	Can be used with the <code>options</code> attribute.
<code>Chef::Provider::Package::Rubygems</code>	<code>gem_package</code>	Can be used with the <code>options</code> attribute.
<code>Chef::Provider::Package::Rubygems</code>	<code>chef_gem</code>	Can be used with the <code>options</code> attribute.
<code>Chef::Provider::Package::Smartos</code>	<code>smartos_package</code>	
<code>Chef::Provider::Package::Solaris</code>	<code>solaris_package</code>	
<code>Chef::Provider::Package::Yum</code>	<code>yum_package</code>	
<code>Chef::Provider::Package::Zypper</code>	<code>package</code>	The provider that is used with the SuSE platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a gems file for use in recipes

```
chef_gem "right_aws" do
 action :install
end

require 'right_aws'
```

Install a gems file from the local file system

```
gem_package "right_aws" do
 source "/tmp/right_aws-1.11.0.gem"
 action :install
end
```

Install a package

```
package "tar" do
 action :install
end
```

#### Install a package version

```
package "tar" do
 version "1.16.1-1"
 action :install
end
```

#### Install a package with options

```
package "debian-archive-keyring" do
 action :install
 options "--force-yes"
end
```

#### Install a package with a response\_file

Use of a `response_file` is only supported on Debian and Ubuntu at this time. Providers need to be written to support the use of a `response_file`, which contains debconf answers to questions normally asked by the package manager on installation. Put the file in `/files/default` of the cookbook where the package is specified and the chef-client will use the `cookbook_file` resource to retrieve it.

To install a package with a `response_file`:

```
package "sun-java6-jdk" do
 response_file "java.seed"
end
```

#### Install a package using a specific provider

```
package "tar" do
 action :install
 source "/tmp/tar-1.16.1-1.rpm"
 provider Chef::Provider::Package::Rpm
end
```

#### Install a specified architecture using a named provider

```
yum_package "glibc-devel" do
 arch "i386"
end
```

#### Purge a package

```
package "tar" do
 action :purge
end
```

#### Remove a package

```
package "tar" do
 action :remove
end
```

#### Upgrade a package

```
package "tar" do
 action :upgrade
end
```

#### Avoid unnecessary string interpolation

Do this:

```
package "mysql-server" do
 version node['mysql']['version']
 action :install
end
```

and not this:

```
package "mysql-server" do
 version "#{node['mysql']['version']}"
 action :install
end
```

#### Install a package in a platform

The following example shows how to use the `package` resource to install an application named "app" and ensure that the correct packages are installed for the correct platform:

```
package "app_name" do
 action :install
end

case node[:platform]
when "ubuntu", "debian"
 package "app_name-doc" do
 action :install
 end
end
```

```

 action :install
 end
when "centos"
 package "app_name-html" do
 action :install
 end
end
end

```

### Install sudo, then configure /etc/sudoers/ file

The following example shows how to install sudo and then configure the `/etc/sudoers` file:

*# the following code sample comes from the ``default`` recipe in the ``sudo`` cookbook: <https://github.c>*

```

package 'sudo' do
 action :install
end

if node['authorization']['sudo']['include_sudoers_d']
 directory '/etc/sudoers.d' do
 mode '0755'
 owner 'root'
 group 'root'
 action :create
 end

 cookbook_file '/etc/sudoers.d/README' do
 source 'README'
 mode '0440'
 owner 'root'
 group 'root'
 action :create
 end
end

template '/etc/sudoers' do
 source 'sudoers.erb'
 mode '0440'
 owner 'root'
 group platform?('freebsd') ? 'wheel' : 'root'
 variables({
 :sudoers_groups => node['authorization']['sudo']['groups'],
 :sudoers_users => node['authorization']['sudo']['users'],
 :passwordless => node['authorization']['sudo']['passwordless']
 })
end

```

where

- the **package** resource is used to install sudo
- the **if** statement is used to ensure availability of the `/etc/sudoers.d` directory
- the **template** resource tells the chef-client where to find the `sudoers` template
- the **variables** attribute is a hash that passes values to template files (that are located in the `templates/` directory for the cookbook)

### Use a case statement to specify the platform

The following example shows how to use a case statement to tell the chef-client which platforms and packages to install using cURL.

```

package "curl"
case node[:platform]
when "redhat", "centos"
 package "package_1"
 package "package_2"
 package "package_3"
when "ubuntu", "debian"
 package "package_a"
 package "package_b"
 package "package_c"
end
end

```

where `node[:platform]` for each node is identified by Ohai during every chef-client run. For example:

```

package "curl"
case node[:platform]
when "redhat", "centos"
 package "zlib-devel"
 package "openssl-devel"
 package "libc6-dev"
when "ubuntu", "debian"
 package "openssl"
 package "pkg-config"
 package "subversion"
end
end

```

### Use symbols to reference attributes

Do this:

```

package "mysql-server" do
 version node['mysql']['version']
 action :install
end

```

and not this:

```
package "mysql-server" do
 version node[:mysql][:version]
 action :install
end
```

### Use a whitespace array to simplify a recipe

The following examples show different ways of doing the same thing. The first shows a series of packages that will be upgraded:

```
package "package-a" do
 action :upgrade
end

package "package-b" do
 action :upgrade
end

package "package-c" do
 action :upgrade
end

package "package-d" do
 action :upgrade
end
```

and the next uses a single **package** resource and a whitespace array (%w):

```
%w{package-a package-b package-c package-d}.each do |pkg|
 package pkg do
 action :upgrade
 end
end
```

where `|pkg|` is used to define the name of the resource, but also to ensure that each item in the whitespace array has its own name.

## pacman\_package

The **pacman\_package** resource is used to manage packages (using pacman) on the Arch Linux platform.

Note
In many cases, it is better to use the <b>package</b> resource instead of this one. This is because when the <b>package</b> resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the <b>package</b> resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

### Syntax

The syntax for using the **pacman\_package** resource in a recipe is as follows:

```
pacman_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- **pacman\_package** tells the chef-client to use the `Chef::Provider::Pacman` provider during the chef-client run
- **name** is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, **name** is also the name of the package
- **attribute** is zero (or more) of the attributes that are available for this resource
- **:action** is the step that the resource will ask the provider to take during the chef-client run

### Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

### Attributes

This resource has the following attributes:



Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Pacman</code>	<code>pacman_package</code>	The provider that is used with the Arch Linux platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package

```
pacman_package "name of package" do
 action :install
end
```

perl

The `perl` resource is used to execute scripts using the Perl interpreter and includes all of the actions and attributes that are available to the `execute` resource.

Note

The `perl` script resource (which is based on the `script` resource) is different from the `ruby_block` resource because Ruby code that is run with this resource is created as a temporary file and executed like other script resources, rather than run inline. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the `not_if` and `only_if` meta parameters to guard the use of this resource for idempotence.

Syntax

The syntax for using the `perl` resource in a recipe is as follows:

```
perl "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `perl` tells the chef-client to use the `Chef::Resource::Script::Perl` provider during the chef-client run
- `name` is the name of the resource block; when the `command` attribute is not specified as part of a recipe, `name` is also the name of the command to be executed
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:run</code>	Default. Use to run a script.
<code>:nothing</code>	Indicates that the command should not be run. This action is used to specify that a command is run only when another resource notifies it.

Attributes

This resource has the following attributes:

Attribute	Description
<code>code</code>	A quoted ( " ") string of code to be executed.
<code>command</code>	The name of the command to be executed. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>creates</code>	Indicates that a command to create a file will not be run when that file already exists.
<code>cwd</code>	The current working directory.
<code>environment</code>	A Hash of environment variables in the form of <code>{ "ENV_VARIABLE" =&gt; "VALUE" }</code> . (These variables must exist for a command to be run successfully.)
<code>flags</code>	One (or more) command line flags that are passed to the interpreter when a command is invoked.
<code>group</code>	The group name or group ID that must be changed before running a command.
<code>path</code>	An array of paths to use when searching for a command. These paths are not added to the command's environment \$PATH. The default value uses the system path.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>returns</code>	The return value for a command. This may be an array of accepted values. An exception is raised when the return value(s) do not match. Default value: <code>0</code> .
<code>timeout</code>	The amount of time (in seconds) a command will wait before timing out. Default value: <code>3600</code> .
<code>user</code>	The user name or user ID that should be changed before running a command.
<code>umask</code>	The file mode creation mask, or umask.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Script</code>	<code>script</code>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<code>Chef::Provider::Script::Perl</code>	<code>perl</code>	The provider that is used with the Perl command interpreter.

## Examples

None.

## portage\_package

The **portage\_package** resource is used to manage packages for the Gentoo platform.

### Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

## Syntax

The syntax for using the **portage\_package** resource in a recipe is as follows:

```
portage_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `portage_package` tells the chef-client to use the `Chef::Provider::Portage` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

Attributes

This resource has the following attributes:

Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Portage</code>	<code>portage_package</code>	The provider that is used with the Gentoo platform. Can be used with the <code>options</code> attribute.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package

```
portage_package "name of package" do
 action :install
end
```

powershell\_script

The **powershell\_script** resource is a resource for the Microsoft Windows platform that is used to execute a script using the Windows PowerShell interpreter, much like how the **script** and **script**-based resources—**bash**, **csh**, **perl**, **python**, and **ruby**—are used. The **powershell\_script** is specific to the Microsoft Windows platform and the Windows PowerShell interpreter. This resource creates and executes a temporary file (similar to how the **script** resource behaves), rather than running the command inline. This resource includes actions (`:run` and `:nothing`) and attributes (`creates`, `cwd`, `environment`, `group`, `path`, `timeout`, and `user`) that are inherited from the **execute** resource. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the `not_if` and `only_if` meta parameters to guard the use of this resource for idempotence.

Syntax

The syntax for using the **powershell\_script** resource in a recipe is as follows:

```
powershell_script "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `powershell_script` tells the chef-client to use the `Chef::Provider::PowershellScript` provider during the chef-client run

- `name` is the name of the resource block; when the `command` attribute is not specified as part of a recipe, `name` is also the command to be executed
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the `powershell_script` resource can work when used in a recipe:

```
powershell_script "name_of_script" do
 cwd Chef::Config[:file_cache_path]
 code <<-EOH
 # some script goes here
EOH
end
```

Actions

This resource has the following actions:

Action	Description
<code>:run</code>	Default. Use to run the script.

Attributes

This resource has the following attributes:

Attribute	Description
<code>architecture</code>	The architecture of the process under which a script is executed. Possible values: <code>:x86</code> (for 32-bit processes) and <code>:x86_64</code> (for 64-bit processes). If these values are not provided in a recipe, the chef-client will default to the correct value for the architecture, as determined by Ohai. An exception will be raised when anything other than <code>:x86</code> is specified for a 32-bit process.
<code>code</code>	A quoted ( " ") string of code to be executed.
<code>command</code>	The name of the command to be executed. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>flags</code>	One (or more) command line flags that are passed to the interpreter when a command is invoked. Default value: [ <code>-NoLogo</code> , <code>-NonInteractive</code> , <code>-NoProfile</code> , <code>-ExecutionPolicy RemoteSigned</code> , <code>-InputFormat None</code> , <code>-File</code> ].
<code>interpreter</code>	The script interpreter to be used during code execution.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::PowershellScript</code>	<code>powershell_script</code>	The default provider for all platforms.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Write to an interpolated path

```
powershell_script "write-to-interpolated-path" do
 code <<-EOH
 $stream = [System.IO.StreamWriter] "#{Chef::Config[:file_cache_path]}/powershell-test.txt"
 $stream.WriteLine("In #{Chef::Config[:file_cache_path]}...word.")
 $stream.close()
EOH
end
```

Change the working directory

```
powershell_script "cwd-then-write" do
 cwd Chef::Config[:file_cache_path]
 code <<-EOH
 $stream = [System.IO.StreamWriter] "C:/powershell-test2.txt"
 $pwd = pwd
 $stream.WriteLine("This is the contents of: $pwd")
 $dirs = dir
 foreach ($dir in $dirs) {
 $stream.WriteLine($dir.fullname)
 }
 $stream.close()
EOH
end
```

Change the working directory in Microsoft Windows

```
powershell_script "cwd-to-win-env-var" do
 cwd "%TEMP%"
 code <<-EOH
 $stream = [System.IO.StreamWriter] "./temp-write-from-chef.txt"
 $stream.WriteLine("chef on windows rox yo!")
 $stream.close()
EOH
end
```

Pass an environment variable to a script

```
powershell_script "read-env-var" do
 cwd Chef::Config[:file_cache_path]
 environment ({'foo' => 'BAZ'})
 code <<-EOH
 $stream = [System.IO.StreamWriter] "./test-read-env-var.txt"
 $stream.WriteLine("FOO is $foo")
 $stream.close()
EOH
end
```

python

The **python** resource is used to execute scripts using the Python interpreter and includes all of the actions and attributes that are available to the **execute** resource.

Note

The **python** script resource (which is based on the **script** resource) is different from the **ruby\_block** resource because Ruby code that is run with this resource is created as a temporary file and executed like other script resources, rather than run inline. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the `not_if` and `only_if` meta parameters to guard the use of this resource for idempotence.

Syntax

The syntax for using the **python** resource in a recipe is as follows:

```
python "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `python` tells the chef-client to use the `Chef::Resource::Script::Python` provider during the chef-client run
- `name` is the name of the resource block; when the `command` attribute is not specified as part of a recipe, `name` is also the name of the command to be executed
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:run</code>	Default. Use to run a script.
<code>:nothing</code>	Indicates that the command should not be run. This action is used to specify that a command is run only when another resource notifies it.

Attributes

This resource has the following attributes:

Attribute	Description
<code>code</code>	A quoted ( " ") string of code to be executed.
<code>command</code>	The name of the command to be executed. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>creates</code>	Indicates that a command to create a file will not be run when that file already exists.
<code>cwd</code>	The current working directory.
<code>environment</code>	A Hash of environment variables in the form of { "ENV_VARIABLE" => "VALUE" }. (These variables must exist for a command to be run successfully.)
<code>flags</code>	One (or more) command line flags that are passed to the interpreter when a command is invoked.

<code>group</code>	The group name or group ID that must be changed before running a command.
<code>path</code>	An array of paths to use when searching for a command. These paths are not added to the command's environment \$PATH. The default value uses the system path.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>returns</code>	The return value for a command. This may be an array of accepted values. An exception is raised when the return value(s) do not match. Default value: <code>0</code> .
<code>timeout</code>	The amount of time (in seconds) a command will wait before timing out. Default value: <code>3600</code> .
<code>user</code>	The user name or user ID that should be changed before running a command.
<code>umask</code>	The file mode creation mask, or umask.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Script</code>	<code>script</code>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<code>Chef::Provider::Script::Python</code>	<code>python</code>	The provider that is used with the Python command interpreter.

Examples

None.

registry\_key

The `registry_key` resource is used to create and delete registry keys in Microsoft Windows.

64-bit versions of Microsoft Windows have a 32-bit compatibility layer in the registry that reflects and redirects certain keys (and their sub-keys) into specific locations. By default, the registry functionality will default to the machine architecture of the system that is being configured. The chef-client can access any reflected or redirected registry key. The chef-client can write to any 64-bit registry location. (This behavior is not affected by the chef-client running as a 32-bit application.) For more information, see: [http://msdn.microsoft.com/en-us/library/windows/desktop/aa384235\(v=vs.85\).aspx](http://msdn.microsoft.com/en-us/library/windows/desktop/aa384235(v=vs.85).aspx).

Syntax

The syntax for using the `registry_key` resource in a recipe is as follows:

```
registry_key "name" do
 attribute "value" # see attributes section below
 ...
 values [{
 :name => "name",
 :type => :string,
 :data => "data"
 },
 {
 :name => "name",
 :type => :string,
 :data => "data"
 },
 ...
]
 action :action # see actions section below
end
```

where

- `registry_key` tells the chef-client to use the `Chef::Provider::Windows::Registry` provider during the chef-client run
- `name` is the name of the resource block; when the `key` attribute is not specified as part of a recipe, `name` is also path to the location in which a registry key is created or from which a registry key is deleted
- `attribute` is zero (or more) of the attributes that are available for this resource
- `values` is a hash that contains at least one registry key to be created or deleted. Each registry key in the hash is grouped by brackets in which the `:name`, `:type`, and `:data` values for that registry key are specified.
- `:type` represents the values available for registry keys in Microsoft Windows. Use `:binary` for REG\_BINARY, `:string` for REG\_SZ, `:multi_string` for REG\_MULTI\_SZ, `:expand_string` for REG\_EXPAND\_SZ, `:dword` for REG\_DWORD, `:dword_big_endian` for REG\_DWORD\_BIG\_ENDIAN, or `:qword` for REG\_QWORD.
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the `registry_key` resource can work when used in a recipe. In this example, a Microsoft Windows registry key named "System" will get a new value called "NewRegistryKeyValue" and a multi-string value named "foo bar":

```
registry_key "HKEY_LOCAL_MACHINE\\...\\System" do
 values [{
 :name => "NewRegistryKeyValue",
```

```
 .type => :multi_string,
 :data => ['foo\0bar\0\0']
 }
 end
 action :create
end
```

and the following example shows how multiple registry key entries can be configured using a single resource block with key values based on node attributes:

```
registry_key 'HKEY_LOCAL_MACHINE\SOFTWARE\Microsoft\Windows\CurrentVersion\name_of_registry_key' do
 values [{:name => 'key_name', :type => :string, :data => 'C:\Windows\System32\file_name.bmp'},
 {:name => 'key_name', :type => :string, :data => node['node_name']['attribute']['value']},
 {:name => 'key_name', :type => :string, :data => node['node_name']['attribute']['value']}
]
 action :create
end
```

Registry Key Path Separators

A Microsoft Windows registry key can be used as a string in Ruby code, such as when a registry key is used as the name of a recipe. In Ruby, when a registry key is enclosed in a double-quoted string (" "), the same backslash character (\) that is used to define the registry key path separator is also used in Ruby to define an escape character. Therefore, the registry key path separators must be escaped. For example, the following registry key:

```
HKCU\SOFTWARE\Policies\Microsoft\Windows\CurrentVersion\Themes
```

will not work when it is defined like this:

```
registry_key "HKCU\SOFTWARE\Policies\Microsoft\Windows\CurrentVersion\Themes" do
 ...
 action :some_action
end
```

but will work when the path separators are escaped properly:

```
registry_key "HKCU\\SOFTWARE\\Policies\\Microsoft\\Windows\\CurrentVersion\\Themes" do
 ...
 action :some_action
end
```

Actions

This resource has the following actions:

Action	Description
<a href="#">:create</a>	Default. Use to create a registry key.
<a href="#">:create_if_missing</a>	Use to create a registry key if it does not exist. Also, use to create a registry key value if it does not exist.
<a href="#">:delete</a>	Use to delete the specified values for a registry key.
<a href="#">:delete_key</a>	Use to delete the specified registry key and all of its subkeys.

Note

Be careful when using the `:delete_key` action with the `recursive` attribute. This will delete the registry key, all of its subkeys and all of the values associated with them. This cannot be undone by the chef-client.

Attributes

This resource has the following attributes:

Attribute	Description
<a href="#">architecture</a>	<p>The architecture of the node for which keys will be created or deleted. Possible values: <code>:i386</code> (for nodes with a 32-bit registry), <code>:x86_64</code> (for nodes with a 64-bit registry), and <code>:machine</code> (to have the chef-client determine the architecture during the chef-client run). Default value: <code>:machine</code>.</p> <p>In order to read or write 32-bit registry keys on 64-bit machines running Microsoft Windows, the <code>architecture</code> attribute must be set to <code>:i386</code>. The <code>:x86_64</code> value can be used to force writing to a 64-bit registry location, but this value is less useful than the default (<code>:machine</code>) because the chef-client will return an exception if <code>:x86_64</code> is used and the machine turns out to be a 32-bit machine (whereas with <code>:machine</code>, the chef-client will be able to access the registry key on the 32-bit machine).</p>

Note

The ARCHITECTURE attribute should only specify `:x86_64` or `:i386` when it is necessary to write 32-bit (`:i386`) or 64-bit (`:x86_64`) values on a 64-bit machine. ARCHITECTURE will default to `:machine` unless a specific value is given.

<a href="#">key</a>	The path to the location in which a registry key will be created or from which a registry key will be
---------------------	-------------------------------------------------------------------------------------------------------

deleted. Default value: the name of the resource block (see Syntax section above).

The path must include the registry hive, which can be specified either as its full name or as the 3- or 4-letter abbreviation. For example, both `HKLM\SECURITY` and `HKEY_LOCAL_MACHINE\SECURITY` are both valid and equivalent. The following hives are valid: `HKEY_LOCAL_MACHINE`, `HKLM`, `HKEY_CURRENT_CONFIG`, `HKCC`, `HKEY_CLASSES_ROOT`, `HKCR`, `HKEY_USERS`, `HKU`, `HKEY_CURRENT_USER`, and `HKCU`.

`provider` Optional. Use to specify a provider by using its long name. For example: `provider Chef::Provider::Long::Name`. See the Providers section below for the list of providers available to this resource.

`recursive` When creating a key, this value indicates whether the required keys for the specified path will be created. When using the `:delete_key` action in a recipe, and if the registry key has subkeys, then the value for this attribute should be set to `true`.

Note

Be careful when using the `:delete_key` action with the `recursive` attribute. This will delete the registry key, all of its subkeys and all of the values associated with them. This cannot be undone by the chef-client.

`values` An array of hashes, where each Hash contains the values that will be set under a registry key. Each Hash must contain `:name`, `:type`, and `:data` (and must contain no other key values).

`:type` represents the values available for registry keys in Microsoft Windows. Use `:binary` for `REG_BINARY`, `:string` for `REG_SZ`, `:multi_string` for `REG_MULTI_SZ`, `:expand_string` for `REG_EXPAND_SZ`, `:dword` for `REG_DWORD`, `:dword_big_endian` for `REG_DWORD_BIG_ENDIAN`, or `:qword` for `REG_QWORD`.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Windows::Registry</code>	<code>registry_key</code>	The default provider for the Microsoft Windows platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Create a registry key

```
registry_key "HKEY_LOCAL_MACHINE\\SOFTWARE\\Microsoft\\Windows\\CurrentVersion\\Policies\\System" do
 values [{
 :name => "EnableLUA",
 :type => :dword,
 :data => 0
 }]
 action :create
end
```

Delete a registry key value

```
registry_key "HKEY_LOCAL_MACHINE\\SOFTWARE\\Policies\\Microsoft\\Windows\\WindowsUpdate\\AU" do
 values [{
 :name => "NoAutoRebootWithLoggedOnUsers",
 :type => :dword
 }]
 action :delete
end
```

Delete a registry key and its subkeys, recursively

```
registry_key "HKCU\\SOFTWARE\\Policies\\Microsoft\\Windows\\CurrentVersion\\Themes" do
 recursive true
 action :delete_key
end
```

Note

Be careful when using the `:delete_key` action with the `recursive` attribute. This will delete the registry key, all of its subkeys and all of the values associated with them. This cannot be undone by the chef-client.

Use re-directed keys

In 64-bit versions of Microsoft Windows, `HKEY_LOCAL_MACHINE\SOFTWARE\Example` is a re-directed key. In the following examples, because `HKEY_LOCAL_MACHINE\SOFTWARE\Example` is a 32-bit key, the output will be "Found 32-bit key" if they are run on a version of Microsoft Windows that is 64-bit:

```
registry_key "HKEY_LOCAL_MACHINE\\SOFTWARE\\Example" do
```



```

architecture :i386
recursive true
action :create
end

```

or:

```

registry_key "HKEY_LOCAL_MACHINE\\SOFTWARE\\Example" do
 architecture :x86_64
 recursive true
 action :delete_key
end

```

or:

```

ruby_block "check 32-bit" do
 block do
 puts "Found 32-bit key"
 end
 only_if { registry_key_exists?("HKEY_LOCAL_MACHINE\\SOFTWARE\\Example", :i386) }
end

```

or:

```

ruby_block "check 64-bit" do
 block do
 puts "Found 64-bit key"
 end
 only_if { registry_key_exists?("HKEY_LOCAL_MACHINE\\SOFTWARE\\Example", :x86_64) }
end

```

**Set proxy settings to be the same as those used by the chef-client**

```

proxy = URI.parse(Chef::Config[:http_proxy])
registry_key 'HKCU\\Software\\Microsoft\\Windows\\CurrentVersion\\Internet Settings' do
 values [{:name => 'ProxyEnable', :type => :reg_dword, :data => 1},
 {:name => 'ProxyServer', :data => "#{proxy.host}:#{proxy.port}"},
 {:name => 'ProxyOverride', :type => :reg_string, :data => <local>},
]
 action :create
end

```

## remote\_directory

The **remote\_directory** resource is used incrementally transfer a directory from a cookbook to a node. The directory that is copied from the cookbook should be located under `COOKBOOK_NAME/files/default/REMOTE_DIRECTORY`. The **remote\_directory** resource will obey file specificity.

### Syntax

The syntax for using the **remote\_directory** resource in a recipe is as follows:

```

remote_directory "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end

```

where

- `remote_directory` tells the chef-client to use the `Chef::Provider::Directory::RemoteDirectory` provider during the chef-client run
- `name` is the path to the location below which the chef-client will manage directories
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

### Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a directory and/or the contents of that directory.
<code>:create_if_missing</code>	Use to create a directory and/or the contents of that directory, but only if it does not exist.
<code>:delete</code>	Use to delete a directory, including the contents of that directory.

### Attributes

This resource has the following attributes:

Attribute	Description
<code>cookbook</code>	The cookbook in which a file is located (if it is not located in the current cookbook). The default value is the current cookbook.

<code>files_backup</code>	The number of backup copies to keep for files in the directory. Default value: 5.
<code>files_group</code>	Use to configure group permissions for files. A string or ID that identifies the group owner by group name, including fully qualified group names such as <code>domain\group</code> or <code>group@domain</code> . If this value is not specified, existing groups will remain unchanged and new group assignments will use the default <code>POSIX</code> group (if available).
<code>files_mode</code>	<p>The octal mode for a file.</p> <p>UNIX- and Linux-based systems: The octal mode that is passed to <code>chmod</code>. If the value is specified as a quoted string, it will work exactly as if the <code>chmod</code> command was passed. If the value is specified as an integer, prepend a zero (0) to the value to ensure it is interpreted as an octal number. For example, to assign read, write, and execute rights for all users, use <code>0777</code> or <code>'777'</code>; for the same rights, plus the sticky bit, use <code>01777</code> or <code>'1777'</code>.</p> <p>Microsoft Windows: The octal mode that is translated into rights for Microsoft Windows security. Values up to <code>0777</code> are allowed (no sticky bits) and mean the same in Microsoft Windows as they do in UNIX, where <code>4</code> equals <code>GENERIC_READ</code>, <code>2</code> equals <code>GENERIC_WRITE</code>, and <code>1</code> equals <code>GENERIC_EXECUTE</code>. This attribute cannot be used to set <code>:full_control</code>. This attribute has no effect if not specified, but when this attribute and <code>rights</code> are both specified, the effects will be cumulative.</p>
<code>files_owner</code>	Use to configure owner permissions for files. A string or ID that identifies the group owner by user name, including fully qualified user names such as <code>domain\user</code> or <code>user@domain</code> . If this value is not specified, existing owners will remain unchanged and new owner assignments will use the current user (when necessary).
<code>group</code>	Use to configure permissions for directories. A string or ID that identifies the group owner by group name, including fully qualified group names such as <code>domain\group</code> or <code>group@domain</code> . If this value is not specified, existing groups will remain unchanged and new group assignments will use the default <code>POSIX</code> group (if available).
<code>inherits</code>	Microsoft Windows only. Indicates that a file inherits rights from its parent. Default value: <code>true</code> .
<code>mode</code>	<p>The octal mode for a directory. If <code>mode</code> is not specified and if the directory already exists, the existing mode on the directory is used. If <code>mode</code> is not specified, the directory does not exist, and the <code>:create</code> action is specified, the chef-client will assume a mask value of <code>0777</code> and then apply the umask for the system on which the directory will be created to the mask value. For example, if the umask on a system is <code>022</code>, the chef-client would use the default value of <code>0755</code>.</p> <p>The behavior is different depending on the platform.</p> <p>UNIX- and Linux-based systems: The octal mode that is passed to <code>chmod</code>. If the value is specified as a quoted string, it will work exactly as if the <code>chmod</code> command was passed. If the value is specified as an integer, prepend a zero (0) to the value to ensure it is interpreted as an octal number. For example, to assign read, write, and execute rights for all users, use <code>0777</code> or <code>'777'</code>; for the same rights, plus the sticky bit, use <code>01777</code> or <code>'1777'</code>.</p> <p>Microsoft Windows: The octal mode that is translated into rights for Microsoft Windows security. Values up to <code>0777</code> are allowed (no sticky bits) and mean the same in Microsoft Windows as they do in UNIX, where <code>4</code> equals <code>GENERIC_READ</code>, <code>2</code> equals <code>GENERIC_WRITE</code>, and <code>1</code> equals <code>GENERIC_EXECUTE</code>. This attribute cannot be used to set <code>:full_control</code>. This attribute has no effect if not specified, but when this attribute and <code>rights</code> are both specified, the effects will be cumulative.</p>
<code>overwrite</code>	Indicates that a file (if different) will be overwritten. Default value: <code>true</code> .
<code>owner</code>	Use to configure permissions for directories. A string or ID that identifies the group owner by user name, including fully qualified user names such as <code>domain\user</code> or <code>user@domain</code> . If this value is not specified, existing owners will remain unchanged and new owner assignments will use the current user (when necessary).
<code>path</code>	The path to the directory. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>purge</code>	Indicates that extra files found in the target directory will be purged. Default value: <code>false</code> .
<code>recursive</code>	Use to create or delete directories recursively. Default value: <code>true</code> ; the chef-client must be able to create the directory structure, including parent directories (if missing), as defined in <code>COOKBOOK_NAME/files/default/REMOTE_DIRECTORY</code> .
<code>rights</code>	Microsoft Windows only. The permissions for users and groups in a Microsoft Windows environment. For example: <code>rights &lt;permissions&gt;, &lt;principal&gt;, &lt;options&gt;</code> where <code>&lt;permissions&gt;</code> specifies the rights granted to the principal, <code>&lt;principal&gt;</code> is the group or user name, and <code>&lt;options&gt;</code> is a Hash with one (or more) advanced rights options.

sourceThe base name of the source file (and inferred from the path attribute).**Recursive Directories**

The **remote\_directory** resource can be used to create directory structures on a target node based on a directory structure that is defined in a cookbook. When the required directory structure does not exist, the **remote\_directory** resource will create that structure explicitly.

A directory structure:

```
/foo
 /bar
 /baz
```

The following example shows a way create a file in the /baz directory:

```
remote_directory "/foo/bar/baz" do
 owner "root"
 group "root"
 mode 0755
 action :create
end
```

With this example, the group, mode, and owner attribute values will be applied to /baz. If the directory structure were:

```
/foo
```

the **remote\_directory** resource would first create the required directory structure:

```
/foo
 /bar
 /baz
```

and apply the group, mode, and owner attribute values to the entire directory structure.

**Providers**

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<u>Chef::Provider::Directory::RemoteDirectory</u>	<u>remote_directory</u>	The default provider for all platforms.

**Examples**

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

**Recursively transfer a directory from a remote location**

```
create up to 10 backups of the files, set the files owner different from the directory.
remote_directory "/tmp/remote_something" do
 source "something"
 files_backup 10
 files_owner "root"
 files_group "root"
 files_mode 00644
 owner "nobody"
 group "nobody"
 mode 00755
end
```

**Use with the chef\_handler lightweight resource**

The following example shows how to use the **remote\_directory** resource and the **chef\_handler** lightweight resource to reboot a handler named WindowsRebootHandler:

```
the following code sample comes from the ``reboot_handler`` recipe in the ``windows`` cookbook: https:
remote_directory node['chef_handler']['handler_path'] do
 source 'handlers'
 recursive true
 action :create
end

chef_handler 'WindowsRebootHandler' do
 source "#{node['chef_handler']['handler_path']}/windows_reboot_handler.rb"
 arguments node['windows']['allow_pending_reboots']
 supports :report => true, :exception => false
 action :enable
end
```

**remote\_file**

The **remote\_file** resource is used to transfer a file from a remote location using file specificity. This resource is similar to the **file** resource.

Note

Fetching files from the `files/` directory in a cookbook should be done with the `cookbook_file` resource.

Syntax

The syntax for using the `remote_file` resource in a recipe is as follows:

```
remote_file "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `remote_file` tells the chef-client to use the `Chef::Provider::File::RemoteFile` provider during the chef-client run
- `name` is the name of the resource block; when the `path` attribute is not specified as part of a recipe, `name` is also the path to the remote file
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example of how the `remote_file` resource can work when used in a recipe:

```
remote_file "#{Chef::Config[:file_cache_path]}/large-file.tar.gz" do
 source "http://www.example.org/large-file.tar.gz"
end
```

Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a file.
<code>:create_if_missing</code>	Use to create a file only if the file does not exist. (When the file exists, nothing happens.)
<code>:delete</code>	Use to delete a file.
<code>:touch</code>	Use to touch a file. This updates the access ( <code>atime</code> ) and file modification ( <code>mtime</code> ) times for a file. (This action may be used with this resource, but is typically only used with the <code>file</code> resource.)

Attributes

This resource has the following attributes:

Attribute	Description
<code>atomic_update</code>	Indicates whether atomic file updates are used on a per-resource basis. Set to <code>true</code> for atomic file updates. Set to <code>false</code> for non-atomic file updates. (This setting overrides <code>file_atomic_update</code> , which is a global setting found in the <code>client.rb</code> file.) Default value: <code>true</code> .
<code>backup</code>	The number of backups to be kept. Set to <code>false</code> to prevent backups from being kept. Default value: <code>5</code> .
<code>checksum</code>	Optional, see <code>use_conditional_get</code> . The SHA-256 checksum of the file. Use to prevent the <code>remote_file</code> resource from re-downloading a file. When the local file matches the checksum, the chef-client will not download it.
<code>force_unlink</code>	Use to specify how the chef-client handles certain situations when the target file turns out not to be a file. For example, when a target file is actually a symlink. Set to <code>true</code> to have the chef-client delete the non-file target and replace it with the specified file. Set to <code>false</code> for the chef-client to raise an error. Default value: <code>false</code> .
<code>ftp_active_mode</code>	Indicates whether the chef-client will use active or passive FTP. Set to <code>true</code> to use active FTP. Default value: <code>false</code> .
<code>group</code>	A string or ID that identifies the group owner by group name, including fully qualified group names such as <code>domain\group</code> or <code>group@domain</code> . If this value is not specified, existing groups will remain unchanged and new group assignments will use the default <code>POSIX</code> group (if available).
<code>headers</code>	A Hash of custom headers. Default value: <code>{}</code> .
<code>inherits</code>	Microsoft Windows only. Indicates that a file inherits rights from its parent. Default value: <code>true</code> .
<code>manage_symlink_source</code>	Indicates that the chef-client will detect and manage the source file for a symlink. Possible values: <code>nil</code> , <code>true</code> , or <code>false</code> . When this value is set to <code>nil</code> , the chef-client will manage a symlink's source file and emit a warning. When this value is set to <code>true</code> , the chef-client will manage a symlink's source file and not emit a warning. Default value: <code>nil</code> . The default value will be changed to <code>false</code> in a future version.
<code>mode</code>	The octal mode for a file. If <code>mode</code> is not specified and if the file already exists, the existing mode on the file is used. If <code>mode</code> is not specified, the file does not exist, and the <code>:create</code> action is specified, the chef-client will assume a mask value of <code>0777</code> and then apply the <code>umask</code> for the system on which the file will be created to the <code>mask</code> value. For example, if the <code>umask</code> on a system is <code>022</code> , the chef-client would use the default value of <code>0755</code> .  The behavior is different depending on the platform.

UNIX- and Linux-based systems: The octal mode that is passed to `chmod`. If the value is specified as a quoted string, it will work exactly as if the `chmod` command was passed. If the value is specified as an integer, prepend a zero (0) to the value to ensure it is interpreted as an octal number. For example, to assign read, write, and execute rights for all users, use `0777` or `'777'`; for the same rights, plus the sticky bit, use `01777` or `'1777'`.

Microsoft Windows: The octal mode that is translated into rights for Microsoft Windows security. Values up to `0777` are allowed (no sticky bits) and mean the same in Microsoft Windows as they do in UNIX, where `4` equals `GENERIC_READ`, `2` equals `GENERIC_WRITE`, and `1` equals `GENERIC_EXECUTE`. This attribute cannot be used to set `:full_control`. This attribute has no effect if not specified, but when this attribute and `rights` are both specified, the effects will be cumulative.

<code>owner</code>	A string or ID that identifies the group owner by user name, including fully qualified user names such as <code>domain\user</code> or <code>user@domain</code> . If this value is not specified, existing owners will remain unchanged and new owner assignments will use the current user (when necessary).
<code>path</code>	The path to the file. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>rights</code>	Microsoft Windows only. The permissions for users and groups in a Microsoft Windows environment. For example: <code>rights &lt;permissions&gt;, &lt;principal&gt;, &lt;options&gt;</code> where <code>&lt;permissions&gt;</code> specifies the rights granted to the principal, <code>&lt;principal&gt;</code> is the group or user name, and <code>&lt;options&gt;</code> is a Hash with one (or more) advanced rights options.
<code>source</code>	Required. The location (URI) of the source file. This value may also specify HTTP ( <code>http://</code> ), FTP ( <code>ftp://</code> ), or local ( <code>file://</code> ) source file locations. There are many ways to define the location of a source file. By using a path:

```
source "http://couchdb.apache.org/img/sketch.png"
```

By using a node attribute:

```
source node['nginx']['foo123']['url']
```

By using attributes to define paths:

```
source "#{node['python']['url']}/#{version}/Python-#{version}.tar.bz2"
```

By defining multiple paths for multiple locations:

```
source "http://seapower/spring.png", "http://seapower/has.png", "http://seapower/sprung.png"
```

By defining those same multiple paths as an array:

```
source ["http://seapower/spring.png", "http://seapower/has.png", "http://seapower/sprung.png"]
```

When multiple paths are specified, the chef-client will attempt to download the files in the order listed, stopping after the first successful download.

<code>use_conditional_get</code>	Use to enable conditional HTTP requests by using a conditional GET (with the If-Modified-Since header) or an opaque identifier (ETag). To use If-Modified-Since headers, <code>use_last_modified</code> must also be set to <code>true</code> . To use ETag headers, <code>use_etag</code> must also be set to <code>true</code> . Default value: <code>true</code> .
<code>use_etag</code>	Indicates that ETag headers are enabled. Set to <code>false</code> to disable ETag headers. To use this setting, <code>use_conditional_get</code> must also be set to <code>true</code> . Default value: <code>true</code> .
<code>use_last_modified</code>	Indicates that If-Modified-Since headers are enabled. Set to <code>false</code> to disable If-Modified-Since headers. To use this setting <code>use_conditional_get</code> must also be set to <code>true</code> . Default value: <code>true</code> .

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::File::RemoteFile</code>	<code>remote_file</code>	The default provider for all platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Transfer a file from a URL

```
remote_file "/tmp/testfile" do
 source "http://www.example.com/tempfiles/testfile"
 mode 00644
 checksum "3a7dac00b1" # A SHA256 (or portion thereof) of the file.
end
```

### Transfer a file only when the source has changed

```

remote_file "/tmp/couch.png" do
 source "http://couchdb.apache.org/img/sketch.png"
 action :nothing
end

http_request "HEAD http://couchdb.apache.org/img/sketch.png" do
 message ""
 url "http://couchdb.apache.org/img/sketch.png"
 action :head
 if File.exists?("/tmp/couch.png")
 headers "If-Modified-Since" => File.mtime("/tmp/couch.png").httpdate
 end
 notifies :create, "remote_file[/tmp/couch.png]", :immediately
end

```

#### Install a file from a remote location using bash

The following is an example of how to install the `foo123` module for Nginx. This module adds shell-style functionality to an Nginx configuration file and does the following:

- Declares three variables
- Gets the Nginx file from a remote location
- Installs the file using Bash to the path specified by the `src_filepath` variable

*# the following code sample is similar to the `upload_progress_module` recipe in the `nginx` cookbook.*

```

src_filename = "foo123-nginx-module-v#{node['nginx']['foo123']['version']}.tar.gz"
src_filepath = "#{Chef::Config['file_cache_path']}/#{src_filename}"
extract_path = "#{Chef::Config['file_cache_path']/nginx_foo123_module}/#{node['nginx']['foo123']['checksum']}"

remote_file src_filepath do
 source node['nginx']['foo123']['url']
 checksum node['nginx']['foo123']['checksum']
 owner 'root'
 group 'root'
 mode 00644
end

bash 'extract_module' do
 cwd ::File.dirname(src_filepath)
 code <<-EOH
 mkdir -p #{extract_path}
 tar xzf #{src_filename} -C #{extract_path}
 mv #{extract_path}/*/* #{extract_path}/
 EOH
 not_if { ::File.exists?(extract_path) }
end

```

#### Store certain settings

The following recipe shows how an attributes file can be used to store certain settings. An attributes file is located in the `attributes/` directory in the same cookbook as the recipe which calls the attributes file. In this example, the attributes file specifies certain settings for Python that are then used across all nodes against which this recipe will run.

Python packages have versions, installation directories, URLs, and checksum files. An attributes file that exists to support this type of recipe would include settings like the following:

```

default['python']['version'] = '2.7.1'

if python['install_method'] == 'package'
 default['python']['prefix_dir'] = '/usr'
else
 default['python']['prefix_dir'] = '/usr/local'
end

default['python']['url'] = 'http://www.python.org/ftp/python'
default['python']['checksum'] = '80e387...85fd61'

```

and then the methods in the recipe may refer to these values. A recipe that is used to install Python will need to do the following:

- Identify each package to be installed (implied in this example, not shown)
- Define variables for the package version and the `install_path`
- Get the package from a remote location, but only if the package does not already exist on the target system
- Use the **bash** resource to install the package on the node, but only when the package is not already installed

*# the following code sample comes from the `oc-nginx` cookbook on [github](https://github.com/cookbooks/oc-nginx): <https://github.com/cookbooks/oc-nginx>*

```

version = node['python']['version']
install_path = "#{node['python']['prefix_dir']/lib/python#{version.split(/(\^d+\.\d+)/)[1]}}"

remote_file "#{Chef::Config[:file_cache_path]/Python-#{version}.tar.bz2" do
 source "#{node['python']['url']}/#{version}/Python-#{version}.tar.bz2"
 checksum node['python']['checksum']
 mode "0644"
 not_if { ::File.exists?(install_path) }
end

bash "build-and-install-python" do
 cwd Chef::Config[:file_cache_path]
 code <<-EOF
 tar -jxvf Python-#{version}.tar.bz2
 (cd Python-#{version} && ./configure #{configure_options})
 (cd Python-#{version} && make && make install)
 EOF
 not_if { ::File.exists?(install_path) }
end

```

```
end
```

### Use the `platform_family?` method

The following is an example of using the `platform_family?` method in the Recipe DSL to create a variable that can be used with other resources in the same recipe. In this example, `platform_family?` is being used to ensure that a specific binary is used for a specific platform before using the `remote_file` resource to download a file from a remote location, and then using the `execute` resource to install that file by running a command.

```
if platform_family?("rhel")
 pip_binary = "/usr/bin/pip"
else
 pip_binary = "/usr/local/bin/pip"
end

remote_file "#{Chef::Config[:file_cache_path]}/distribute_setup.py" do
 source "http://python-distribute.org/distribute_setup.py"
 mode "0644"
 not_if { ::File.exists?(pip_binary) }
end

execute "install-pip" do
 cwd Chef::Config[:file_cache_path]
 command <<-EOF
 # command for installing Python goes here
 EOF
 not_if { ::File.exists?(pip_binary) }
end
```

where a command for installing Python might look something like:

```
#{node['python']['binary']} distribute_setup.py
#{::File.dirname(pip_binary)}/easy_install pip
```

## route

The `route` resource is used to manage the system routing table in a Linux environment.

### Syntax

The syntax for using the `route` resource in a recipe is as follows:

```
route "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `route` tells the chef-client to use the `Chef::Provider::Route` provider during the chef-client run
- `name` is the name of the resource block; when the `target` attribute is not specified as part of a recipe, `name` is also the IP address of the target route
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

### Actions

This resource has the following actions:

Action	Description
<code>:add</code>	Default. Use to add a route.
<code>:delete</code>	Use to delete a route.

### Attributes

This resource has the following attributes:

Attribute	Description
<code>device</code>	The network interface to which the route applies.
<code>gateway</code>	The gateway for the route.
<code>netmask</code>	The decimal representation of the network mask. For example: <code>255.255.255.0</code> .
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>target</code>	The IP address of the target route. Default value: the <code>name</code> of the resource block (see Syntax section above).

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Route</code>	<code>route</code>	The default provider for all platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Add a host route

```
route "10.0.1.10/32" do
 gateway "10.0.0.20"
 device "eth1"
end
```

### Delete a network route

```
route "10.1.1.0/24" do
 gateway "10.0.0.20"
 action :delete
end
```

## rpm\_package

The **rpm\_package** resource is used to manage packages for the RPM Package Manager platform.

### Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

## Syntax

The syntax for using the **rpm\_package** resource in a recipe is as follows:

```
rpm_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `rpm_package` tells the chef-client to use the `Chef::Provider::Rpm` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:remove</code>	Use to remove a package.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.



<u>response_file</u>	Optional. The direct path to the file used to pre-seed a package.
<u>source</u>	Optional. The package source for providers that use a local file.
<u>version</u>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<u>Chef::Provider::Package</u>	<u>package</u>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<u>Chef::Provider::Package::Rpm</u>	<u>rpm_package</u>	The provider that is used with the RPM Package Manager platform. Can be used with the <u>options</u> attribute.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package

```
rpm_package "name of package" do
 action :install
end
```

ruby

The **ruby** resource is used to execute scripts using the Ruby interpreter and includes all of the actions and attributes that are available to the **execute** resource.

Note

The **ruby** script resource (which is based on the **script** resource) is different from the **ruby\_block** resource because Ruby code that is run with this resource is created as a temporary file and executed like other script resources, rather than run inline. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the not\_if and only\_if meta parameters to guard the use of this resource for idempotence.

Syntax

The syntax for using the **ruby** resource in a recipe is as follows:

```
ruby "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- **ruby** tells the chef-client to use the `Chef::Resource::Script::Ruby` provider during the chef-client run
- **name** is the name of the resource block; when the `command` attribute is not specified as part of a recipe, **name** is also the name of the command to be executed
- **attribute** is zero (or more) of the attributes that are available for this resource
- **:action** is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<u>:run</u>	Default. Use to run a script.
<u>:nothing</u>	Indicates that the command should not be run. This action is used to specify that a command is run only when another resource notifies it.

Attributes

This resource has the following attributes:

Attribute	Description
<u>code</u>	A quoted ( " ") string of code to be executed.
<u>command</u>	The name of the command to be executed. Default value: the <u>name</u> of the resource block (see Syntax section above).

<a href="#"><u>creates</u></a>	Indicates that a command to create a file will not be run when that file already exists.
<a href="#"><u>cwd</u></a>	The current working directory.
<a href="#"><u>environment</u></a>	A Hash of environment variables in the form of { "ENV_VARIABLE" => "VALUE" }. (These variables must exist for a command to be run successfully.)
<a href="#"><u>flags</u></a>	One (or more) command line flags that are passed to the interpreter when a command is invoked.
<a href="#"><u>group</u></a>	The group name or group ID that must be changed before running a command.
<a href="#"><u>path</u></a>	An array of paths to use when searching for a command. These paths are not added to the command's environment \$PATH. The default value uses the system path.
<a href="#"><u>provider</u></a>	Optional. Use to specify a provider by using its long name. For example: <a href="#"><u>provider Chef::Provider::Long::Name</u></a> . See the Providers section below for the list of providers available to this resource.
<a href="#"><u>returns</u></a>	The return value for a command. This may be an array of accepted values. An exception is raised when the return value(s) do not match. Default value: <code>0</code> .
<a href="#"><u>timeout</u></a>	The amount of time (in seconds) a command will wait before timing out. Default value: <code>3600</code> .
<a href="#"><u>user</u></a>	The user name or user ID that should be changed before running a command.
<a href="#"><u>umask</u></a>	The file mode creation mask, or umask.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<a href="#"><u>Chef::Provider::Script</u></a>	<a href="#"><u>script</u></a>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<a href="#"><u>Chef::Provider::Script::Ruby</u></a>	<a href="#"><u>ruby</u></a>	The provider that is used with the Ruby command interpreter.

## Examples

None.

## ruby\_block

The **ruby\_block** resource is used to execute Ruby code during a chef-client run. Ruby code in the **ruby\_block** resource is evaluated with other resources during convergence, whereas Ruby code outside of a **ruby\_block** resource is evaluated before other resources, as the recipe is compiled.

## Syntax

The syntax for using the **ruby\_block** resource in a recipe is as follows:

```
ruby_block "name" do
 block do
 # some Ruby code
 end
 action :action # see actions section below
end
```

where

- [ruby\\_block](#) tells the chef-client to use the [Chef::Provider::RubyBlock](#) provider during the chef-client run
- [name](#) is the name of the resource block; when the [block\\_name](#) attribute is not specified as part of a recipe, [name](#) is also the name of the Ruby block
- [block](#) is the attribute that is used to define the Ruby block
- [:action](#) is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:

Action	Description
<a href="#"><u>:create</u></a>	Default. Use to create a Ruby block.

## Attributes

This resource has the following attributes:

Attribute	Description
-----------	-------------

<code>block</code>	A block of Ruby code.
<code>block_name</code>	The name of the Ruby block. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::RubyBlock</code>	<code>ruby_block</code>	The default provider for all platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Re-read configuration data

```
ruby_block "reload_client_config" do
 block do
 Chef::Config.from_file("/etc/chef/client.rb")
 end
 action :create
end
```

### Install repositories from a file, trigger a command, and force the internal cache to reload

The following example shows how to install new yum repositories from a file, where the installation of the repository triggers a creation of the yum cache that forces the internal cache for the chef-client to reload:

```
execute "create-yum-cache" do
 command "yum -q makecache"
 action :nothing
end

ruby_block "reload-internal-yum-cache" do
 block do
 Chef::Provider::Package::Yum::YumCache.instance.reload
 end
 action :nothing
end

cookbook_file "/etc/yum.repos.d/custom.repo" do
 source "custom"
 mode 00644
 notifies :run, "execute[create-yum-cache]", :immediately
 notifies :create, "ruby_block[reload-internal-yum-cache]", :immediately
end
```

### Use an if statement with the platform recipe DSL method

The following example shows how an if statement can be used with the `platform?` method in the Recipe DSL to run code specific to Microsoft Windows. The code is defined using the `ruby_block` resource:

```
the following code sample comes from the `client` recipe in the following cookbook: https://github.c
if platform?("windows")
 ruby_block "copy libmysql.dll into ruby path" do
 block do
 require 'fileutils'
 FileUtils.cp "#{node['mysql']['client']['lib_dir']}\\libmysql.dll",
 node['mysql']['client']['ruby_dir']
 end
 not_if { File.exist?("#{node['mysql']['client']['ruby_dir']}\\libmysql.dll") }
 end
end
```

### Stash a file in a data bag

The following example shows how to use the `ruby_block` resource to stash a BitTorrent file in a data bag so that it can be distributed to nodes in the organization.

```
the following code sample comes from the `seed` recipe in the following cookbook: https://github.com
ruby_block "share the torrent file" do
 block do
 f = File.open(node['bittorrent']['torrent'], 'rb')
 #read the .torrent file and base64 encode it
 enc = Base64.encode64(f.read)
 data = {
 'id'=>bittorrent_item_id(node['bittorrent']['file']),
 'seed'=>node.ipaddress,
 'torrent'=>enc
 }
 end
end
```

```

item = Chef::DataBagItem.new
item.data_bag('bittorrent')
item.raw_data = data
item.save
end
action :nothing
subscribes :create, "bittorrent_torrent[#{node['bittorrent']['torrent']}]"
end

```

### Update the /etc/hosts file

The following example shows how the `ruby_block` resource can be used to update the `/etc/hosts` file:

*# the following code sample comes from the ``ec2`` recipe in the following cookbook: <https://github.com/>*

```

ruby_block "edit etc hosts" do
 block do
 rc = Chef::Util::FileEdit.new("/etc/hosts")
 rc.search_file_replace_line(/^127\.0\.0\.1 localhost$/,
 "127.0.0.1 #{new_fqdn} #{new_hostname} localhost")
 rc.write_file
 end
end

```

### Set environment variables

The following example shows how to use variables within a Ruby block to set environment variables using `rbenv`.

```

node.set[:rbenv][:root] = rbenv_root
node.set[:ruby_build][:bin_path] = rbenv_binary_path

ruby_block "initialize" do
 block do
 ENV['RBENV_ROOT'] = node[:rbenv][:root]
 ENV['PATH'] = "#{node[:rbenv][:root]}/bin:#{node[:ruby_build][:bin_path]}:#{ENV['PATH']}"
 end
end

```

### Set JAVA\_HOME

The following example shows how to use a variable within a Ruby block to set the `java_home` environment variable:

```

ruby_block "set-env-java-home" do
 block do
 ENV['JAVA_HOME'] = java_home
 end
end

```

### Run specific blocks of Ruby code on specific platforms

The following example shows how the `platform?` method and an if statement can be used in a recipe along with the `ruby_block` resource to run certain blocks of Ruby code on certain platforms:

```

if platform?("ubuntu", "debian", "redhat", "centos", "fedora", "scientific", "amazon")
 ruby_block "update-java-alternatives" do
 block do
 if platform?("ubuntu", "debian") and version == 6
 run_context = Chef::RunContext.new(node, {})
 r = Chef::Resource::Execute.new("update-java-alternatives", run_context)
 r.command "update-java-alternatives -s java-6-openjdk"
 r.run_action(:create)
 else
 require "fileutils"
 arch = node['kernel']['machine'] =~ /x86_64/ ? "x86_64" : "i386"
 Chef::Log.debug("glob is #{java_home_parent}/java*#{version}*openjdk*")
 jdk_home = Dir.glob("#{java_home_parent}/java*#{version}*openjdk*")[0]
 Chef::Log.debug("jdk_home is #{jdk_home}")

 if File.exists? java_home
 FileUtils.rm_f java_home
 end
 FileUtils.ln_sf jdk_home, java_home

 cmd = Chef::ShellOut.new(
 %Q[update-alternatives --install /usr/bin/java java #{java_home}/bin/java 1;
 update-alternatives --set java #{java_home}/bin/java]
).run_command
 unless cmd.exitstatus == 0 or cmd.exitstatus == 2
 Chef::Application.fatal!("Failed to update-alternatives for openjdk!")
 end
 end
 end
 end
 action :nothing
end

```

### Reload the configuration

The following example shows how to reload the configuration of a chef-client using the `remote_file` resource to:

- using an `if` statement to check whether the plugins on a node are the latest versions
- identify the location from which Ohai plugins are stored
- using the `notifies` attribute and a `ruby_block` resource to trigger an update (if required) and to then reload the client.rb file.

```

directory node[:ohai][:plugin_path] do
 owner "chef"
 recursive true
end

ruby_block "reload_config" do
 block do
 Chef::Config.from_file("/etc/chef/client.rb")
 end
 action :nothing
end

if node[:ohai].key?(:plugins)
 node[:ohai][:plugins].each do |plugin|
 remote_file node[:ohai][:plugin_path] + "/" + plugin do
 source plugin
 owner "chef"
 notifies :create, resources(:ruby_block => "reload_config")
 end
 end
end
end

```

## script

The **script** resource is used to execute scripts using the specified interpreter (Bash, csh, Perl, Python, or Ruby) and includes all of the actions and attributes that are available to the **execute** resource.

### Note

The **script** resource is different from the **ruby\_block** resource because Ruby code that is run with this resource is created as a temporary file and executed like other script resources, rather than run inline. Commands that are executed with this resource are (by their nature) not idempotent, as they are typically unique to the environment in which they are run. Use the `not_if` and `only_if` meta parameters to guard the use of this resource for idempotence.

## Syntax

The syntax for using the **script** resource in a recipe is as follows:

```

script "name" do
 some_attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end

```

where

- `script` tells the chef-client to use one of the following providers during the chef-client run: `Chef::Resource::Script`, `Chef::Resource::Script::Bash`, `Chef::Resource::Script::Csh`, `Chef::Resource::Script::Perl`, `Chef::Resource::Script::Python`, or `Chef::Resource::Script::Ruby`. The provider that is used by the chef-client depends on the platform of the machine on which the run is taking place
- `name` is the name of the resource block; when the `command` attribute is not specified as part of a recipe, `name` is also the name of the command to be executed
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

## Actions

This resource has the following actions:

Action	Description
<code>:run</code>	Default. Use to run a script.
<code>:nothing</code>	Indicates that the command should not be run. This action is used to specify that a command is run only when another resource notifies it.

## Attributes

This resource has the following attributes:

Attribute	Description
<code>code</code>	A quoted (") string of code to be executed.
<code>command</code>	The name of the command to be executed. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>creates</code>	Indicates that a command to create a file will not be run when that file already exists.
<code>cwd</code>	The current working directory.
<code>environment</code>	A Hash of environment variables in the form of <code>{"ENV_VARIABLE" =&gt; "VALUE"}</code> . (These variables must exist for a command to be run successfully.)
<code>flags</code>	One (or more) command line flags that are passed to the interpreter when a command is invoked.

<code>group</code>	The group name or group ID that must be changed before running a command.
<code>interpreter</code>	The script interpreter to be used during code execution.
<code>path</code>	An array of paths to use when searching for a command. These paths are not added to the command's environment \$PATH. The default value uses the system path.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>returns</code>	The return value for a command. This may be an array of accepted values. An exception is raised when the return value(s) do not match. Default value: <code>0</code> .
<code>timeout</code>	The amount of time (in seconds) a command will wait before timing out. Default value: <code>3600</code> .
<code>user</code>	The user name or user ID that should be changed before running a command.
<code>umask</code>	The file mode creation mask, or umask.

Providers

The following providers are available. Use the short name to use the provider in a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Script</code>	<code>script</code>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<code>Chef::Provider::Script::Bash</code>	<code>bash</code>	The provider that is used with the Bash command interpreter.
<code>Chef::Provider::Script::Csh</code>	<code>csh</code>	The provider that is used with the csh command interpreter.
<code>Chef::Provider::Script::Perl</code>	<code>perl</code>	The provider that is used with the Perl command interpreter.
<code>Chef::Provider::Script::Python</code>	<code>python</code>	The provider that is used with the Python command interpreter.
<code>Chef::Provider::Script::Ruby</code>	<code>ruby</code>	The provider that is used with the Ruby command interpreter.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Use a named provider to run a script

```
bash "install_something" do
 user "root"
 cwd "/tmp"
 code <<-EOH
 wget http://www.example.com/tarball.tar.gz
 tar -zxf tarball.tar.gz
 cd tarball
 ./configure
 make
 make install
 EOH
end
```

Run a script

```
script "install_something" do
 interpreter "bash"
 user "root"
 cwd "/tmp"
 code <<-EOH
 wget http://www.example.com/tarball.tar.gz
 tar -zxf tarball.tar.gz
 cd tarball
 ./configure
 make
 make install
 EOH
end
```

or something like:

```
bash "openvpn-server-key" do
 environment("KEY_CN" => "server")
 code <<-EOF
 openssl req -batch -days #{node["openvpn"]["key"]["expire"]} \
 -nodes -new -newkey rsa:#{key_size} -keyout #{key_dir}/server.key \
 -out #{key_dir}/server.csr -extensions server \
 -config #{key_dir}/openssl.cnf
 EOF
 not_if { ::File.exists?("#{key_dir}/server.crt") }
end
```

where `code` contains the OpenSSL command to be run. The `not_if` method tells the chef-client not to run the command if the file already exists.

#### Install a file from a remote location using bash

The following is an example of how to install the `foo123` module for Nginx. This module adds shell-style functionality to an Nginx configuration file and does the following:

- Declares three variables
- Gets the Nginx file from a remote location
- Installs the file using Bash to the path specified by the `src_filepath` variable

```
the following code sample is similar to the ``upload_progress_module`` recipe in the ``nginx`` cookbook

src_filename = "foo123-nginx-module-v#{node['nginx']['foo123']['version']}.tar.gz"
src_filepath = "#{Chef::Config[:file_cache_path]}/#{src_filename}"
extract_path = "#{Chef::Config[:file_cache_path]}/nginx_foo123_module/#{node['nginx']['foo123']['checksum']}"

remote_file src_filepath do
 source node['nginx']['foo123']['url']
 checksum node['nginx']['foo123']['checksum']
 owner 'root'
 group 'root'
 mode 00644
end

bash 'extract_module' do
 cwd ::File.dirname(src_filepath)
 code <<-EOH
 mkdir -p #{extract_path}
 tar xzf #{src_filename} -C #{extract_path}
 mv #{extract_path}/*/* #{extract_path}/
 EOH
 not_if { ::File.exists?(extract_path) }
end
```

#### Install an application from git using bash

The following example shows how Bash can be used to install a plug-in for rbnb named "ruby-build", which is located in git version source control. First, the application is synchronized, and then Bash changes its working directory to the location in which "ruby-build" is located, and then runs a command.

```
git "#{Chef::Config[:file_cache_path]}/ruby-build" do
 repository "git://github.com/sstephenson/ruby-build.git"
 reference "master"
 action :sync
end

bash "install_ruby_build" do
 cwd "#{Chef::Config[:file_cache_path]}/ruby-build"
 user "rbenv"
 group "rbenv"
 code <<-EOH
 ./install.sh
 EOH
 environment 'PREFIX' => "/usr/local"
end
```

To read more about `ruby-build`, see here: <https://github.com/sstephenson/ruby-build>.

#### Store certain settings

The following recipe shows how an attributes file can be used to store certain settings. An attributes file is located in the `attributes/` directory in the same cookbook as the recipe which calls the attributes file. In this example, the attributes file specifies certain settings for Python that are then used across all nodes against which this recipe will run.

Python packages have versions, installation directories, URLs, and checksum files. An attributes file that exists to support this type of recipe would include settings like the following:

```
default['python']['version'] = '2.7.1'

if python['install_method'] == 'package'
 default['python']['prefix_dir'] = '/usr'
else
 default['python']['prefix_dir'] = '/usr/local'
end

default['python']['url'] = 'http://www.python.org/ftp/python'
default['python']['checksum'] = '80e387...85fd61'
```

and then the methods in the recipe may refer to these values. A recipe that is used to install Python will need to do the following:

- Identify each package to be installed (implied in this example, not shown)
- Define variables for the package `version` and the `install_path`
- Get the package from a remote location, but only if the package does not already exist on the target system
- Use the `bash` resource to install the package on the node, but only when the package is not already installed

```
the following code sample comes from the ``oc-nginx`` cookbook on [github]: https://github.com/cookbook

version = node['python']['version']
install_path = "#{node['python']['prefix_dir']/lib/python#{version.split(/(^d+\\.d+)/)[1]}}"

remote_file "#{Chef::Config[:file_cache_path]}/Python-#{version}.tar.bz2" do
```

```

source #node[python] [['0'], #version]; python #version; cat <<<
checksum node['python']['checksum']
mode "0644"
not_if { ::File.exists?(install_path) }
end

bash "build-and-install-python" do
 cwd Chef::Config[:file_cache_path]
 code <<-EOF
 tar -jxvf Python-#{version}.tar.bz2
 (cd Python-#{version} && ./configure #{configure_options})
 (cd Python-#{version} && make && make install)
 EOF
 not_if { ::File.exists?(install_path) }
end

```

## service

The **service** resource is used to manage a service.

### Syntax

The syntax for using the **service** resource in a recipe is as follows:

```

service "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end

```

where

- [service](#) tells the chef-client to use one of the following providers during the chef-client run: [Chef::Provider::Service::Init](#), [Chef::Provider::Service::Init::Debian](#), [Chef::Provider::Service::Upstart](#), [Chef::Provider::Service::Init::Freebsd](#), [Chef::Provider::Service::Init::Gentoo](#), [Chef::Provider::Service::Init::Redhat](#), [Chef::Provider::Service::Solaris](#), [Chef::Provider::Service::Windows](#), or [Chef::Provider::Service::Macosx](#). The chef-client will detect the platform at the start of the run based on data collected by Ohai. After the platform is identified, the chef-client will determine the correct provider
- [name](#) is the name of the resource block; when the [service\\_name](#) attribute is not specified as part of a recipe, [name](#) is also the name of the service
- [attribute](#) is zero (or more) of the attributes that are available for this resource
- [:action](#) is the step that the resource will ask the provider to take during the chef-client run

### Actions

This resource has the following actions:

Action	Description
<a href="#">:enable</a>	Use to enable a service at boot.
<a href="#">:disable</a>	Use to disable a service.
<a href="#">:nothing</a>	Default. Use to do nothing with a service.
<a href="#">:start</a>	Use to start a service (and keep it running until stopped or disabled).
<a href="#">:stop</a>	Use to stop a service.
<a href="#">:restart</a>	Use to restart a service.
<a href="#">:reload</a>	Use to reload the configuration for this service.

### Attributes

This resource has the following attributes:

Attribute	Description
<a href="#">init_command</a>	The path to the init script associated with the service. This is typically <code>/etc/init.d/SERVICE_NAME</code> . The <a href="#">init_command</a> attribute can be used to prevent the need to specify overrides for the <a href="#">start_command</a> , <a href="#">stop_command</a> , and <a href="#">restart_command</a> attributes. Default value: <code>nil</code> .
<a href="#">pattern</a>	The pattern to look for in the process table. Default value: <a href="#">service_name</a> .
<a href="#">priority</a>	Debian platform only. The relative priority of the program for start and shutdown ordering. May be an integer or a hash. An integer is used to define the start run levels; stop run levels are then 100-integer. A hash is used to define values for specific run levels. For example, <code>{ 2 =&gt; [:start, 20], 3 =&gt; [:stop, 55] }</code> will set a priority of twenty for run level two and a priority of fifty-five for run level three.
<a href="#">provider</a>	Optional. Use to specify a provider by using its long name. For example: <a href="#">provider</a>



`chef::Provider::Long::Name`. See the `Providers` section below for the list of providers available to this resource.

<code>reload_command</code>	The command used to tell a service to reload its configuration.
<code>restart_command</code>	The command used to restart a service.
<code>service_name</code>	The name of the service. Default value: the <code>name</code> of the resource block (see <code>Syntax</code> section above).
<code>start_command</code>	The command used to start a service.
<code>status_command</code>	The command used to check the run status for a service.
<code>stop_command</code>	The command used to stop a service.
<code>supports</code>	A list of attributes that controls how the chef-client will attempt to manage a service: <code>:restart</code> , <code>:reload</code> , <code>:status</code> . For <code>:restart</code> , the init script or other service provider can use a restart command; if <code>:restart</code> is not specified, the chef-client will attempt to stop and then start a service. For <code>:reload</code> , the init script or other service provider can use a reload command. For <code>:status</code> , the init script or other service provider can use a status command to determine if the service is running; if <code>:status</code> is not specified, the chef-client will attempt to match the <code>service_name</code> against the process table as a regular expression, unless a pattern is specified as a parameter attribute. Default value: { <code>:restart =&gt; false</code> , <code>:reload =&gt; false</code> , <code>:status =&gt; false</code> } for all platforms (except for the Red Hat platform family, which defaults to { <code>:restart =&gt; false</code> , <code>:reload =&gt; false</code> , <code>:status =&gt; true</code> }.)

Providers

The **service** resource does not have service-specific short names. This is because the chef-client identifies the platform at the start of every chef-client run based on data collected by Ohai. The chef-client looks up the platform in the `provider_mapping.rb` file, and then determines the correct provider for that platform. In certain situations, such as when more than one init system is available on a node, a specific provider may need to be identified by using the `provider` attribute and the long name for that provider.

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Service::Init</code>	<code>service</code>	When this short name is used, the chef-client will determine the correct provider during the chef-client run.
<code>Chef::Provider::Service::Init::Debian</code>	<code>service</code>	The provider that is used with the Debian and Ubuntu platforms.
<code>Chef::Provider::Service::Upstart</code>	<code>service</code>	The provider that is used when Upstart is available on the platform.
<code>Chef::Provider::Service::Init::Freebsd</code>	<code>service</code>	The provider that is used with the FreeBSD platform.
<code>Chef::Provider::Service::Init::Gentoo</code>	<code>service</code>	The provider that is used with the Gentoo platform.
<code>Chef::Provider::Service::Init::Redhat</code>	<code>service</code>	The provider that is used with the Red Hat and CentOS platforms.
<code>Chef::Provider::Service::Solaris</code>	<code>service</code>	The provider that is used with the Solaris platform.
<code>Chef::Provider::Service::Windows</code>	<code>service</code>	The provider that is used with the Microsoft Windows platform.
<code>Chef::Provider::Service::Macosx</code>	<code>service</code>	The provider that is used with the Mac OS X platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Start a service

```
service "example_service" do
 action :start
end
```

Start a service, enable it

```
service "example_service" do
 supports :status => true, :restart => true, :reload => true
 action [:enable, :start]
end
```

Use a pattern

```
service "samba" do
 pattern "smbd"
 action [:enable, :start]
end
```

**manage a service, depending on the node platform**

```

service "example_service" do
 case node["platform"]
 when "centos", "redhat", "fedora"
 service_name "redhat_name"
 else
 service_name "other_name"
 end
 supports :restart => true
 action [:enable, :start]
end

```

**Change a service provider, depending on the node platform**

```

service "example_service" do
 case node["platform"]
 when "ubuntu"
 if node["platform_version"].to_f >= 9.10
 provider Chef::Provider::Service::Upstart
 end
 end
 action [:enable, :start]
end

```

**Set an IP address using variables and a template**

The following example shows how the **template** resource can be used in a recipe to combine settings stored in an attributes file, variables within a recipe, and a template to set the IP addresses that are used by the Nginx service. The attributes file contains the following:

```
default['nginx']['dir'] = "/etc/nginx"
```

The recipe then does the following to:

- Declare two variables at the beginning of the recipe, one for the remote IP address and the other for the authorized IP address
- Use the **service** resource to restart and reload the Nginx service
- Load a template named "authorized\_ip.erb" from the `/templates` directory that is used to set the IP address values based on the variables specified in the recipe

```

node.default['nginx']['remote_ip_var'] = "remote_addr"
node.default['nginx']['authorized_ips'] = ["127.0.0.1/32"]

service "nginx" do
 supports :status => true, :restart => true, :reload => true
end

template "authorized_ip" do
 path "#{node['nginx']['dir']}/authorized_ip"
 source "modules/authorized_ip.erb"
 owner "root"
 group "root"
 mode 00644
 variables(
 :remote_ip_var => node['nginx']['remote_ip_var'],
 :authorized_ips => node['nginx']['authorized_ips']
)

 notifies :reload, resources(:service => "nginx")
end

```

where the `variables` attribute tells the template to use the variables set at the beginning of the recipe and the `source` attribute is used to call a template file located in the cookbook's `/templates` directory. The template file looks something like:

```

geo $<%= @remote_ip_var %> $authorized_ip {
 default no;
 <% @authorized_ips.each do |ip| %>
 <%= "#{ip} yes;" %>
 <% end %>
}

```

**Use a cron timer to manage a service**

The following example shows how to install the `crond` application using two resources and a variable:

```

the following code sample comes from the ``cron`` cookbook: https://github.com/opscode-cookbooks/cron

cron_package = case node['platform']
when "redhat", "centos", "scientific", "fedora", "amazon"
 node['platform_version'].to_f >= 6.0 ? "cronie" : "vixie-cron"
else
 "cron"
end

package cron_package do
 action :install
end

service "crond" do
 case node['platform']
 when "redhat", "centos", "scientific", "fedora", "amazon"
 service_name "crond"
 when "debian", "ubuntu", "suse"
 service_name "cron"
 end
end

```

```

 action [:start, :enable]
 end

```

where

- `cron_package` is a variable that is used to identify which platforms apply to which install packages
- the **package** resource uses the `cron_package` variable to determine how to install the `crond` application on various nodes (with various platforms)
- the **service** resource enables the `crond` application on nodes that have Red Hat, CentOS, Red Hat Enterprise Linux, Fedora, or Amazon Web Services, and the `cron` service on nodes that run Debian, Ubuntu, or SuSE.

#### Restart a service, and then notify a different service

The following example shows how start a service named "example\_service" and immediately notify the Nginx service to restart.

```

service "example_service" do
 action :start
 provider Chef::Provider::Service::Init
 notifies :restart, "service[nginx]", :immediately
end

```

where by using the default `provider` for the **service**, the recipe is telling the chef-client to determine the specific provider to be used during the chef-client run based on the platform of the node on which the recipe will run.

#### Stop a service, do stuff, and then restart it

The following example shows how to use the **execute**, **service**, and **mount** resources together to ensure that a node running on Amazon EC2 is running MySQL. This example does the following:

- Checks to see if the Amazon EC2 node has MySQL
- If the node has MySQL, stops MySQL
- Installs MySQL
- Mounts the node
- Restarts MySQL

*# the following code sample comes from the ``server\_ec2`` recipe in the following cookbook: <https://github.com>*

```

if (node.attribute?('ec2') && ! FileTest.directory?(node['mysql']['ec2_path']))

 service "mysql" do
 action :stop
 end

 execute "install-mysql" do
 command "mv #{node['mysql']['data_dir']} #{node['mysql']['ec2_path']}"
 not_if do FileTest.directory?(node['mysql']['ec2_path']) end
 end

 [node['mysql']['ec2_path'], node['mysql']['data_dir']].each do |dir|
 directory dir do
 owner "mysql"
 group "mysql"
 end
 end

 mount node['mysql']['data_dir'] do
 device node['mysql']['ec2_path']
 fstype "none"
 options "bind,rw"
 action [:mount, :enable]
 end

 service "mysql" do
 action :start
 end
end

```

where

- the two **service** resources are used to stop, and then restart the MySQL service
- the **execute** resource is used to install MySQL
- the **mount** resource is used to mount the node and enable MySQL

#### Control a service using the execute resource

##### Warning

This is an example of something that should NOT be done. Use the **service** resource to control a service, not the **execute** resource.

Do something like this:

```

service "tomcat" do
 action :start
end

```

and NOT something like this:

```

execute "start-tomcat" do
 command "/etc/init.d/tomcat6 start"
 action :run
end

```

There is no reason to use the **execute** resource to control a service because the **service** resource exposes the `start_command` attribute directly, which gives a recipe full control over the command issued in a much cleaner, more direct manner.

### smartos\_package

The **smartos\_package** resource is used to manage packages for the SmartOS platform.

Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

### Syntax

The syntax for using the **smartos\_package** resource in a recipe is as follows:

```
smartos_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `smartos_package` tells the chef-client to use the `Chef::Provider::Smartos` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

### Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:remove</code>	Use to remove a package.

### Attributes

This resource has the following attributes:

Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

### Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Smartos</code>	<code>smartos_package</code>	The provider that is used with the SmartOS platform.

### Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package

```
smartos_package "name of package" do
 action :install
end
```

solaris\_package

The **solaris\_package** resource is used to manage packages for the Solaris platform.

Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

Syntax

The syntax for using the **solaris\_package** resource in a recipe is as follows:

```
solaris_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `solaris_package` tells the chef-client to use the `Chef::Provider::Solaris` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:remove</code>	Use to remove a package.

Attributes

This resource has the following attributes:

Attribute	Description
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	The name of the package. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>response_file</code>	Optional. The direct path to the file used to pre-seed a package.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Solaris</code>	<code>solaris_package</code>	The provider that is used with the Solaris platform.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Install a package

```
solaris_package "name of package" do
 action :install
end
```

subversion

The **subversion** resource is used to manage source control resources that exist in a Subversion repository.

Note
This resource is often used in conjunction with the <b>deploy</b> resource.

Syntax

The syntax for using the subversion resource in a recipe is as follows:

```
subversion "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `subversion` tells the chef-client to use the `Chef::Provider::Subversion` provider during the chef-client run.
- `"name"` is the location in which the source files will be placed and/or synchronized with the files under source control management
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

The following is an example shows the subversion resource:

```
subversion "CouchDB Edge" do
 repository "http://svn.apache.org/repos/asf/couchdb/trunk"
 revision "HEAD"
 destination "/opt/mysources/couch"
 action :sync
end
```

where

- the name of the resource is `CouchDB Edge`
- the `repository` and `reference` nodes tell the chef-client which repository and revision to use

Actions

This resource has the following actions:

Action	Description
<code>:sync</code>	Default. Use to update the source to the specified version, or to get a new clone or checkout.
<code>:checkout</code>	Use to clone or check out the source. When a checkout is available, this provider does nothing.
<code>:export</code>	Use to export the source, excluding or removing any version control artifacts.
<code>:force_export</code>	Use to export the source, excluding or removing any version control artifacts and to force an export of the source that is overwriting the existing copy (if it exists).

Attributes

This resource has the following attributes:

Attribute	Description
<code>destination</code>	The path to the location to which the source will be cloned, checked out, or exported. Default value: the <code>name</code> of the resource block (see Syntax section above).
<code>group</code>	The system group that is responsible for the checked-out code.
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>repository</code>	The URI for the Subversion repository.
<code>revision</code>	The revision to be checked out. This can be symbolic, like <code>HEAD</code> or it can be a source control management-specific revision identifier. Default value: <code>HEAD</code> .

<a href="#">svn_arguments</a>	The extra arguments that are passed to the Subversion command.
<a href="#">svn_info_args</a>	Use when the <code>svn_info</code> command is used by the chef-client and arguments need to be passed. (The <code>svn_arguments</code> command does not work when the <code>svn_info</code> command is used.)
<a href="#">svn_password</a>	The password for the user that has access to the Subversion repository.
<a href="#">svn_username</a>	The user name for a user that has access to the Subversion repository.
<a href="#">user</a>	The system user that is responsible for the checked-out code.

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Subversion</code>	<code>subversion</code>	This provider work only with Subversion.

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

Get the latest version of an application

```
subversion "CouchDB Edge" do
 repository "http://svn.apache.org/repos/asf/couchdb/trunk"
 revision "HEAD"
 destination "/opt/mysources/couch"
 action :sync
end
```

template

The **template** resource is used to manage file contents with an embedded Ruby (erb) template. This resource includes actions and attributes from the **file** resource. Template files managed by the **template** resource follow the same file specificity rules as the **remote\_file** and **file** resources.

Syntax

The syntax for using the **template** resource in a recipe is as follows:

```
template "name" do
 source "template_name.erb"
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- `template` tells the chef-client to use the `Chef::Provider::File::Template` provider during the chef-client run
- `name` is the path to the location in which a file will be created and the name of the file to be managed; for example: `/var/www/html/index.html`, where `/var/www/html/` is the path to the location and `index.html` is the name of the file
- `source` is the template file that will be used to create the file on the node, for example: `index.html.erb`; the template file is located in the `/templates` directory of a cookbook
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a file.
<code>:create_if_missing</code>	Use to create a file only if the file does not exist. (When the file exists, nothing happens.)
<code>:delete</code>	Use to delete a file.
<code>:touch</code>	Use to touch a file. This updates the access (atime) and file modification (mtime) times for a file. (This action may be used with this resource, but is typically only used with the <b>file</b> resource.)

Attributes

This resource has the following attributes:

Attribute	Description
<code>atomic_update</code>	Indicates whether atomic file updates are used on a per-resource basis. Set to <code>true</code> for atomic file

updates. Set to `false` for non-atomic file updates. (This setting overrides `file_atomic_update`, which is a global setting found in the `client.rb` file.) Default value: `true`.

<a href="#"><code>backup</code></a>	The number of backups to be kept. Set to <code>false</code> to prevent backups from being kept. Default value: <code>5</code> .
<a href="#"><code>cookbook</code></a>	The cookbook in which a file is located (if it is not located in the current cookbook). The default value is the current cookbook.
<a href="#"><code>force_unlink</code></a>	Use to specify how the chef-client handles certain situations when the target file turns out not to be a file. For example, when a target file is actually a symlink. Set to <code>true</code> to have the chef-client delete the non-file target and replace it with the specified file. Set to <code>false</code> for the chef-client to raise an error. Default value: <code>false</code> .
<a href="#"><code>group</code></a>	A string or ID that identifies the group owner by group name, including fully qualified group names such as <code>domain\group</code> or <code>group@domain</code> . If this value is not specified, existing groups will remain unchanged and new group assignments will use the default POSIX group (if available).
<a href="#"><code>helper</code></a>	Use to define a helper method inline. For example: <code>helper(:hello_world) { "hello world" }</code> or <code>helper(:app) { node["app"] }</code> or <code>helper(:app_conf) { [setting] node["app"] [setting] }</code> . Default value: <code>{}</code> .
<a href="#"><code>helpers</code></a>	Use to define a helper module inline or in a library. For example, an inline module: <code>helpers do</code> , which is then followed by a block of Ruby code. And for a library module: <code>helpers(MyHelperModule)</code> . Default value: <code>[]</code> .
<a href="#"><code>inherits</code></a>	Microsoft Windows only. Indicates that a file inherits rights from its parent. Default value: <code>true</code> .
<a href="#"><code>local</code></a>	Use to load a template from a local path. By default, the chef-client loads templates from a cookbook's <code>/templates</code> directory. When this attribute is set to <code>true</code> , use the <code>source</code> attribute specify the path to a template on the local node. Default value: <code>false</code> .
<a href="#"><code>manage_symlink_source</code></a>	Indicates that the chef-client will detect and manage the source file for a symlink. Possible values: <code>nil</code> , <code>true</code> , or <code>false</code> . When this value is set to <code>nil</code> , the chef-client will manage a symlink's source file and emit a warning. When this value is set to <code>true</code> , the chef-client will manage a symlink's source file and not emit a warning. Default value: <code>nil</code> . The default value will be changed to <code>false</code> in a future version.
<a href="#"><code>mode</code></a>	<p>The octal mode for a file. If mode is not specified and if the file already exists, the existing mode on the file is used. If mode is not specified, the file does not exist, and the <code>:create</code> action is specified, the chef-client will assume a mask value of <code>0777</code> and then apply the umask for the system on which the file will be created to the <code>mask</code> value. For example, if the umask on a system is <code>022</code>, the chef-client would use the default value of <code>0755</code>.</p> <p>The behavior is different depending on the platform.</p> <p>UNIX- and Linux-based systems: The octal mode that is passed to <code>chmod</code>. If the value is specified as a quoted string, it will work exactly as if the <code>chmod</code> command was passed. If the value is specified as an integer, prepend a zero (<code>0</code>) to the value to ensure it is interpreted as an octal number. For example, to assign read, write, and execute rights for all users, use <code>0777</code> or <code>'777'</code>; for the same rights, plus the sticky bit, use <code>01777</code> or <code>'1777'</code>.</p> <p>Microsoft Windows: The octal mode that is translated into rights for Microsoft Windows security. Values up to <code>0777</code> are allowed (no sticky bits) and mean the same in Microsoft Windows as they do in UNIX, where <code>4</code> equals <code>GENERIC_READ</code>, <code>2</code> equals <code>GENERIC_WRITE</code>, and <code>1</code> equals <code>GENERIC_EXECUTE</code>. This attribute cannot be used to set <code>:full_control</code>. This attribute has no effect if not specified, but when this attribute and <code>rights</code> are both specified, the effects will be cumulative.</p>
<a href="#"><code>owner</code></a>	A string or ID that identifies the group owner by user name, including fully qualified user names such as <code>domain\user</code> or <code>user@domain</code> . If this value is not specified, existing owners will remain unchanged and new owner assignments will use the current user (when necessary).
<a href="#"><code>path</code></a>	<p>The path to the file.</p> <p>Microsoft Windows: A path that begins with a forward slash (<code>/</code>) will point to the root of the current working directory of the chef-client process. This path can vary from system to system. Therefore, using a path that begins with a forward slash (<code>/</code>) is not recommended.</p>
<a href="#"><code>provider</code></a>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<a href="#"><code>rights</code></a>	Microsoft Windows only. The permissions for users and groups in a Microsoft Windows environment. For example: <code>rights &lt;permissions&gt;, &lt;principal&gt;, &lt;options&gt;</code> where <code>&lt;permissions&gt;</code> specifies the rights granted to the principal, <code>&lt;principal&gt;</code> is the group or user name, and <code>&lt;options&gt;</code> is a Hash with one (or more) advanced rights options.
<a href="#"><code>source</code></a>	The location of a template file. By default, the chef-client will look for a template file in the



`/templates` directory of a cookbook. When the `local` attribute is set to `true`, use this attribute to specify path to a template on the local node. This attribute may also be used to distribute specific files to specific platforms (see the section "File Specificity", below). Default value: the `name` of the resource block (see Syntax section above).

`variables`

A Hash of variables that are passed into a Ruby template file.

This attribute can also be used to reference a partial template file by using a Hash. For example:

```
template "/file/name.txt" do
 variables :partials => {
 "partial_name_1.txt.erb" => "message",
 "partial_name_2.txt.erb" => "message",
 "partial_name_3.txt.erb" => "message"
 },
end
```

where each of the partial template files can then be combined using normal Ruby template patterns within a template file, such as:

```
<% @partials.each do |partial, message| %>
 Here is <%= partial %>
<%= render partial, :variables => {:message => message} %>
<% end %>
```

Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::File::Template</code>	<code>template</code>	The default provider for all platforms.

File Specificity

A cookbook will frequently be designed to work across many platforms and will often be required to distribute a specific file to a specific platform. A cookbook can be designed to support distributing files across platforms, but ensuring that the right file ends up on each system.

Pattern

The pattern for file specificity is as follows:

1. `host-node[:fqdn]`
2. `node[:platform]-node[:platform_version]`
3. `node[:platform]-version_components`: The version string is split on decimals and searched from greatest specificity to least; for example, if the location from the last rule was `centos-5.7.1`, then `centos-5.7` and `centos-5` would also be searched.
4. `node[:platform]`
5. `default`

Example

A cookbook may have a `/templates` directory structure like this:

```
templates/
 windows-6.2
 windows-6.1
 windows-6.0
 windows
 default
```

and a resource that looks something like the following:

```
template "C:\path\to\file\text_file.txt" do
 source "text_file.txt"
 mode 0755
 owner "root"
 group "root"
end
```

This resource would be matched in the same order as the `/templates` directory structure. For a node named "host-node-desktop" that is running Windows 7, the second item would be the matching item and the location:

```
/templates
 windows-6.2/text_file.txt
 windows-6.1/text_file.txt
 windows-6.0/text_file.txt
 windows/text_file.txt
 default/text_file.txt
```

Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

**Configure a file from a template**

```
template "/tmp/config.conf" do
 source "config.conf.erb"
end
```

**Configure a file from a local template**

```
template "/tmp/config.conf" do
 local true
 source "/tmp/config.conf.erb"
end
```

**Configure a file using a variable map**

```
template "/tmp/config.conf" do
 source "config.conf.erb"
 variables(
 :config_var => node["configs"]["config_var"]
)
end
```

**Use the ``not\_if`` condition**

The following example shows how to use the `not_if` condition to create a file based on a template and using the presence of an attribute on the node to specify the condition:

```
template "/tmp/somefile" do
 mode 00644
 source "somefile.erb"
 not_if { node[:some_value] }
end
```

The following example shows how to use the `not_if` condition to create a file based on a template and then Ruby code to specify the condition:

```
template "/tmp/somefile" do
 mode 00644
 source "somefile.erb"
 not_if do
 File.exists?("/etc/passwd")
 end
end
```

The following example shows how to use the `not_if` condition to create a file based on a template and using a Ruby block (with curly braces) to specify the condition:

```
template "/tmp/somefile" do
 mode 00644
 source "somefile.erb"
 not_if { File.exists?("/etc/passwd") }
end
```

The following example shows how to use the `not_if` condition to create a file based on a template and using a string to specify the condition:

```
template "/tmp/somefile" do
 mode 00644
 source "somefile.erb"
 not_if "test -f /etc/passwd"
end
```

**Use the ``only\_if`` condition**

The following example shows how to use the `only_if` condition to create a file based on a template and using the presence of an attribute on the node to specify the condition:

```
template "/tmp/somefile" do
 mode 00644
 source "somefile.erb"
 only_if { node[:some_value] }
end
```

The following example shows how to use the `only_if` condition to create a file based on a template, and then use Ruby to specify a condition:

```
template "/tmp/somefile" do
 mode 00644
 source "somefile.erb"
 only_if do ! File.exists?("/etc/passwd") end
end
```

The following example shows how to use the `only_if` condition to create a file based on a template and using a string to specify the condition:

```
template "/tmp/somefile" do
 mode 00644
 source "somefile.erb"
 only_if "test -f /etc/passwd"
end
```

**Use a whitespace array ( %w )**

The following example shows how to use a Ruby whitespace array to define a list of configuration tools, and then use that list of tools within the **template** resource to ensure that all of these configuration tools are using the same RSA key:

```
%w{openssl.cnf pkitool vars Rakefile}.each do |f|
 template "/etc/openvpn/easy-rsa/#{f}" do
 source "#{f}.erb"
 owner "root"
 group "root"
 mode 0755
 end
end
```

**Use a relative path**

```
template "#{ENV['HOME']}/chef-getting-started.txt" do
 source "chef-getting-started.txt.erb"
 mode 00644
end
```

**Delay notifications**

```
template "/etc/nagios3/configures-nagios.conf" do
 # other parameters
 notifies :run, "execute[test-nagios-config]", :delayed
end
```

**Notify immediately**

By default, notifications are `:delayed`, that is they are queued up as they are triggered, and then executed at the very end of a chef-client run. To run an action immediately, use `:immediately`:

```
template "/etc/nagios3/configures-nagios.conf" do
 # other parameters
 notifies :run, "execute[test-nagios-config]", :immediately
end
```

and then the chef-client would immediately run the following:

```
execute "test-nagios-config" do
 command "nagios3 --verify-config"
 action :nothing
end
```

**Notify multiple resources**

```
template "/etc/chef/server.rb" do
 source "server.rb.erb"
 owner "root"
 group "root"
 mode "644"
 notifies :restart, "service[chef-solr]", :delayed
 notifies :restart, "service[chef-solr-indexer]", :delayed
 notifies :restart, "service[chef-server]", :delayed
end
```

**Reload a service**

```
template "/tmp/somefile" do
 mode "0644"
 source "somefile.erb"
 notifies :reload, "service[apache]"
end
```

**Restart a service when a template is modified**

```
template "/etc/www/configures-apache.conf" do
 notifies :restart, "service[apache]"
end
```

**Send notifications to multiple resources**

To send notifications to multiple resources, just use multiple attributes. Multiple attributes will get sent to the notified resources in the order specified.

```
template "/etc/netatalk/netatalk.conf" do
 notifies :restart, "service[afpd]", :immediately
 notifies :restart, "service[cnid]", :immediately
end

service "afpd"
service "cnid"
```

**Execute a command using a template**

The following example shows how to set up IPv4 packet forwarding using the **execute** resource to run a command named "forward\_ipv4" that uses a template defined by the **template** resource:

```
execute "forward_ipv4" do
 command "echo > /proc/.../ipv4/ip_forward"
 action :nothing
end
```

```

end

template "/etc/file_name.conf" do
 source "routing/file_name.conf.erb"
 notifies :run, 'execute[forward_ipv4]', :delayed
end

```

where the `command` attribute for the **execute** resource contains the command that is to be run and the `source` attribute for the **template** resource specifies which template to use. The `notifies` attribute for the **template** specifies that the `execute[forward_ipv4]` (which is defined by the **execute** resource) should be queued up and run at the end of the chef-client run.

#### Set an IP address using variables and a template

The following example shows how the **template** resource can be used in a recipe to combine settings stored in an attributes file, variables within a recipe, and a template to set the IP addresses that are used by the Nginx service. The attributes file contains the following:

```
default['nginx']['dir'] = "/etc/nginx"
```

The recipe then does the following to:

- Declare two variables at the beginning of the recipe, one for the remote IP address and the other for the authorized IP address
- Use the **service** resource to restart and reload the Nginx service
- Load a template named "authorized\_ip.erb" from the `/templates` directory that is used to set the IP address values based on the variables specified in the recipe

```

node.default['nginx']['remote_ip_var'] = "remote_addr"
node.default['nginx']['authorized_ips'] = ["127.0.0.1/32"]

service "nginx" do
 supports :status => true, :restart => true, :reload => true
end

template "authorized_ip" do
 path "#{node['nginx']['dir']}/authorized_ip"
 source "modules/authorized_ip.erb"
 owner "root"
 group "root"
 mode 00644
 variables(
 :remote_ip_var => node['nginx']['remote_ip_var'],
 :authorized_ips => node['nginx']['authorized_ips']
)

 notifies :reload, resources(:service => "nginx")
end

```

where the `variables` attribute tells the template to use the variables set at the beginning of the recipe and the `source` attribute is used to call a template file located in the cookbook's `/templates` directory. The template file looks something like:

```

geo $<%= @remote_ip_var %> $authorized_ip {
 default no;
 <%= @authorized_ips.each do |ip| %>
 <%= "#{ip} yes;" %>
 <%= end %>
}

```

#### Add a rule to an IP table

The following example shows how to add a rule named "test\_rule" to an IP table using the **execute** resource to run a command using a template that is defined by the **template** resource:

```

execute 'test_rule' do
 command "command_to_run"
 --option value
 ...
 --option value
 --source #{node[:name_of_node][:ipsec][:local][:subnet]}
 -j test_rule"
 action :nothing
end

template "/etc/file_name.local" do
 source "routing/file_name.local.erb"
 notifies :run, 'execute[test_rule]', :delayed
end

```

where the `command` attribute for the **execute** resource contains the command that is to be run and the `source` attribute for the **template** resource specifies which template to use. The `notifies` attribute for the **template** specifies that the `execute[test_rule]` (which is defined by the **execute** resource) should be queued up and run at the end of the chef-client run.

#### Apply proxy settings consistently across a Chef organization

The following example shows how a template can be used to apply consistent proxy settings for all nodes of the same type:

```

template "#{node[:matching_node][:dir]}/sites-available/site_proxy.conf" do
 source "site_proxy.matching_node.conf.erb"
 owner "root"
 group "root"
 mode "0644"
 variables(
 :ssl_certificate => "#{node[:matching_node][:dir]}/shared/certificates/site_proxy.crt",
 :ssl_key => "#{node[:matching_node][:dir]}/shared/certificates/site_proxy.key",

```

```

:listen_port =>
:server_name =>
:fqdn =>
:server_options =>
:proxy_options =>
)
end

```

where `matching_node` represents a type of node (like Nginx) and `site_proxy` represents the type of proxy being used for that type of node (like Nexus).

#### Get template settings from a local file

The **template** resource can be used to render a template based on settings contained in a local file on disk or to get the settings from a template in a cookbook. Most of the time, the settings are retrieved from a template in a cookbook. The following example shows how the **template** resource can be used to retrieve these settings from a local file.

The following example is based on a few assumptions:

- The environment is a Ruby on Rails application that needs render a file named `database.yml`
- Information about the application—the user, their password, the server—is stored in a data bag on the Chef server
- The application is already deployed to the system and that only requirement in this example is to render the `database.yml` file

The application source tree looks something like:

```

myapp/
-> config/
 -> database.yml.erb

```

#### Note

There should not be a file named `database.yml` (without the `.erb`), as the `database.yml` file is what will be rendered using the **template** resource.

The deployment of the app will end up in `/srv`, so the full path to this template would be something like `/srv/myapp/current/config/database.yml.erb`.

The content of the template itself may look like this:

```

<%= @rails_env %>:
adapter: <%= @adapter %>
host: <%= @host %>
database: <%= @database %>
username: <%= @username %>
password: <%= @password %>
encoding: 'utf8'
reconnect: true

```

The recipe will be similar to the following:

```

results = search(:node, "role:myapp_database_master AND environment:#{node.chef_environment}")
db_master = results[0]

template "/srv/myapp/shared/database.yml" do
 source "/srv/myapp/current/config/database.yml.erb"
 local true
 variables(
 :rails_env => node.chef_environment,
 :adapter => db_master['myapp']['db_adapter'],
 :host => db_master['fqdn'],
 :database => "myapp_#{node.chef_environment}",
 :username => "myapp",
 :password => "SUPERSECRET",
)
end

```

where:

- the `search` method in the Recipe DSL is used to find the first node that is the database master (of which there should only be one)
- the `:adapter` attribute may also require an attribute to have been set on a role, which then determines the correct adapter

The template will render similar to the following:

```

production:
adapter: mysql
host: domU-12-31-39-14-F1-C3.compute-1.internal
database: myapp_production
username: myapp
password: SUPERSECRET
encoding: utf8
reconnect: true

```

This example showed how to use the **template** resource to render a template based on settings contained in a local file. Some other issues that should be considered when using this type of approach include:

- Should the `database.yml` file be in a `.gitignore` file?
- How do developers run the application locally?
- How does this work with chef-solo?

#### user

The **user** resource is used to add users, update existing users, remove users, and to lock/unlock user passwords.

Note

System attributes are collected by Ohai at the start of every chef-client run. By design, the actions available to the **user** resource are processed **after** the start of the chef-client run. This means that attributes added or modified by the **user** resource during the chef-client run must be reloaded before they can be available to the chef-client. These attributes can be reloaded in two ways: by picking up the values at the start of the (next) chef-client run or by using the ohai resource to reload these attributes during the current chef-client run.

Syntax

The syntax for using the **user** resource in a recipe is as follows:

```
user "name" do
 attribute "value" # see attributes section below
 ...
 action :action # see actions section below
end
```

where

- user** tells the chef-client to use one of the following providers during the chef-client run: `Chef::Provider::User::Useradd`, `Chef::Provider::User::Pw`, `Chef::Provider::User::Dscl`, or `Chef::Provider::User::Windows`. The provider that is used by the chef-client depends on the platform of the machine on which the chef-client run is taking place
- name** is the name of the resource block; when the `username` attribute is not specified as part of a recipe, **name** is also the name of the user
- attribute** is zero (or more) of the attributes that are available for this resource
- :action** is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:create</code>	Default. Use to create a user with given attributes. If the user already exists, use to ensure that the resource is in the correct state (which is effectively the same as <code>:modify</code> ).
<code>:remove</code>	Use to remove a user.
<code>:modify</code>	Use to modify an existing user. This action will raise an exception if the user does not exist.
<code>:manage</code>	Use to manage an existing user. This action will do nothing if the user does not exist.
<code>:lock</code>	Use to lock a user's password.
<code>:unlock</code>	Use to unlock a user's password.

Attributes

This resource has the following attributes:

Attribute	Description
<code>comment</code>	One (or more) comments about the user.
<code>gid</code>	The identifier for the group.
<code>home</code>	The location of the home directory.
<code>password</code>	The password shadow hash. This attribute requires that <code>ruby-shadow</code> be installed. This is part of the Debian package: <code>libshadow-ruby1.8</code> .
<code>provider</code>	Optional. Use to specify a provider by using its long name. For example: <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>shell</code>	The login shell.
<code>supports</code>	A Mash where keys represent features and values are booleans that indicate if that feature is supported. Default value: <code>:manage_home =&gt; false, :non_unique =&gt; false</code> .
<code>system</code>	Indicates that a system user should be created. This attribute may be used with <code>useradd</code> as the provider to create a system user which passes the <code>-r</code> flag to <code>useradd</code> .
<code>uid</code>	The numeric user identifier.
<code>username</code>	The name of the user. Default value: the <code>name</code> of the resource block (see Syntax section above).

Supported Features

The `supports` attribute allows a list of supported features to be identified. There are two features of note:

- `:manage_home` indicates whether a user's home directory will be created when the user is created. When the `useradd` provider is used, `-dm` will be passed to `useradd` (when the `:create` action is used) and `-d` will be passed to `usermod` (when the `:manage` or `:modify` actions are used). If supports `:manage_home=>true`, the `user` resource does not pass the `-d` and `-m` parameters together (i.e. `-dm`) to `usermod`.

When the `Windows` provider is used, Microsoft Windows does not create a home directory for a user until that user logs on for the first time; specifying the home directory does not have any effect as to where Microsoft Windows ultimately places the home directory.

- `:non_unique` indicates whether non-unique UIDs are allowed. This option is currently unused by the existing providers.

## Password Shadow Hash

There are a number of encryption options and tools that can be used to create a password shadow hash. In general, using a strong encryption method like SHA-512 and the `passwd` command in the OpenSSL toolkit is a good approach, however the encryption options and tools that are available may be different from one distribution to another. The following examples show how the command line can be used to create a password shadow hash. When using the `passwd` command in the OpenSSL tool:

```
openssl passwd -1 "theplaintextpassword"
```

When using `mkpasswd`:

```
mkpasswd -m sha-512
```

For more information:

- <http://www.openssl.org/docs/apps/passwd.html>
- Check the local documentation or package repository for the distribution that is being used. For example, on Ubuntu 9.10-10.04, the `mkpasswd` package is required and on Ubuntu 10.10+ the `whois` package is required.

## Providers

The following providers are available. Use the short name to use the provider in a recipe:

Long name	Short name	Notes
<code>Chef::Provider::User::Useradd</code>	<code>user</code>	The default provider for the <code>user</code> resource.
<code>Chef::Provider::User::Pw</code>	<code>user</code>	The provider that is used with the FreeBSD platform.
<code>Chef::Provider::User::Dsc1</code>	<code>user</code>	The provider that is used with the Mac OS X platform.
<code>Chef::Provider::User::Windows</code>	<code>user</code>	The provider that is used with all Microsoft Windows platforms.

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Create a random user

```
user "random" do
 supports :manage_home => true
 comment "Random User"
 uid 1234
 gid "users"
 home "/home/random"
 shell "/bin/bash"
 password "1JJsvHslV$szsCjVEroftprNn4JHtDi."
end
```

### Create a system user

```
user "systemguy" do
 comment "system guy"
 system true
 shell "/bin/false"
end
```

### Create a system user with a variable

The following example shows how to create a system user using a variable called `user_home` where the matching nodes have a group identifier that is the same as the node, and the login shell is `/bin/bash`:

```
user_home = "#{node[:matching_node][:user]}"

user node[:matching_node][:group] do
 gid node[:matching_node][:group]
 shell "/bin/bash"
 home user_home
 system true
 action :create
end
```

where `matching_node` represents a type of node. For example, if the `user_home` variable specified `{node[:nginx]...}`, a recipe might look something like this:

```
user_home = "#{node[:nginx][:user]}"

user node[:nginx][:group] do
 gid node[:nginx][:group]
 shell "/bin/bash"
 home user_home
 system true
 action :create
end
```

yum\_package

The **yum\_package** resource is used to install, upgrade, and remove packages with yum for the Red Hat and CentOS platforms. The **yum\_package** resource is able to resolve provides data for packages much like yum can do when it is run from the command line. This allows a variety of options for installing packages, like minimum versions, virtual provides, and library names.

Note

Support for using file names to install packages (as in `yum_package "/bin/sh"`) is not available because the volume of data required to parse for this is excessive.

Note

In many cases, it is better to use the **package** resource instead of this one. This is because when the **package** resource is used in a recipe, the chef-client will use details that are collected by Ohai at the start of the chef-client run to determine the correct package application. Using the **package** resource allows a recipe to be authored in a way that allows it to be used across many platforms. That said, there are scenarios where using an application-specific package is preferred.

Syntax

The syntax for using the **yum\_package** resource in a recipe is as follows:

```
yum_package "name" do
 attribute "value" # see attributes section below
 ...
 action :action
end
```

where

- `yum_package` tells the chef-client to use the `Yum` provider during the chef-client run
- `name` is the name of the resource block; when the `package_name` attribute is not specified as part of a recipe, `name` is also the name of the package
- `attribute` is zero (or more) of the attributes that are available for this resource
- `:action` is the step that the resource will ask the provider to take during the chef-client run

Actions

This resource has the following actions:

Action	Description
<code>:install</code>	Default. Use to install a package. If a version is specified, use to install the specified version of a package.
<code>:upgrade</code>	Use to install a package and/or to ensure that a package is the latest version.
<code>:remove</code>	Use to remove a package.
<code>:purge</code>	Use to purge a package. This action typically removes the configuration files as well as the package.

Attributes

This resource has the following attributes:

Attribute	Description
<code>allow_downgrade</code>	<b>yum_package</b> resource only. Indicates that yum can downgrade a package to satisfy requested version requirements.
<code>arch</code>	The architecture of the package that will be installed or upgraded. (This value can also be passed as part of the package name.)
<code>flush_cache</code>	An array that indicates whether the yum cache should be flushed before or after a yum operation that installs, upgrades, or removes a package. Possible values: <code>:before</code> and <code>:after</code> . Default value: <code>{ :before =&gt; false, :after =&gt; false }</code> .
<code>options</code>	One (or more) additional options that are passed to the command.
<code>package_name</code>	One of the following: the name of a package, the name of a package and its architecture; the name of a dependency. Default value: the <code>name</code> of the resource block (see Syntax section above).



<code>provider</code>	Optional. Use to specify a provider by using its long name. For example, <code>provider Chef::Provider::Long::Name</code> . See the Providers section below for the list of providers available to this resource.
<code>source</code>	Optional. The package source for providers that use a local file.
<code>version</code>	The version of a package to be installed or upgraded.

## Providers

The following providers are available. Use the short name to call the provider from a recipe:

Long name	Short name	Notes
<code>Chef::Provider::Package</code>	<code>package</code>	When this short name is used, the chef-client will attempt to determine the correct provider during the chef-client run.
<code>Chef::Provider::Package::Yum</code>	<code>yum_package</code>	

## Examples

The following examples demonstrate various approaches for using resources in recipes. If you want to see examples of how Chef uses resources in recipes, take a closer look at the cookbooks that Chef authors and maintains: <https://github.com/opscode-cookbooks>.

### Install an exact version

```
yum_package "netpbm = 10.35.58-8.el5"
```

### Install a minimum version

```
yum_package "netpbm >= 10.35.58-8.el5"
```

### Install a minimum version using the default action

```
yum_package "netpbm"
```

### To install a package

```
yum_package "netpbm" do
 action :install
end
```

### To install a partial minimum version

```
yum_package "netpbm >= 10"
```

### To install a specific architecture

```
yum_package "netpbm" do
 arch "i386"
end
```

or:

```
yum_package "netpbm.x86_64"
```

### To install a specific version-release

```
yum_package "netpbm" do
 version "10.35.58-8.el5"
end
```

### To install a specific version (even when older than the current)

```
yum_package "tzdata" do
 version "2011b-1.el5"
 allow_downgrade true
end
```

### Handle `cookbook_file` and `yum_package` resources in the same recipe

When a `cookbook_file` resource and a `yum_package` resource are both called from within the same recipe, dump the cache and use the new repository immediately to ensure that the correct package is installed:

```
cookbook_file "/etc/yum.repos.d/custom.repo" do
 source "custom"
 mode 00644
end

yum_package "only-in-custom-repo" do
 action :install
 flush_cache [:before]
end
```